

Introduction to differential gene expression analysis using RNA-seq

Written by Friederike Dündar, Luce Skrabanek, Paul Zumbo

September, 2015
updated October, 2016

Contents

1	Introduction to RNA-seq	4
1.1	RNA extraction	4
1.1.1	Quality control of RNA preparation (RIN)	5
1.2	Library preparation methods	5
1.3	Sequencing (Illumina)	6
1.4	Experimental Design	9
2	Raw Data (Sequencing Reads)	10
2.1	Download from ENA	10
2.2	Storing sequencing reads: FASTQ format	11
2.3	Quality control of raw sequencing data	14
3	Read Alignment	16
3.1	Reference genomes and annotation	16
3.1.1	File formats for defining genomic regions	17
3.2	Aligning reads using STAR	20
3.3	Storing aligned reads: SAM/BAM file format	22
3.3.1	The SAM file header section	23
3.3.2	The SAM file alignment section	24
3.3.3	Manipulating SAM/BAM files	28
3.4	Quality control of aligned reads	30
3.4.1	Basic alignment assessments	30
3.4.2	Bias identification	34
3.4.3	Quality control with QoRTs	37
4	Read Quantification	38
4.1	Gene-based read counting	38
4.2	Isoform counting methods	39
5	Normalizing and Transforming Read Counts	41
5.1	Normalization for sequencing depth differences	41
5.2	Transformation of sequencing-depth-normalized read counts	43
5.2.1	\log_2 transformation of read counts	43
5.2.2	Transformation of read counts including variance shrinkage	44
5.3	Read count correlations	46
6	Differential Gene Expression Analysis	48
6.1	Running DGE analysis tools	49
6.1.1	DESeq2 workflow	49
6.1.2	Exploratory plots following DGE analysis	50
6.1.3	Exercise suggestions	53
6.1.4	edgeR	53

6.1.5 <code>limma-voom</code>	54
6.2 Judging DGE results	55
7 Appendix	57
7.1 Additional tables	57
7.2 Installing bioinformatics tools on a UNIX server	67

List of Tables

1 Sequencing depth recommendations.	8
2 Illumina's different base call quality score schemes.	13
3 The fields of the alignment section of SAM files.	25
4 The FLAG field of SAM files.	26
5 Programs for DGE.	49
6 Biases and artifacts of Illumina sequencing data.	57
7 FASTQC test modules.	58
8 Optional entries in the header section of SAM files.	60
9 Overview of RSeQC scripts.	61
10 Overview of QoRTs QC functions.	63
11 Normalizing read counts between different conditions.	65
12 Normalizing read counts within the same sample.	66

List of Figures

1 RNA integrity assessment (RIN).	5
2 RNA quality controls before sequencing.	6
3 The different steps of sequencing by synthesis.	7
4 Sequence data repositories.	10
5 Phred score ranges.	13
6 Typical bioinformatics workflow of differential gene expression analysis.	14
7 RNA-seq read alignment.	16
8 Schematic representation of a SAM file.	23
9 CIGAR strings of aligned reads.	28
10 Different modes of counting read-transcript overlaps.	38
11 Schema of a simple <i>deBruijn</i> graph-based transcript determination.	40
12 Effects of different read count normalization methods.	43
13 Comparison of the read distribution plots for untransformed and \log_2 -transformed values.	44
14 Comparison of \log_2 - and $rlog$ -transformed read counts.	45
15 Dendrogram of $rlog$ -transformed read counts.	47
16 PCA on raw counts and $rlog$ -transformed read counts.	47
17 Histogram and MA plot after DGE analysis.	50
18 Heatmaps of \log -transformed read counts.	51
19 Read counts for two genes in two conditions.	53

Technical Prerequisites

Command-line interface The first steps of the analyses – that are the most computationally demanding – will be performed directly on our servers. The interaction with our servers is completely text-based, i.e., there will be no graphical user interface. We will instead be communicating entirely via the command line using the UNIX shell scripting language `bash`. You can find a good introduction into the `shell` basics at http://linuxcommand.org/lc3_learning_the_shell.php (for our course, chapters 2, 3, 5, 7, and 8 are probably most relevant).

To start using the command line, Mac users should use the App called `Terminal`. Windows users need to install `putty`, a Terminal emulator (<http://www.chiark.greenend.org.uk/~sgtatham/putty/download.html>). You probably want the bits under the *A Windows installer for everything except PuTTYtel* heading. Putty will allow you to establish a connection with a UNIX server and interact with it.

Programs that we will be using via the command line:

FASTQC	http://www.bioinformatics.babraham.ac.uk/projects/fastqc/
featureCounts	http://bioinf.wehi.edu.au/subread-package/)
RSeQC	http://rseqc.sourceforge.net/
samtools	http://samtools.sourceforge.net
STAR	https://github.com/alexdobin/STAR
UCSC tools	https://hgdownload.soe.ucsc.edu/admin/exe/

Details on how to install these programs via the command line can be found in the Appendix.

The only program with a graphical user interface will be `IGV`. Go to <https://www.broadinstitute.org/igv/> → “Downloads”, register with your academic email address and launch the Java web start (for Windows machines, you should go for the 1.2 GB version).

R The second part of the analyses – where we will need support for statistics and visualization more than pure computation power – will mostly be done on the individual computers using the programming language R. You can download R for both MacOS and Windows from <http://cran.rstudio.com/>. After you have installed R, we highly recommend to install RStudio (<http://www.rstudio.com/products/rstudio/download/>), which will provide you with an interface to write commands at a prompt, construct a script and view plots all in a single integrated environment.

R packages that will be used throughout the course: `DESeq2`, `edgeR`, `ggplot2`, `limma`, `vsn`

1 Introduction to RNA-seq

The original goal of RNA sequencing was to identify which genomic loci are expressed in a cell (population) at a given time over the entire expression range, i.e., to offer a superior alternative to cDNA microarrays. Indeed, RNA-seq was shown to detect lowly expressed transcripts while suffering from strongly reduced false positive rates in comparison to microarray based expression quantification (Illumina, 2011; Nookaew et al., 2012; Zhao et al., 2014). Since RNA-seq does not rely on a pre-specified selection of cDNA probes, there are numerous additional applications of RNA-seq that go beyond the counting of expressed transcripts of known genes, such as the detection and quantification of non-genic transcripts, splice isoforms, novel transcripts and protein-RNA interaction sites. However, the detection of gene expression changes (i.e., mRNA levels) between different cell populations and/or experimental conditions remains the most common application of RNA-seq.

The general workflow of a differential gene expression analysis is:

1. Sequencing (biochemistry)

- (a) RNA extraction
- (b) Library preparation (including mRNA enrichment)
- (c) Sequencing

2. Bioinformatics

- (a) Processing of sequencing reads (including alignment)
- (b) Estimation of individual gene expression levels
- (c) Normalization
- (d) Identification of differentially expressed (DE) genes

1.1 RNA extraction

Before RNA can be sequenced, it must first be extracted and separated from its cellular environment, which consists primarily of proteins and DNA. The most prevalent methods for RNA isolation are silica-gel based membranes or liquid-liquid extractions with *acidic* phenol-chloroform. In the former case, RNA exclusively binds to a silica-gel membrane, while the remaining cellular components are washed away. Silica-gel membranes require ethanol for binding. The volume of ethanol influences which transcripts are bound to the membrane: more ethanol results in the retention of RNAs <200 bp, whereas a smaller volume results in their loss. When using phenol-chloroform extraction, the cellular components are dissolved into three phases: the organic phase; the interphase; and the aqueous phase, in which the RNA is retained. Phenol-chloroform extraction is typically followed by an alcohol precipitation to de-salt and concentrate the RNA. An alcohol precipitation can be performed with either ethanol or isopropanol, both of which require the use of a salt. Different salts lead to different precipitation efficiencies and result in different RNA populations; e.g., lithium chloride, a commonly used salt, has been reported to result in the loss of tRNAs, 5S rRNAs, snRNAs, and other RNAs <250–300 bp (Cathala et al., 1983). Given the multitude of factors that can influence the outcome of RNA extraction, it is therefore important to process the RNA in a highly controlled and standardized manner, so that the knowledge of how the RNA was isolated can be appropriately leveraged for one's understanding of the data later on. Additional information on how RNA extraction methods influence RNA-seq data can be found in (Sultan et al., 2014).

Although both extraction methods previously mentioned are designed to eliminate DNA contamination, they are often imperfect. But even small amounts of DNA contamination (as little as 0.01% genomic DNA by weight) can negatively impact results (NuGEN, 2013). Accordingly, it is advisable to take additional measures to ensure DNA-free RNA, e.g., by treating the RNA with DNase.

1.1.1 Quality control of RNA preparation (RIN)

Traditionally, RNA integrity was assessed via gel electrophoresis by visual inspection of the ribosomal RNA bands. Intact eukaryotic total RNA should yield clear 28S and 18S rRNA bands. The 28S rRNA band is approximately twice as intense as the 18S rRNA band (2:1 ratio). As RNA degrades, the 2:1 ratio of high quality RNA decreases, and low molecular weight RNA begins to accumulate (Figure 1a). Since the human interpretation of gel images is subjective and has been shown to be inconsistent, Agilent developed a software algorithm that allows for the calculation of an RNA Integrity Number (RIN) from a digital representation of the size distribution of RNA molecules (which can be obtained from an Agilent Bioanalyzer). The RIN number is based on a numbering system from 1 to 10, with 1 being the most degraded and 10 being the most intact (Figure 1b). This approach facilitates the interpretation and reproducibility, of RNA quality assessments, and provides a means by which samples can be compared in a standardized manner.

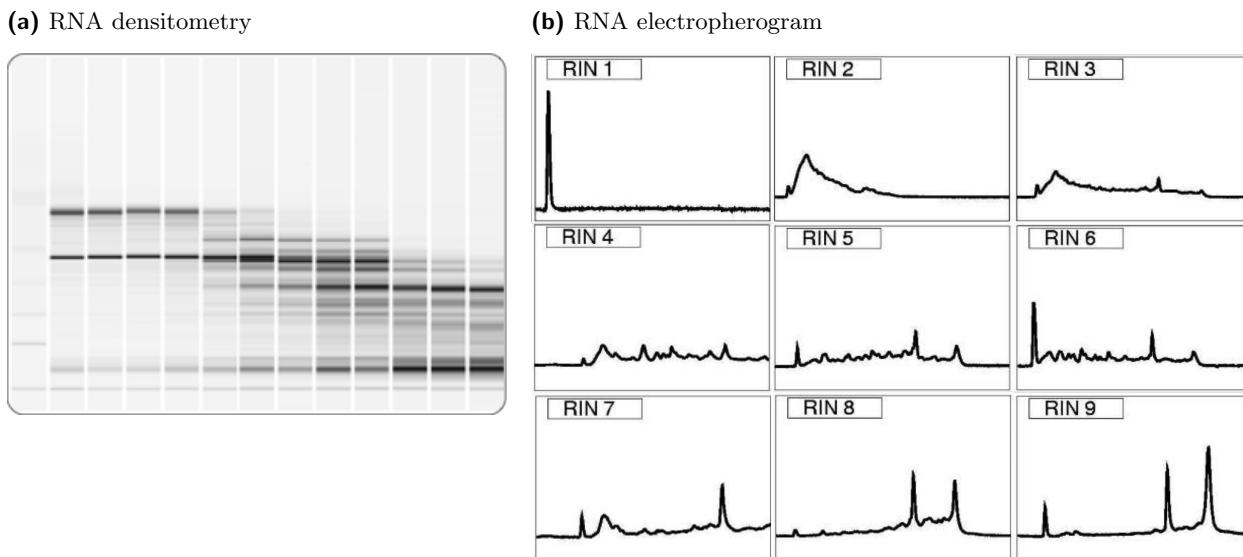


Figure 1: RNA integrity assessment is based on the ratio of $\frac{28S}{18S}$ rRNA, estimated from the band intensity (a) or a densitometry plot (b).

1.2 Library preparation methods

In high-throughput sequencing terms, a *library* is a (preferably random) collection of DNA *fragments* that are ready for sequencing with a specific protocol. For Illumina-based protocols, cDNA fragments will typically be between 150 to 300 bp, and after hybridization to the *flowcell*, the ends (50 to 150 bp) of these fragments will be sequenced.

Due to the numerous types of RNA families, there is a great variety of library preparation protocols. Since the quantification of mRNA is by far the most commonly used application of RNA-seq experiments, we will focus on protocols that are typically applied in this context. Keep in mind that the library preparation can seriously affect the outcome of the sequencing in terms of quality as well as coverage. For more details on library preparation protocols including single-cell RNA-seq, CLIP-seq and more, see Head et al. (2014).

mRNA enrichment Since extracted RNA contains 80–85% rRNA and 10–15% tRNA (Farrell, 2010), mRNA needs to be enriched for. This is typically done either via enrichment of poly(A)-tail containing nucleic acids using oligo(dT) beads or by removal of ribosomal RNA via complementary sequences. Note that various non-polyadenylated RNAs such as histone transcripts and immature mRNAs will not be captured with the poly(A)-enrichment protocol while the “ribo-minus” approach does not exclude unspliced RNAs.

strand-specific sequencing If you need to distinguish overlapping transcripts, e.g., when sequencing prokaryotic transcriptomes or because the aim of the RNA-seq experiment includes the identification of

anti-sense transcripts, the information about which strand a fragment originated from needs to be preserved during library preparation. The most commonly used method incorporates deoxy-UTP during the synthesis of the second cDNA strand (for details see Levin et al. (2010)).

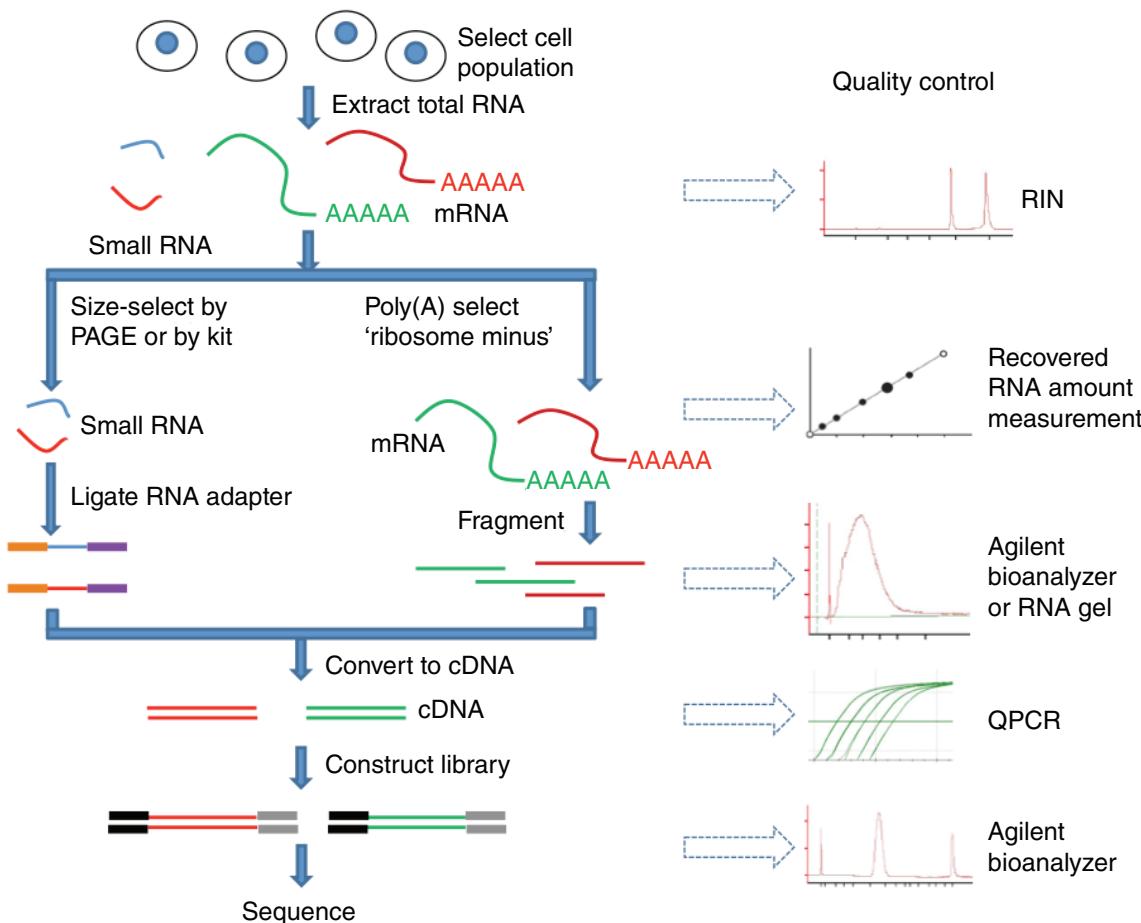


Figure 2: Quality controls during a typical RNA-seq library preparation. After RNA extraction and measuring its integrity, rRNA is depleted (either using poly(A)-selection or rRNA depletion) and the remaining RNA molecules are fragmented, ideally achieving a uniform size distribution. Double-stranded cDNA is synthesized and the adapters for sequencing are added to construct the final library whose fragment size distribution should be unimodal and well-defined. Image taken from Zeng and Mortazavi (2012).



The goal of your RNA-seq experiment will determine the most appropriate library preparation protocol.

1.3 Sequencing (Illumina)

After hybridization of the DNA fragments to the flowcell through means of *adapters*, each fragment is massively and clonally amplified, forming *clusters* of double-stranded DNA. This step is necessary to ensure that the sequencing signal will be strong enough to be detected unambiguously for each base of each fragment. The most commonly used Illumina sequencing protocols will only cover 50 to 100 bp of each fragment (depending on the *read* length that was chosen). The sequencing of the fragment ends is based on fluorophore-labelled dNTPs with reversible terminator elements that will become incorporated and excited by a laser one at a time and thereby enable the optical identification of single bases (Figure 3, Table 3).

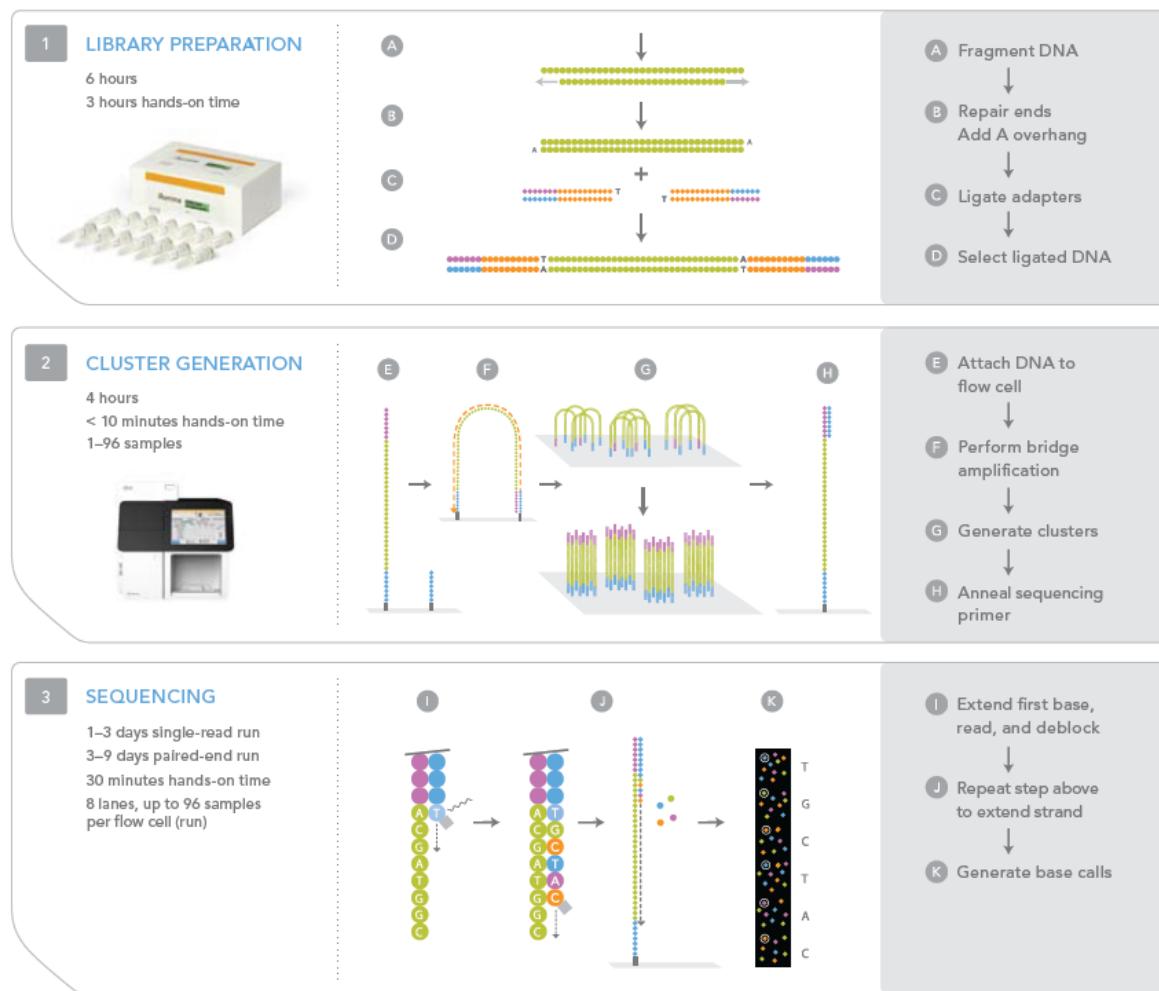


Figure 3: The different steps of sequencing with Illumina's *sequencing by synthesis* method. *Library preparation*: Adapters are ligated to each cDNA fragment to eventually attach them to the flowcell on which they are going to be sequenced. To increase the signal of the sequencing step, every fragment is first clonally amplified after the hybridization onto the flowcell (*cluster generation*). Finally, the nucleotide order of each fragment is revealed through PCR with fluorophore-labelled nucleotides: Images are taken after each round of nucleotide incorporation and bases are identified based on the recorded excitation spectra. Figure from Illumina.

Sequencing depth and coverage Technically, *coverage* refers to the number of reads being sequenced in relation to the genome size, i.e., it is an estimate of how many times each base of the genome is sequenced. For experiments based on the sequencing of the genome, the Lander-Waterman equation is commonly cited as it describes the relationship between coverage, the number of sequenced reads and the genome size:

$$\text{coverage} = \frac{\text{read length} * \text{number of reads}}{\text{haploid genome length}}$$

To identify sequencing errors (and possibly distinguish them from genomic variants), every base should be covered more than once. The coverage value will always be an estimate as the genome is usually not covered uniformly since, for example, euchromatic fragments tend to be overrepresented and standard PCR protocols will favor GC-rich regions and impede AT-rich regions (see Table 6 for more examples of biases that occur with Illumina sequencing platforms).

For RNA-seq, the coverage estimation has rather little practical value as the size of the transcriptome is not known as accurately as the size of the genome, and, more importantly, the per-base coverage will vary drastically between different transcripts depending, most importantly, on their expression. Thus, the number of required reads is determined by the least abundant RNA species of interest. However, it is impossible to know before sequencing how many reads are going to be needed to capture enough fragments of the most

lowly expressed genes. In order to estimate the sequencing depth (= read numbers) needed for a specific RNA-seq experiment, consider the following parameters:

- guidelines from the literature/references (e.g., ENCODE (2011), Sims et al. (2014))
- type of experiment and the type of biological question
- transcriptome size and complexity (many repetitive regions present?)
- error rate of the sequencing platform

See Table 1 for recommended numbers of reads for typical RNA-seq applications. Be aware that, depending on your application, you may want to sequence deeper – consider increasing the number of reads if your goal is to:

- identify lowly expressed genes
- identify very small fold changes between different conditions
- quantify alternative splicing/different isoforms
- detect chimeric transcripts
- detect novel transcripts, transcription start and end sites
- perform *de novo* transcript assembly

Keep in mind that strongly expressed genes and residual rRNA will always account for a large fraction of all reads.

If you are interested in performing power analyses for differential gene expression detection using RNA-seq data, you can have a look at the publication and R code provided by Ching et al. (2014).



In most cases of differential gene expression analysis, it is more important to increase the number of biological replicates than the sequencing depth of single samples (Rapaport et al., 2013; Ching et al., 2014; Liu et al., 2014; Gierliński et al., 2015).

Single read vs. paired-end Single read (SR) sequencing determines the DNA sequence of just one end of each DNA fragment.

Paired-end (PE) sequencing yields both ends of each DNA fragment. PE runs are more expensive (you are generating twice as many DNA reads as with SR), but they increase the mappability for repetitive regions and allow for easier identification of structural variations and indels. They may also increase the precision of studies investigating splicing variants or chimeric transcripts.

Table 1: Recommended sequencing depths for typical RNA-seq experiments for different genome sizes (Genohub, 2015). DGE = differential gene expression, SR = single read, PE = paired-end.

	Small (bacteria)	Intermediate (fruit fly, worm)	Large (mouse, human)
No. of reads for DGE ($\times 10^6$)	5 SR	10 SR	20–50 SR
No. of reads for <i>de novo</i> transcriptome assembly ($\times 10^6$)	30–65 PE	70–130 PE	100–200 PE
Read length (bp)	50	50–100	>100

1.4 Experimental Design

Most RNA-seq experiments aim to identify genes whose expression varies between two or more experimental settings. To distinguish transcription changes caused by the condition being studied from transcription variation caused by differences between individual organisms, cell populations, or experimenters, it is important to perform RNA-seq experiments with sufficient numbers of replicates. There should be enough replicates to a) identify outlier samples and b) be able to remove them without losing too much information about the background variation between transcripts of the same sample type.

Technical replicates Technical replicates can be defined as *different library preparations from the same RNA sample*. They should account for batch effects from the library preparation such as reverse transcription and PCR amplification. To avoid possible lane effects (e.g., differences in the sample loading, cluster amplification, and efficiency of the sequencing reaction), it is good practice to multiplex the same sample over different lanes of the same flowcell. In most cases, technical variability introduced by the sequencing protocol is quite low and well controlled.

Biological replicates There is an on-going debate over what kinds of samples represent true biological replicates. Obviously, the variability between different samples will be greater between RNA extracted from two unrelated humans than between RNA extracted from two different batches of the same cell line. In the latter case, most of the variation that will eventually be detected was probably introduced by the experimenter (e.g., slightly differing media and plating conditions). Nevertheless, this is variation the researcher is typically not interested in assessing, therefore the ENCODE consortium defines biological replicates as *RNA from an independent growth of cells/tissue* (ENCODE (2011)).

Numbers of replicates Currently, most published RNA-seq experiments contain three biological replicates. Based on one of the most exhaustive RNA-seq experiment reported to-date (48 replicates per condition), Schurch et al. (2016) recommend the use of at least six replicates per condition if the focus is on a reliable description of one condition's transcriptome or strongly changing genes between two conditions. If the goal of the experiment is to identify as many differentially expressed genes as possible (including slightly changing ones and those that are lowly expressed), as many as twelve replicates are recommended.

! As a general rule, the more genes with low fold changes that are to be detected, the more replicates are needed to increase the precision of the estimation of the biological variability.

Artificial RNA spike-in If it is important to you to accurately quantify absolute transcript concentrations, you may want to consider to use spike-ins of artificial RNA (such as the ERCC spike-in standard (Risso et al., 2014)). These RNA of known quantities can be used for the calibration of the RNA concentrations in each sample and assess the sensitivity, coverage and linearity of your experiment. Note that different spike-in controls are needed for each type of RNA, but standards are not yet available for all RNA types (ENCODE, 2011).

- ?**
1. What is the main advantage of stranded RNA-seq libraries?
 2. What are advantages of paired-end sequencing for RNA-seq experiments?

2 Raw Data (Sequencing Reads)

Most journals require that any sequencing data related to a manuscript is deposited in a publicly accessible data base. The Sequence Read Archive (SRA) is the main repository for nucleic acid sequences and it has been growing tremendously in the past years. There are three copies of the SRA which are maintained by the NCBI, the European Bioinformatics Institute, and the DNA Databank of Japan (Figure 4), respectively.

International Nucleotide Sequence Database Collaboration

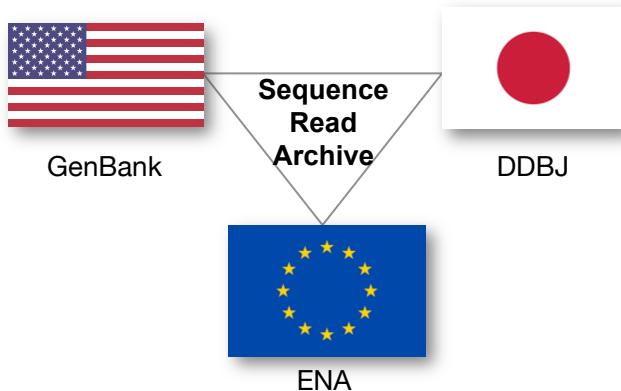


Figure 4: The Sequence Read Archive (SRA) is the largest data base of nucleic acid sequences and all three members of the International Nucleotide Sequencing Database Collaboration (GenBank, the European Nucleotide Archive, the DNA Databank of Japan) maintain instances of it.

2.1 Download from ENA

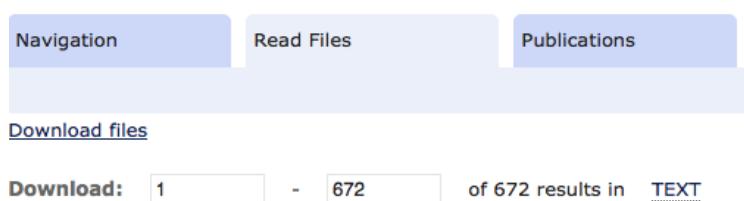
Here, we show you how to download raw sequence data from the European instance of the SRA, which can be accessed via <https://www.ebi.ac.uk/ena>. At ENA, the sequencing reads are directly available in FASTQ format, which will be explained below.

To download a set of FASTQ files:

1. Go to <https://www.ebi.ac.uk/ena>.
2. Search for the accession number of the project, e.g., ERP004763 (should be indicated in the published paper).
3. There are several ways to start the download:
 - (a) Click on the link within the column “Fastq files (ftp)” and save the file of interest. Done.
 - (b) If you prefer the command line, copy the link’s address of the “Fastq files” column (right mouse click), go to the command line, move to the target directory, type:

```
$ wget <link copied from the ENA website>
```

- (c) If there are many samples within one project, you can download the summary of the sample information from ENA by right-clicking on “TEXT” and copying the link location.



```
$ `wget -O samples_at_ENA.txt "<LINK>"` # the quotation marks are important
```

Once you have done this, go to the folder where you will store the data and use the 11th column of the TEXT file ("Fastq files (ftp)") to feed the individual FTP URLs of the different samples to the `wget` command:

```
$ cut -f11 samples_at_ENA.txt > | xargs wget
```

All sequencing data submitted to the SRA (i.e., with an SRA accession number) can also be retrieved through NCBI (<https://www.ncbi.nlm.nih.gov/sra>). Details about the download procedure with NCBI's SRA instance can be found here: <https://www.ncbi.nlm.nih.gov/books/NBK242621/>.

Example data Throughout the course, we will be working with sequencing reads from the most comprehensive RNA-seq dataset to date that contains mRNA from 48 replicates of two *S. cerevisiae* populations: wildtype and *snf2* knock-out mutants (Gierlinski et al., 2015; Schurch et al., 2016). All 96 samples were sequenced on one flowcell (Illumina HiSeq 2000); each sample was distributed over seven lanes, which means that there are seven technical replicates per sample. The accession number for the entire data set (consisting of 7 x 2 x 48 (= 672) raw read files) is ERP004763.



Use the information from the file `ERP004763_sample_mapping.tsv` (from <https://ndownloader.figshare.com/files/2194841>) to download all FASTQ files related to the biological replicates no. 1 of sample type "SNF2" as well as of sample type "WT". Try to do it via the command line and make sure to create two folders (e.g., `SNF2_rep1` and `WT_rep1`) of which each should contain seven FASTQ files in the end.

2.2 Storing sequencing reads: FASTQ format

Currently, raw reads are most commonly stored as FASTQ files. However, details of the file formats may vary widely depending on the sequencing platform, the lab that released the data, or the data repository. For a more comprehensive overview of possible file formats of raw sequencing data, see the NCBI's file format guide: <https://www.ncbi.nlm.nih.gov/books/NBK242622/>.

The FASTQ file format was derived from the simple text format for nucleic acid or protein sequences, FASTA. FASTQ bundles the sequence of every single read produced during a sequencing run together with the quality scores. FASTQ files are uncompressed and quite large because they contain the following information for every single sequencing read:

1. @ followed by the read ID and possibly information about the sequencing run
2. sequenced bases
3. + (perhaps followed by the read ID again, or some other description)
4. quality scores for each base of the sequence (ASCII-encoded, see below)

Again: be aware that this is not a strictly defined file format – variations do exist and may cause havoc!

Here's a real-life example snippet of a FASTQ file downloaded from ENA:

```
1 $ zcat ERR459145.fastq.gz | head
2 @ERR459145.1 DHKW5DQ1:219:DOPT7ACXX:2:1101:1590:2149/1
3 GATCGGAAGAGCGGTTCAGCAGGAATGCCGAGATCGGAAGAGCGGTTCA
4 +
5 @7<DBADDBH?DH>HHHEGHIIIGGIGFGIBFAAGAFHA '5?B@D
6 @ERR459145.2 DHKW5DQ1:219:DOPT7ACXX:2:1101:2652:2237/1
7 GCAGCATCGGCCCTTGCTTGAAGGCAATGTCTTCAGGATCTAAG
8 +
9 @C ; BDDEFGHHHIIIGBHHEHCCHGCGIGGHIGHGIGIIGHIIAHIIIGI
10 @ERR459145.3 DHKW5DQ1:219:DOPT7ACXX:2:1101:3245:2163/1
11 TGCATCTGCATGATCTCAACCATGTCTAAATCCAAATTGTCAGCCTGCGCG
```



For **paired-end** (PE) sequencing runs, there will always be **two** FASTQ files – one for the forward reads, one for the backward reads.

Once you have downloaded the files for a PE run, make sure you understand how the origin of each read (forward or reverse read) is encoded in the read name information as some downstream analysis tools may require you to combine the two files into one.



1. Count the number of reads stored in a FASTQ file.
2. Extract just the quality scores of the first 10 reads of a FASTQ file.
3. Concatenate the two FASTQ files of a PE run.

Sequence identifier The first line of each FASTQ read snippet contains the read ID. Earlier Illumina sequencing platforms (< version 1.8) generated read IDs with the following format:

```
@<machine_id>:<lane>:<tile>:<x_coord>:<y_coord>#<index>/<read_#>
```

Starting from version 1.8 the sequence identifier line has the following format:

```
@<machine_id>:<run number>:<flowcell ID>:<lane>:<tile>:<x-pos>:<y-pos>
<read>:<is filtered>:<control number>:<index sequence>
```

Base call quality scores Illumina sequencing is based on identifying the individual nucleotides by the fluorescence signal emitted upon their incorporation into the growing sequencing read (Figure 3). Once the sequencing run is complete, images taken during each DNA synthesis step are analyzed and the read clusters' fluorescence intensities are extracted and translated into the four letter code. The deduction of nucleotide sequences from the images acquired during sequencing is commonly referred to as *base calling*.

Due to the imperfect nature of the sequencing process and limitations of the optical instruments (see Table 6), base calling will always have inherent uncertainty. This is the reason why FASTQ files store the DNA sequence of each read together with a position-specific quality score that represents the error probability, i.e., how likely it is that an individual base call may be incorrect. The score is called Phred score, Q , which is proportional to the probability p that a base call is incorrect, where $Q = -10 * \log_{10}(p)$. For example, a Phred score of 10 corresponds to one error in every ten base calls ($Q = -10 * \log_{10}(0.1)$), or 90% accuracy; a Phred score of 20 corresponds to one error in every 100 base calls, or 99% accuracy. A higher Phred score thus reflects higher confidence in the reported base.

To assign each base a unique score identifier (instead of numbers of varying character length), Phred scores are typically represented as ASCII characters. At <http://ascii-code.com/> you can see which characters are assigned to what number.

For raw reads, the range of scores will depend on the sequencing technology and the base caller used (Illumina, for example, used a tool called **Bustard**, or, more recently, **RTA**). Unfortunately, Illumina has been anything but consistent in how they a) calculated and b) ASCII-encoded the Phred score (see Table 2 and Figure 5 for the different conventions)! In addition, Illumina now allows Phred scores for base calls with as high as 45, while 41 used to be the maximum score until the HiSeq X. This may cause issues with downstream applications that expect an upper limit of 41.



Note that different base quality assignments exist (Table 2). Try to always make sure you know which version of the Phred score you are dealing with.

The converting of an Illumina FASTQ file version 1.3 (Phred+64) to version 1.8 (Phred+33) can be accomplished with the following **sed** command:

```
1 $ sed -e '4~4y/@ABCDEFGHIJKLMNPQRSTUVWXYZ[\ ]^_`abcdefghijklmnopqrstuvwxyz/!"#$%&'\\'()'*
**,-.\0123456789:;<=>?@ABCDEFGHIJ/' originalFile.fastq
```

Table 2: Base call quality scores are represented with the Phred range. Different Illumina (formerly Solexa) versions used different scores and ASCII offsets. Starting with Illumina format 1.8, the score now represents the standard Sanger/Phred format that is also used by other sequencing platforms and the sequencing archives. Also see Figure 5.

Description	ASCII characters	Quality score		
	Range	Offset	Type	Range
Solexa/early Illumina (1.0)	59 to 126 (; to ~)	64	Solexa	-5 to 62
Illumina 1.3+	64 to 126 (@ to ~)	64	Phred	0 to 62
Sanger standard/Illumina 1.8+	33 to 126 (! to ~)	33	Phred	0 to 93

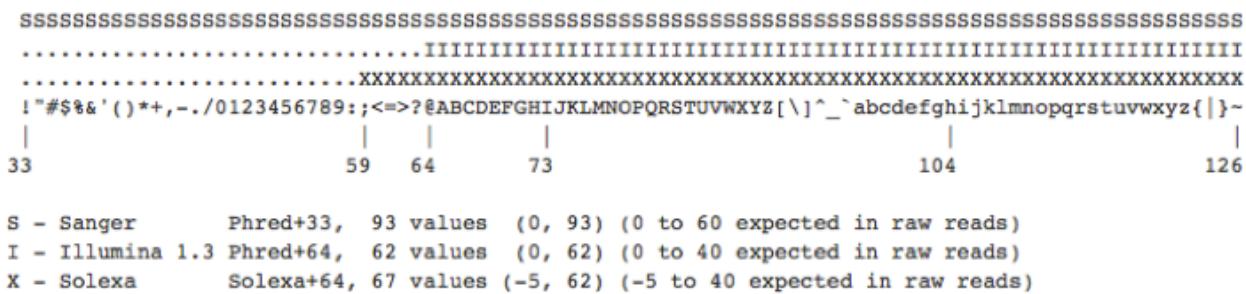


Figure 5: The ASCII interpretation and ranges of the different Phred score notations used by Illumina and the original Sanger interpretation (also see Table 2). Although the Sanger format allows a theoretical score of 93, raw sequencing reads typically do not exceed a Phred score of 60. In fact, most Illumina-based sequencing will result in maximum scores of 41 to 45.

- ?
- Which base call is more likely to be incorrect – one with a Phred score of “#” or one with a Phred score of “;”?
 - Can you guess why the Phred scores are always transformed to ASCII with an offset of at least 33?
 - What is the baseline uncertainty that Illumina grants to its base calling algorithm?
 - How can you find out which Phred score encoding was used in a given FASTQ file?

2.3 Quality control of raw sequencing data

Quality controls should be done at every analysis step. Ideally, quality control should be proactive and comprehensive — see it as a chance to get to know your data which will enable you to perform downstream analyses with appropriate assumptions and parameters. Even if flaws and biases are identified, you may be able to correct those problems *in silico*.

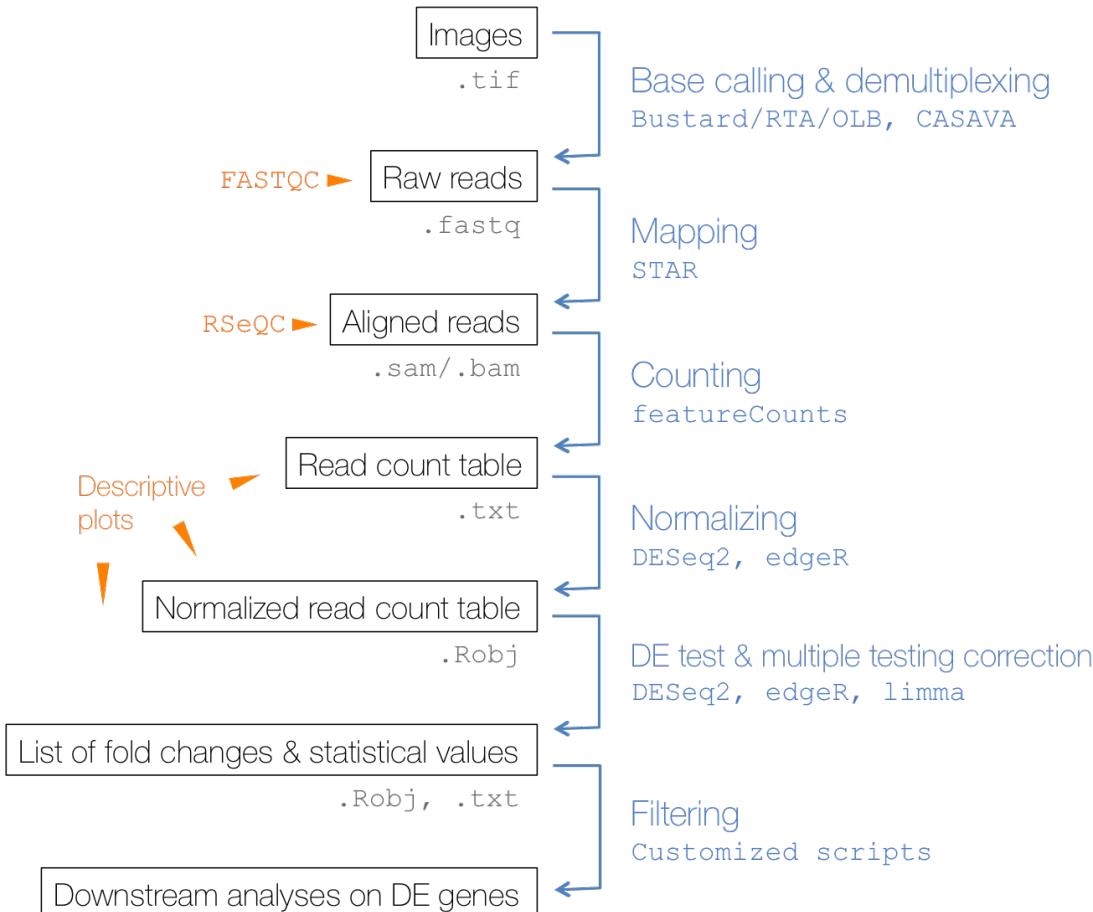


Figure 6: Typical bioinformatics workflow of differential gene expression analysis with commonly used tools. Quality controls are marked in orange, the most commonly used file formats to store the results of each bioinformatic step are indicated in grey.

Since an analysis typically starts with the raw reads (stored in FASTQ files), your first step should be to check the overall quality of the sequenced reads. A poor RNA-seq run will be characterized by one or more of the following types of uninformative sequences:

- PCR duplicates**
- adapter contamination
- rRNA and tRNA reads
- unmappable reads, e.g. from contaminating nucleic acids

All but the last category of possible problems can be detected using a program called **FASTQC**. **FASTQC** is released by the Babraham Institute and can be freely downloaded at <http://www.bioinformatics.babraham.ac.uk/projects/fastqc/>. It performs multiple tests to evaluate the quality scores as well as the sequence composition of the reads stored in a given FASTQ file. Each test is flagged with either “pass”, “warning”, or “fail”, depending on how far the sample deviates from a hypothetical dataset without significant

*It is impossible to distinguish whether identical reads represent PCR duplicates or independent occurrences of the exact same transcript fragment.

bias (see Table 7 for the different tests carried out by FASTQC). Keep in mind though that some sample types are *expected* to have certain biases, so not all “fail” verdicts mean that the sequencing should be repeated.

To run FASTQC, use the following command:

```
1 $ fastqc ERR458493.fastq.gz --extract
```

You can specify as many files to process in a single run as you like (separated by spaces). The results are by default stored in a folder that contains the sample name.

```
1 $ ls ERR458493_fastqc/
2   fastqc_data.txt
3   fastqc fo
4   fastqc_report.html # open this to get a quick visual impression of the
5     results
6   Icons/
7   Images/
8   summary.txt # textual summary
9
10 $ cat ERR458493_fastqc/summary.txt
11 PASS Basic Statistics ERR458493.fastq.gz
12 PASS Per base sequence quality ERR458493.fastq.gz
13 FAIL Per tile sequence quality ERR458493.fastq.gz
14 PASS Per sequence quality scores ERR458493.fastq.gz
15 FAIL Per base sequence content ERR458493.fastq.gz
16 PASS Per sequence GC content ERR458493.fastq.gz
17 PASS Per base N content ERR458493.fastq.gz
18 PASS Sequence Length Distribution ERR458493.fastq.gz
19 WARN Sequence Duplication Levels ERR458493.fastq.gz
20 PASS Overrepresented sequences ERR458493.fastq.gz
21 PASS Adapter Content ERR458493.fastq.gz
22 WARN Kmer Content ERR458493.fastq.gz
```

If you ran FASTQC on more than one file, you may want to combine the plots with the brief text summary to quickly identify outlier samples. The following commands extract all test results that did not pass (`grep -v PASS`) and combines them with all images into a single PNG file using the `montage` tool. All commands are carried out for the sample names stored in `samples.txt` (one sample name per line). `convert` can be used to merge all PNG files into a single PDF file.

```
1 # extract the sample IDs for WT replicate 1
2 $ awk '$3 == "WT" && $4 == 1 {print $1}' data/ERP004763_sample_mapping.tsv >
3   samples.txt
4
5 $ head -n3 samples.txt
6   ERR458493
7   ERR458494
8   ERR458495
9
10 $ while read SAMPLE
11   do
12     grep -v PASS ${SAMPLE}_fastqc/summary.txt | \
13       montage txt:- ${SAMPLE}_fastqc/Images/*png \
14         -tile x3 -geometry +0.1+0.1 -title ${SAMPLE} ${SAMPLE}.png
15   done < samples.txt
16 $ convert *png fastqc_summary.pdf
```

3 Read Alignment

In order to identify the transcripts present in a specific sample as well as their expression levels, the genomic origin of the sequenced cDNA fragments must be determined. The assignment of sequencing reads to the most likely locus of origin is called *read alignment* or *mapping* and it is a crucial step in many of high-throughput sequencing experiments. The general challenge of short read alignment is to map millions of reads accurately and in a reasonable time, despite the presence of sequencing errors, genomic variation and repetitive elements. The different alignment programs employ various strategies that are meant to speed up the process (e.g., by indexing the reference genome) and find a balance between mapping fidelity and error tolerance.

The main challenge of RNA-seq data in particular is the spliced alignment of exon-exon-spanning reads (Figure 7) and the presence of multiple different transcripts (isoforms) of the same gene. Some alignment programs tried to mitigate this problem by aligning to the transcriptome, but this approach is limited to known transcripts and thus heavily dependent on the annotation. Moreover, many reads will overlap with more than one isoform, introducing mapping ambiguity. Thus, the most popular RNA-seq alignment programs (e.g., STAR, TopHat, GSNAP; see Engström et al. (2013) for a review of RNA-seq aligners) nowadays use existing gene annotation for the placement of spliced reads in addition to attempting to identify novel splice events based on reads that cannot be aligned to the reference genome (or transcriptome). The identification of novel splice junctions is based on certain assumptions about transcript structures that may or may not be correct (e.g., most algorithms search for the most parsimonious combination of exons which might not reflect the true biological driving force of isoform generation). Additionally, lowly expressed isoforms may have very few reads that span their specific splice junctions while, conversely, splice junctions that are supported by few reads are more likely to be false positives. Therefore, novel splice junctions will show a bias towards strongly expressed genes. Until reads routinely are sequenced longer, the alignment of spliced reads will therefore remain the most prevalent problems of RNA-seq data.

(a) Aligning to the transcriptome



(b) Aligning to the genome

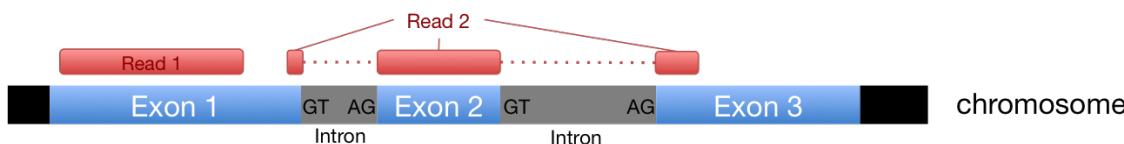


Figure 7: RNA-seq of mRNAs produces two kinds of reads: single exon reads (Read 1) and exon-exon-spanning reads (Read 2). While single exon reads can be aligned equally easily to the genome and to the transcriptome, exon-exon-spanning reads have to be split in order to be aligned properly if only the genome sequence is used as a reference (b).

3.1 Reference genomes and annotation

Genome sequences and annotation are often generated by consortia such as (mod)ENCODE, The Mouse Genome Project, The Berkeley Drosophila Genome Project and many more. The results of these efforts can either be downloaded from individual websites set up by the respective consortia or from more comprehensive data bases such as the one hosted by the University of California, Santa Cruz (UCSC; <https://genome.ucsc.edu/>) or the European genome resource, Ensembl (<http://www.ensembl.org>).

UCSC and Ensembl try to organize, unify and provide access to a wide range of genomes and annotation data. The advantage of downloading data from UCSC (or Ensembl) is that even if you were to work with different species, the file formats and naming conventions will be consistent (and your scripts will be more likely to work). The documentation at <https://genome.ucsc.edu/FAQ/FAQreleases.html> gives a good

overview of the genomes and annotation that are available at UCSC. Unfortunately, UCSC and Ensembl differ in their naming conventions and the frequency of updates.

! Note that UCSC and Ensembl use slightly different naming conventions that can seriously affect downstream analyses. Try to stick to one source.
Always ensure you know exactly which version of a genome and annotation you are working with.

Reference sequences are usually stored in plain text FASTA files that can either be compressed with the generic `gzip` command or, using the tool `faToTwoBit`, into `.2bit` format.

We used the UCSC Genome Browser website to download the reference genome of yeast (go to <https://genome.ucsc.edu/>, click on “Downloads” → “Genome Data” to reach <http://hgdownload.soe.ucsc.edu/downloads.html>, where you will find an overview of all available reference genomes and the respective links).

```

1 # Download genome sequence of S. cerevisiae from UCSC
2 $ wget http://hgdownload.soe.ucsc.edu/goldenPath/sacCer3/bigZips/sacCer3.2bit
3
4 # turning compressed 2bit format into FASTA format
5 $ UCSC-tools/twoBitToFa sacCer3.2bit sacCer3.fa
6
7 $ head sacCer3.fa
>chrI
CCACACCACACCCACACACCCACACACACCACACACACCACACACCACACC
CACACACACACATCTAACACTACCTAACACAGGCCATAATCTAACCTG
GCCAACCTGTCTCTCAACTTACCCCTCATTACCCCTGCCTCCACTCGTTAC
CCTGTCCTATTCAACCATACCACTCCGAACCACCATCCATCCCTCTACTT
ACTACCACTCACCCACCGTTACCCCTCCAATTACCCATATCCAACCCACTG
CCACTTACCCCTACCATACCCCTACCATCCACCATGACCTACTCACCATAC
TGTTCTTCTACCCACCATATTGAAACGCTAACAAATGATCGTAAATAACA
CACACGTGCTTACCCCTACCACTTATACCAACCCACCATGCCATACTCAC
CCTCACTTGTATACTGATTTACGTACGCACACGGATGCTACAGTATATA

```

3.1.1 File formats for defining genomic regions

While the reference sequence is not much more than a very long string of A/T/C/G/N, various file formats exist to store information about the location of transcription start sites, exons, introns etc. All formats agree on having one line per genomic feature, but the nature of the information contained in each row can vary strongly between the formats.

GFF The General Feature Format has nine required fields; the first three fields form the basic `name`, `start`, `end` tuple that allows for the identification of the location in respect to the reference genome (e.g., bases 100 to 1,000 of chromosome 1). Fields must be separated by a single TAB, but no white space. All but the final field in each feature line must contain a value; missing values should be denoted with a ‘?’

There are two versions of the GFF format in use which are similar, but not compatible:

1. GFF version 2 (Sanger Institute; see <http://gmod.org/wiki/GFF2> or <https://www.sanger.ac.uk/resources/software/gff/spec.html>)
2. GFF version 3 (Sequence Ontology Project; see <http://gmod.org/wiki/GFF3>)

GFF2 files use the following fields:

1. **reference sequence**: coordinate system of the annotation (e.g., “Chr1”)
2. **source**: describes how the annotation was derived (e.g., the name of the annotation software)
3. **method**: annotation type (e.g., gene)
4. **start position**: 1-based integer, always less than or equal to the stop position

5. **stop position**: for zero-length features, such as insertion sites, start equals end and the implied site is to the right of the indicated base
6. **score**: e.g., sequence identity
7. **strand**: “+” for the forward strand, “-” for the reverse strand, or “.” for annotations that are not stranded
8. **phase**: codon phase for annotations linked to proteins; 0, 1, or 2, indicating the frame, or the number of bases that should be removed from the beginning of this feature to reach the first base of the next codon
9. **group**: contains the class and ID of an annotation which is the logical parent of the current one (“feature is composed of”)

GFF3 files (asterisk denotes difference to GFF2)

1. **reference sequence**
2. **source**
3. **type***: constrained to be either: (a) a term from the “lite” sequence ontology, SOFA; or (b) a SOFA accession number.
4. **start position**
5. **stop position**
6. **score**
7. **strand**
8. **phase**
9. **attributes***: list of feature attributes as TAG=VALUE pairs; spaces are allowed in this field, multiple TAG=VALUE pairs are separated by semicolons; the TAGS have predefined meanings:
 - ID (must be unique)
 - Name (display name)
 - Alias (secondary name)
 - Parent
 - Target (the format of the value is “target_id start end [strand]”)
 - Gap (in CIGAR format)
 - Derives_from (database cross reference)
 - Ontology_term

```

1 # GFF-version 2
2 IV curated exon 5506900 5506996 . + . Transcript B0273.1
3 IV curated exon 5506026 5506382 . + . Transcript B0273.1
4 IV curated exon 5506558 5506660 . + . Transcript B0273.1
5 IV curated exon 5506738 5506852 . + . Transcript B0273.1
6
7 # GFF-version 3
8 ctg123 . exon 1300 1500 . + . ID=exon00001
9 ctg123 . exon 1050 1500 . + . ID=exon00002
10 ctg123 . exon 3000 3902 . + . ID=exon00003
11 ctg123 . exon 5000 5500 . + . ID=exon00004
12 ctg123 . exon 7000 9000 . + . ID=exon00005

```

GTF The Gene Transfer Format is based on the GFF, but is defined more strictly. (It is sometimes referred to as GFF2.5 because the first eight GTF fields are the same as GFF2, but, as for GFF3, the 9th field has been expanded into a list of attributes.) Contrary to GFF files, the TYPE VALUE pairs of GTF files are separated by one space and must end with a semi-colon (followed by exactly one space if another attribute is added afterwards):

```

1 # example for the 9th field of a GTF file
2 gene_id "Em:U62.C22.6"; transcript_id "Em:U62.C22.6.mRNA"; exon_number 1

```

The **gene_id** and **transcript_id** values are globally unique identifiers for the genomic locus of the transcript or the same transcript itself and must be the first two attributes listed. Textual attributes should be surrounded by double quotes.

```

1 # GTF example
2 chr1 HAVANA gene 11869 14412 . + . gene_id "ENSG00000223972.4";
   transcript_id "ENSG00000223972.4"; gene_type "pseudogene"; gene_status "
   KNOWN"; gene_name "DDX11L1"; transcript_type "pseudogene"; transcript_status
   "KNOWN"; transcript_name "DDX11L1"; level 2; havana_gene "
   OTTHUMG00000000961.2";
3 chr1 HAVANA transcript 11869 14409 . + . gene_id "ENSG00000223972.4";
   transcript_id "ENST00000456328.2"; gene_type "pseudogene"; gene_status "
   KNOWN"; gene_name "DDX11L1"; transcript_type "processed_transcript";
   transcript_status "KNOWN"; transcript_name "DDX11L1-002"; level 2; tag "
   basic"; havana_gene "OTTHUMG00000000961.2"; havana_transcript "
   OTTHUMT00000362751.1";

```

More information on GTF format can be found at <http://mblab.wustl.edu/GTF2.html> (or, for the most recent version: <http://mblab.wustl.edu/GTF22.html>).

The following screenshot illustrates how you can, for example, download a GTF file of yeast transcripts from the UCSC Genome Table Browser (<https://genome.ucsc.edu/cgi-bin/hgTables>).

Table Browser

Use this program to retrieve the data associated with a track in text format, to calculate intersections between tracks, and to retrieve DNA sequence covered by a track. For help in using this application see [Using the Table Browser](#) for a description of the controls in this form, the [User's Guide](#) for general information and sample queries, and the OpenHelix Table Browser [tutorial](#) for a narrated presentation of the software features and usage. For more complex queries, you may want to use [Galaxy](#) or our [public MySQL server](#). To examine the biological function of your set through annotation enrichments, send the data to [GREAT](#). Send data to [GenomeSpace](#) for use with diverse computational tools. Refer to the [Credits](#) page for the list of contributors and usage restrictions associated with these data. All tables can be downloaded in their entirety from the [Sequence and Annotation Downloads](#) page.

The screenshot shows the UCSC Table Browser interface with the following search parameters:

- clade:** Other
- genome:** S. cerevisiae
- assembly:** Apr. 2011 (SacCer_Apr2011/sacCer3)
- group:** Genes and Gene Predictions
- track:** SGD Genes
- table:** sgdGene
- region:** genome (radio button selected)
- region value:** chrIV:765966-775965
- identifiers (names/accessions):** (empty)
- filter:** (empty)
- intersection:** (empty)
- correlation:** (empty)
- output format:** GTF - gene transfer format
- Send output to:** Galaxy, GREAT, GenomeSpace (checkboxes)
- output file:** sacCer3.gtf (text input field)
- file type returned:** plain text (radio button selected)

At the bottom are two buttons: **get output** and **summary/statistics**.

! GTF files downloaded from the UCSC Table Browser have the same entries for `gene_id` and `transcript_id`. This can lead to problems with downstream analysis tools that expect exons of different isoforms to have the same `gene_id`, but different `transcript_ids`.

Here's a way to get a properly formatted GTF file (i.e., with different entries for `gene_name` and `transcript_id`) of RefSeq genes using the UCSC tool `genePredToGtf`:

```
1 # first, download a table for "Genes and Gene Predictions" from the UCSC Table
  Browser indicating as the output format: "all fields from selected table"
2 # NOTE: this may not work for all GTF files downloaded from UCSC! genePredToGtf
  is very finicky and every organism's annotation may have been generated and
  deposited by a different person)
3 $ head -n1 allfields_hg19.txt
4 bin      name      chrom      strand      txStart      txEnd      cdsStart      cdsEnd
        exonCount      exonStarts      exonEnds      score      name2      cdsStartStat
        cdsEndState      exonFrames
5 # remove first column and first line, feed that into genePredToGtf
6 $ cut -f 2- allfields_hg19.txt | sed '1d' | \
  genePredToGtf file stdin hg19_RefSeq.gtf
7 $ head -n1 hg19_RefSeq.gtf
8 chr1  stdin  exon  66999639  67000051  . + .  gene_id "SGIP1"; transcript_id "
  NM_032291"; exon_number "1"; exon_id "NM_032291.1"; gene_name "SGIP1";
```

BED format The BED format is the simplest way to store annotation tracks. It has three required fields (`chromosome`, `start`, `end`) and up to 9 optional fields (`name`, `score`, `strand`, `thickStart`, `thickEnd`, `itemRgb`, `blockCount`, `blockSizes`, `blockStarts`). The number of fields per line can thus vary from three to twelve, but must be consistent within a file and must obey the order, i.e. lower-numbered fields must always be populated if higher-numbered fields are used. Fields seven to twelve are only necessary if regions should be drawn in a Genome Browser with the typical appearance known for gene tracks. Note that the BED format indicates a region with 0-based start position and 1-based end position (GTF/GFF are 1-based in both positions).

```
1 # 6-column BED file defining transcript loci
2 chr1  66999824  67210768  NM_032291  0 +
3 chr1  33546713  33586132  NM_052998  0 +
4 chr1  25071759  25170815  NM_013943  0 +
5 chr1  48998526  50489626  NM_032785  0 -
```

- 1. Which annotation data base is currently recommended for poly(A)-enriched RNA-seq data?
- 2. Which annotation data base would you use for RNA-seq of total RNA?
- 3. How many non-coding RNA transcripts does the Ensembl annotation for the human reference hg19 contain? Find out via the command line.



! Obtaining a correctly GTF file may be one of the most difficult tasks in the entire analysis! Do take this seriously and invest the time to make sure that the GTF file you are using is correctly formatted. Do not take the risk of introducing strange results (which you may not notice) that are due to formatting issues only!

3.2 Aligning reads using STAR

Numerous alignment programs have been published in the past (and will be published in the future), and depending on your specific project, some aligners may be preferred over others. For straight-forward RNA-seq data, **STAR** (Dobin et al., 2013) has been shown to be very efficient and reasonably sensitive. The main caveat is the large number of putative novel splice sites that should be regarded with caution (Engström et al., 2013). Another popular aligner is **TopHat**, which is basically a sophisticated wrapper around the genomic aligner **Bowtie** (Kim et al., 2013).

Shown here are the example commands for the alignment of sequencing reads to the *S. cerevisiae* genome.

- 1. Generate genome index** This step has to be done only once per genome type (and alignment program). The index files will comprise the genome sequence, suffix arrays, chromosome names and lengths, splice junctions coordinates, and information about the genes. Therefore, the main input for this step encompasses the reference genome and an annotation file.

```

1 # create a directory to store the index in
2 $ REF_DIR=mat/referenceGenomes/S_cerevisiae
3 $ mkdir ~/STARindex
4
5 # set a variable for STAR access
6 $ runSTAR=mat/software/STAR-2.5.2a/bin/Linux_x86_64/STAR
7
8 # Run STAR in "genomeGenerate" mode
9 $ ${runSTAR} --runMode genomeGenerate \
10   --genomeDir ~/STARindex \ # index will be stored there
11   --genomeFastaFiles ${REF_DIR}/sacCer3.fa \ # reference genome sequence
12   --sjdbGTFfile ${REF_DIR}/sacCer3.gtf \ # annotation file
13   --sjdbOverhang 49 # should be read length minus 1 ; length of the
14     genomic sequence around the annotated junction to be used for the
       splice junctions database
      --runThreadN 1 \ # can be used to define more processors

```

- 2. Alignment** This step has to be done for each individual FASTQ file.

For the particular data set used here, each sample was distributed over seven flow cell lanes, i.e., each sample has seven separate FASTQ files. Unlike most aligners, STAR will merge those files on the fly if multiple input file names are indicated. The file names must be separated by a comma without whitespaces.

```

1 # make a folder to store the STAR output in
2 $ mkdir alignment_STAR
3
4 # list fastq.gz files separated by comma without whitespaces
5 $ ls -m raw_reads/*fastq.gz | sed 's/ //g'
6
7 # execute STAR in the default runMode "alignReads"
8 # NOTE: the default may not be optimal for your application!
9 # please, read the STAR manual and decide which parameters are suitable for
10   your data set!
11 $ ${runSTAR} --genomeDir ${REF_DIR}/STARindex/ \
12   --readFilesIn raw_reads/ERR458493.fastq.gz,raw_reads/ERR458494.fastq.gz,
13     raw_reads/ERR458495.fastq.gz,raw_reads/ERR458496.fastq.gz,raw_reads/
14       ERR458497.fastq.gz,raw_reads/ERR458498.fastq.gz,raw_reads/ERR458499.
15         fastq.gz \
16   --readFilesCommand zcat \ # necessary because of gzipped fastq files
17   --outFileNamePrefix alignment_STAR/WT_1_ \
18   --outFilterMultimapNmax 1 \ # only reads with 1 match in the reference
19     will be returned as aligned
20   --outReadsUnmapped Fastx \ # will generate an extra output file with the
21     unaligned reads
22   --outSAMtype BAM SortedByCoordinate \
23   --twoPassMode \ # STAR will perform mapping, then extract novel junctions
24     which will be inserted into the genome index which will then be used
25       to re-map all reads
26   --runThreadN 1 # can be increased if sufficient computational power is
27     available

```

- 3. BAM file indexing** Most downstream applications will require a .BAM.BAI file together with every BAM file to quickly access the BAM files without having to load them into memory. To obtain these index files, simply run the `samtools index` command for each BAM file once the mapping is finished.

```
1 # export samtools path (for convenience)
2 $ export PATH=/zenodotus/abc/store/courses/2016_rnaseq/software/samtools
   -1.3.1:$PATH
3
4 # index the BAM file
5 $ samtools index alignment_STAR/WT_1_Aligned.sortedByCoord.out.bam
```

STAR has more than 100 parameters, which are all described in its manual (which comes as part of the software and can be found at <https://github.com/alexdobin/STAR/tree/master/doc>). While the command we show above will work well for most applications (although there's one catch as you will see later on!), we strongly recommend you familiarize yourself with the STAR manual. The most important points are:

- handling of multi-mapped reads (e.g., how the best alignment score is assigned and the number and order in which secondary alignments are reported);
- optimization for very small genomes;
- defining the minimum and maximum intron sizes that are allowed (the default setting for the maximum intron size is 1,000,000 bp);
- handling of genomes with more than 5,000 scaffolds (usually reference genomes in a draft stage);
- using STAR for the detection of chimeric (fusion) and circular transcripts.



Which STAR options shown above:

- ... have to be different for every sample that you map?
- ... should remain consistent for all samples of one analysis?
- ... will affect the number of reads in the final output file?

3.3 Storing aligned reads: SAM/BAM file format

The output option of **STAR** already indicates that the results of the alignment will be stored in a **SAM** or **BAM** file. The Sequence Alignment/Map (**SAM**) format is, in fact, a generic nucleotide alignment format that describes the alignment of sequencing reads (or *query sequences*) to a reference. The human readable, TAB-delimited **SAM** files can be compressed into the Binary Alignment/Map format. These **BAM** files are bigger than simply gzipped **SAM** files, because they have been optimized for fast random access rather than size reduction. Position-sorted **BAM** files can be indexed so that all reads aligning to a locus can be efficiently retrieved without loading the entire file into memory.

To convert a **BAM** file into a **SAM** file, use **samtools view**:

```
1 $ samtools view -h WT_1_Aligned.sortedByCoord.out.bam > WT_1_Aligned.
   sortedByCoord.out.sam
```

As shown in Figure 8, **SAM** files typically contain a short header section and a very long alignment section where each row represents a single read alignment. The following sections will explain the **SAM** format in a bit more detail. For the most comprehensive and updated information go to <https://github.com/samtools/hts-specs>.

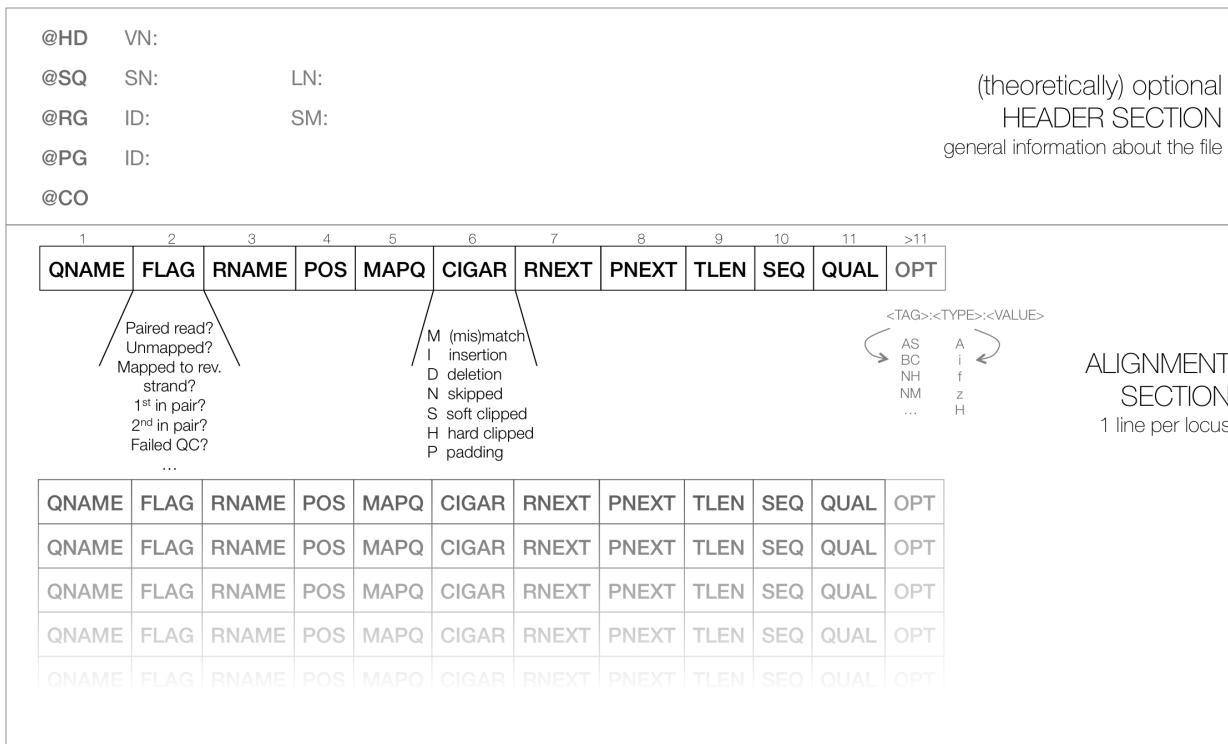


Figure 8: Schematic representation of a SAM file. Each line of the optional header section starts with “@”, followed by the appropriate abbreviation (e.g., SQ for sequence dictionary which lists all chromosomes names (SN) and their lengths (LN)). See Table 8 for all possible entries and tags. The vast majority of lines within a SAM file typically correspond to read alignments where each read is described by the 11 mandatory entries (black font) and a variable number of optional fields (grey font). See Section 3.3.2 for more details.

3.3.1 The SAM file header section

The header section includes information about how the alignment was generated and stored. All lines in the header section are tab-delimited and begin with the “@” character, followed by `tag:value` pairs, where `tag` is a two-letter string that defines the content and the format of `value`. For example, the “@SQ” line in the header section contains the information about the names and lengths of the *reference sequences* to which the reads were aligned. For a hypothetical organism with three chromosomes of length 1,000 bp, the SAM header should contain the following three lines:

```
@SQ SN:chr1 LN:1000
@SQ SN:chr2 LN:1000
@SQ SN:chr3 LN:1000
```

`samtools view` can be used to retrieve a SAM or BAM file’s header. The output from the following example was slightly modified for better readability. See Table 8 for more information about the entries typically stored within the header section.

```

1 # The default behavior of samtools view is to not show the header section.
2 # samtools view -h will show both header and alignment section;
3 # samtools view -H will return the header section only
4
5 $ samtools view -H Sample1_Aligned.sortedByCoord.out.bam
6 @HD VN:1.4
7
8 @SQ SN:chrI LN:230218
9 @SQ SN:chrII LN:813184
10 @SQ SN:chrIII LN:316620
11 @SQ SN:chrIV LN:1531933
12 @SQ SN:chrV LN:576874
13
14 @PG ID:STAR VN:STAR_2.4.0e CL:STAR --runThreadN 8 --genomeDir STAR-sacCer3
   --readFilesIn Lane1.fastq.gz,Lane2.fastq.gz,Lane3.fastq.gz,Lane4.fastq.gz,
   Lane5.fastq.gz,Lane6.fastq.gz,Lane7.fastq.gz --readFilesCommand zcat --
   outFileNamePrefix Sample1_ --outSAMtype BAM SortedByCoordinate --
   outSAMunmapped Within --outFilterMultimapNmax 1
15
16 @CO user command line: STAR --genomeDir STAR-sacCer3 --readFilesIn Lane1.
   fastq.gz,Lane2.fastq.gz,Lane3.fastq.gz,Lane4.fastq.gz,Lane5.fastq.gz,Lane6
   .fastq.gz,Lane7.fastq.gz --readFilesCommand zcat --outFileNamePrefix
   Sample1_ --outFilterMultimapNmax 1 --outSAMunmapped Within --runThreadN 8
   --outSAMtype BAM SortedByCoordinate

```

3.3.2 The SAM file alignment section

The optional header section is followed by the alignment section where each line corresponds to one sequenced read. For each read, there are 11 mandatory fields that always appear in the same order:

<QNAME> <FLAG> <RNAME> <POS> <MAPQ> <CIGAR> <MRNM> <MPOS> <ISIZE> <SEQ> <QUAL>

If the corresponding information is unavailable or irrelevant, field values can be ‘0’ or ‘*’ (depending on the field, see Table 3), but they cannot be missing! After the 11 mandatory fields, a variable number of optional fields can be present (Figure 8).

Here’s an example of one single line of a real-life SAM file:

```

1 ERR458493.552967 16 chrI 140 255 12M61232N37M2S * 0 0
   CCACTCGTCACCAGGGCCGGCGGGCTGATCACTTATCGTCATCTGGC BB?
   HHJJIGHHJIGIIJJIJGIJIIJIIIGHBJJJJHHHHFFDDDA1+B NH:i:1 HI:i:1 AS:i:41 nM:
   i:2

```

The following table explains the format and content of each field. The FLAG, CIGAR, and the optional fields (marked in blue) are explained in more detail below.

Table 3: Overview of the fields that are required for each row of a SAM file’s alignment section. The number of optional fields can vary widely between different SAM files and even between reads within in the same file. The field types marked in blue are explained in more detail in the main text below.

Pos.	Field	Example entry	Description	NA value
1	QNAME	Read1	Query template (= read) name (PE: read pair name)	required
2	FLAG	83	Information about the read’s mapping properties encoded as bit-wise flags (see next section and Table 4).	required
3	RNAME	chrI	Reference sequence name. This should match a @SQ line in the header.	*
4	POS	15364	1-based leftmost mapping position of the first matching base. Set as 0 for an unmapped read without coordinates.	0
5	MAPQ	30	Mapping quality of the alignment. Should be a Phred-scaled posterior probability that the position of the read is incorrect, but the value is completely dependent on the alignment program. Some tools set this to 0 if multiple alignments are found for one read.	0
6	CIGAR	51M	Detailed information about the alignment (see below).	*
7	RNEXT	=	PE reads: reference sequence name of the next read. Set to “=” if both mates are mapped to the same chromosome.	*
8	PNEXT	15535	PE reads: leftmost mapping position of the next read.	0
9	TLEN	232	PE reads: inferred template length (fragment size).	0
10	SEQ	CCA...GGC	The sequence of the aligned read on the forward strand (not including indels).	*
11	QUAL	BBH...1+B	Base quality (same as the quality string in the FASTQ format, but always in Sanger format [ASCII+33]).	*
12ff	OPT	NM:i:0	Optional fields (format: <TAG>:<TYPE>:<VALUE>; see below).	

FLAG field The FLAG field encodes various pieces of information about the individual read, which is particularly important for PE reads. It contains an integer that is generated from a sequence of Boolean bits (0, 1). This way, answers to multiple binary (Yes/No) questions can be compactly stored as a series of bits, where each of the single bits can be addressed and assigned separately.

Table 4 gives an overview of the different properties that can be encoded in the FLAG field. The developers of the SAM format and `samtools` tend to use the hexadecimal encoding as a means to refer to the different bits in their documentation. The value of the FLAG field in a given SAM file, however, will always be the decimal representation of the sum of the underlying binary values (as shown in Table 3, row 2).

Table 4: The FLAG field of SAM files stores several information about the respective read alignment in one single decimal number. The decimal number is the sum of all the answers to the Yes/No questions associated with each binary bit. The hexadecimal representation is used to refer to the individual bits (questions).

Binary (Decimal)	Hex	Description
00000000001 (1)	0x1	Is the read paired?
00000000010 (2)	0x2	Are both reads in a pair mapped “properly” (i.e., in the correct orientation with respect to one another)?
00000000100 (4)	0x4	Is the read itself unmapped?
00000001000 (8)	0x8	Is the mate read unmapped?
00000010000 (16)	0x10	Has the read been mapped to the reverse strand?
00000100000 (32)	0x20	Has the mate read been mapped to the reverse strand?
00001000000 (64)	0x40	Is the read the first read in a pair?
00010000000 (128)	0x80	Is the read the second read in a pair?
00100000000 (256)	0x100	Is the alignment not primary? (A read with split matches may have multiple primary alignment records.)
01000000000 (512)	0x200	Does the read fail platform/vendor quality checks?
10000000000 (1024)	0x400	Is the read a PCR or optical duplicate?

A bit is set if the corresponding state is true. For example, if a read is paired, 0x1 will be set, returning the decimal value of 1. Therefore, all FLAG values associated with paired reads must be uneven decimal numbers. Conversely, if the 0x1 bit is unset (= read is not paired), no assumptions can be made about 0x2, 0x8, 0x20, 0x40 and 0x80.

In a run with single reads, the flags you will most commonly see are:

- 0: This read has been mapped to the forward strand. (None of the bit-wise flags have been set.)
- 4: The read is unmapped (0x4 is set).
- 16: The read is mapped to the reverse strand (0x10 is set).

(0x100, 0x200 and 0x400 are not used by most aligners, but could, in principle be set for single reads.)

Some common FLAG values that you may see in a PE experiment include:

69	$(= 1 + 4 + 64)$	The read is paired, is the first read in the pair, and is unmapped.
77	$(= 1 + 4 + 8 + 64)$	The read is paired, is the first read in the pair, both are unmapped.
83	$(= 1 + 2 + 16 + 64)$	The read is paired, mapped in a proper pair, is the first read in the pair, and it is mapped to the reverse strand.
99	$(= 1 + 2 + 32 + 64)$	The read is paired, mapped in a proper pair, is the first read in the pair, and its mate is mapped to the reverse strand.
133	$(= 1 + 4 + 128)$	The read is paired, is the second read in the pair, and it is unmapped.
137	$(= 1 + 8 + 128)$	The read is paired, is the second read in the pair, and it is mapped while its mate is not.
141	$(= 1 + 4 + 8 + 128)$	The read is paired, is the second read in the pair, but both are unmapped.
147	$(= 1 + 2 + 16 + 128)$	The read is paired, mapped in a proper pair, is the second read in the pair, and mapped to the reverse strand.
163	$(= 1 + 2 + 32 + 128)$	The read is paired, mapped in a proper pair, is the second read in the pair, and its mate is mapped to the reverse strand.

Note that the strand information of the **FLAG** field (0x10) does not necessarily indicate the strand of the original transcript. Unless a strand-specific RNA-seq library protocol was used, this only tells you which strand of the ds-cDNA fragment was sequenced.

A useful website for quickly translating the **FLAG** integers into plain English explanations like the ones shown above is: <https://broadinstitute.github.io/picard/explain-flags.html>

- 
1. How can you retrieve just the alignment section of a **BAM** file?
 2. What does a **MAPQ** value of 20 mean?
 3. What does a **FLAG** value of 2 mean?
 4. Would you be happy or sad if your paired-end read alignments all had **FLAG** values of 77 or 141?
 5. Your favorite read pair has **FLAG** values of 153 and 69. Which read aligned to the forward strand of the reference?

CIGAR [Concise Idiosyncratic Gapped Alignment Report] String The sixth field of a **SAM** file contains a so-called **CIGAR** string indicating which *operations* were necessary to map the read to the reference sequence at that particular locus.

The following operations are defined in **CIGAR** format (also see Figure 9):

- M Alignment (can be a sequence match or mismatch!)
- I Insertion in the read compared to the reference
- D Deletion in the read compared to the reference
- N Skipped region from the reference. For mRNA-to-genome alignments, an N operation represents an intron. For other types of alignments, the interpretation of N is not defined.
- S Soft clipping (clipped sequences are present in read); S may only have H operations between them and the ends of the string
- H Hard clipping (clipped sequences are NOT present in the alignment record); can only be present as the first and/or last operation
- P Padding (silent deletion from padded reference)
- = Sequence match (not widely used)
- X Sequence mismatch (not widely used)

The sum of lengths of the M, I, S, =, X operations must equal the length of the read.

Reference sequence with aligned reads	CIGAR string	Explanation
C T G C A T G T T A G A T A A * * G A T A G C T G T G C T A A A G G A T A * C T G G A T A A * G G A T A T G T T A [red box] a a a C A T G T T A G A A A C A T G T T A G	1M2I4M1D3M 5M1P1I4M 5M15N5M 3S8M 3H8M	Insertion & Deletion Padding & Insertion Spliced read Soft clipping Hard clipping

Figure 9: Image based on a figure from Li et al. (2009).

OPT field(s) Following the eleven mandatory SAM file fields, the optional fields are presented as key-value pairs in the format of `<TAG>:<TYPE>:<VALUE>`, where TYPE is one of:

- A Character
- i Integer
- f Float number
- Z String
- H Hex string

The information stored in these optional fields will vary widely depending on the mapper and new tags can be added freely. In addition, reads within the same SAM file may have different numbers of optional fields, depending on the program that generated the SAM file. Commonly used optional tags include:

- AS:i Alignment score
- BC:Z Barcode sequence
- HI:i Match is *i*-th hit to the read
- NH:i Number of reported alignments for the query sequence
- NM:i Edit distance of the query to the reference
- MD:Z String that contains the exact positions of mismatches (should complement the CIGAR string)
- RG:Z Read group (should match the entry after ID if @RG is present in the header).

Thus, for example, we can use the NM:i:0 tag to select only those reads which map perfectly to the reference (i.e., have no mismatches).

While the optional fields listed above are fairly standardized, tags that begin with X, Y, and Z are reserved for particularly free usage and will never be part of the official SAM file format specifications. XS, for example, is used by TopHat to encode the strand information (e.g., XS:A:+) while Bowtie2 and BWA use XS:i: for reads with multiple alignments to store the alignment score for the next-best-scoring alignment (e.g., XS:i:30).

3.3.3 Manipulating SAM/BAM files

As indicated above, `samtools` is a powerful suite of tools designed to interact with SAM and BAM files (Li et al., 2009).

```

1 # return a peek into a SAM or BAM file (note that a SAM file can also easily be
  inspected using the basic UNIX commands for any text file, such as cat,
  head, less etc.)
2 $ samtools view InFile.bam | head
3
4 # turn a BAM file into the human-readable SAM format (including the header)
5 $ samtools view -h InFile.bam > InFile.sam
6
7 # compress a SAM file into BAM format (-Sb is equivalent to -S -b)
8 $ samtools view -Sb InFile.sam > OutFile.bam
9
10 # generate an index for a BAM file (needed for many downstream tools)
11 $ samtools index InFile.bam

```

To see all the operations that can be done using `samtools`, type `samtools --help`.

The myriad information stored within the alignment files allow you to focus on virtually any subset of read alignments that you may be interested in. The `samtools view` tool has many options that directly interpret some of the mandatory fields of its alignment section (Table 3), such as the mapping quality, the location and the FLAG field values.

```

1 # get only unmapped reads
2 $ samtools view -h \ # show header
3     -b \ # output a BAM file
4     -f 4 \ # include only reads where the 0x4 bit is set
5     Aligned.sortedByCoord.out.bam > unmapped_reads.bam
6
7 # get only mapped reads
8 $ samtools view -hb -F 4 \ # include only reads where the 0x4 bit is NOT set
9     Aligned.sortedByCoord.out.bam > mapped_reads.bam
10
11 # skip read alignments with mapping quality below 20
12 samtools view -h -b -q 20 Aligned.sortedByCoord.out.bam > high_mapq_reads.bam

```

If you would like to filter an alignment file based on any of the optional tags, you will have to resort to means outside `samtools`. Looking for exact matches using `grep` can be particularly helpful here, but you should make sure that you make the regular expression search as stringent as possible.



The number of optional SAM/BAM fields, their value types and the information stored within them completely depend on the alignment program and can thus vary substantially. Before you do any filtering on any flag, make sure you know how the aligner generated that value.

Here is an example for **retrieving reads with only one alignment** (aka uniquely aligned reads), which might be useful if STAR was not run with `--outFilterMultimapNmax 1`:

```

1 # STAR uses the NH:i tag to record the number of alignments found for a read
2 # NH:1 => 1 alignment; NH:2 => 2 alignments etc.
3 $ samtools view -h Aligned.sortedByCoord.out.bam | \ # decompress the BAM file
4     egrep "^\@|\bNH:i:1\b" | \ # lines with either @ at the beginning of the
      line (= header) or exact matches of NH:i:1 are returned
5     samtools view -S -b - > uniquely_aligned_reads.bam # turn the SAM file lines
      from stdin into a BAM file, - indicates standard input for samtools

```

To filter out **reads with insert sizes greater than 1000 bp**, one could make use of the CIGAR string. The following example assume that the alignment program indicated large insertions with the N operator (see Section 3.3.2) – this may not be true for all aligners!

```

1 # for the sake of simplicity, let's work on the SAM file:
2 $ samtools view -h WT_1_Aligned.sortedByCoord.out.bam > WT_1_Aligned.
    sortedByCoord.out.sam
3
4 # here's an example using grep, excluding lines with at least four digits
    followed by N
5 $ egrep -v "[0-9][0-9][0-9][0-9]N" WT_1_Aligned.sortedByCoord.out.sam >
    smallInsert_reads.sam
6
7 # awk can be used to match a regex within a specified column
8 $ awk '!($6 ~ /[0-9][0-9][0-9][0-9]N/)' {print $0}' WT_1_Aligned.sortedByCoord.
    out.sam > smallInsert_reads.sam

```

To retrieve **intron-spanning reads**, the commands will be similar:

```
1 # egrep allows for nicer regex syntax than grep
2 $ egrep "^(@|[0-9]+M[0-9]+N[0-9]+M)" WT_1_Aligned.sortedByCoord.out.sam >
   intron-spanning_reads.sam
3
4 # the same result achieved with awk
5 $ awk '$1 ~ /^[^@/ || $6 ~ /[0-9]+M[0-9]+N[0-9]+M/ {print $0}' WT_1_Aligned.
   sortedByCoord.out.sam > intron-spanning_reads.sam
```



1. How can you extract all reads that were aligned to the reverse strand?
2. Does it make sense to filter the **BAM** files generated by **STAR** using the mapping quality filter as shown above, i.e., do you find any differences after filtering with **-q 40**?

3.4 Quality control of aligned reads

Once the reads have been aligned, the following properties should be assessed before downstream analyses are started:

- Could most reads be aligned?
- Are there any obvious biases of the read distributions?
- Are the replicate samples as similar to each other as expected?

3.4.1 Basic alignment assessments

There are numerous ways to do basic checks of the alignment success. An alignment of RNA-seq reads is usually considered to have succeeded if the mapping rate is $>70\%$.

The very first QC of aligned reads should be to generally check the aligner's output. The **STAR** and **samtools index** commands in Section 3.2 generate the following files:

- *Aligned.sortedByCoord.out.bam
- *Aligned.sortedByCoord.out.bam.bai
- *Log.final.out
- *Log.out
- *Log.progress.out
- *SJ.out.tab
- *Unmapped.out.mate1

Information about the individual output files are given in the **STAR** manual which you can find in the program's directory (e.g., **STAR-STAR_2.4.2a/doc/STARmanual.pdf**) or online (<https://github.com/alexdobin/STAR/tree/master/doc>).



- Do you know what the different **STAR** output files contain? Which one(s) will you need most for your downstream analyses?
- How can you decrease the size of the ***out.mate1** files? What format do they have?
- Which optional **SAM** fields does **STAR** add and what do they represent?

Most aligners will return a summary of the basic stats of the aligned reads, such as the number of mapped reads. For **STAR**, the information is stored in ***Log.final.out**.

```

1 $ cat WT_1_Log.final.out
2                         Started job on | Jul 24 17:53:18
3                         Started mapping on | Jul 24 17:53:22
4                         Finished on | Jul 24 17:53:51
5      Mapping speed, Million of reads per hour | 870.78
6
7                         Number of input reads | 7014609
8                         Average input read length | 51
9                         UNIQUE READS:
10                        Uniquely mapped reads number | 6012470
11                        Uniquely mapped reads % | 85.71%
12                        Average mapped length | 50.73
13                        Number of splices: Total | 50315
14                        Number of splices: Annotated (sjdb) | 47843
15                        Number of splices: GT/AG | 49812
16                        Number of splices: GC/AG | 65
17                        Number of splices: AT/AC | 7
18                        Number of splices: Non-canonical | 431
19                        Mismatch rate per base, % | 0.36%
20                        Deletion rate per base | 0.00%
21                        Deletion average length | 1.37
22                        Insertion rate per base | 0.00%
23                        Insertion average length | 1.04
24                         MULTI-MAPPING READS:
25                        Number of reads mapped to multiple loci | 0
26                        % of reads mapped to multiple loci | 0.00%
27                        Number of reads mapped to too many loci | 796537
28                        % of reads mapped to too many loci | 11.36%
29                         UNMAPPED READS:
30                        % of reads unmapped: too many mismatches | 0.00%
31                        % of reads unmapped: too short | 2.90%
32                        % of reads unmapped: other | 0.04%

```

The number of *uniquely mapped reads* is usually the most important number. If you are handling more than two BAM files, it will certainly be worthwhile to visualize the alignment rate for all files, e.g., using R.

In addition to the log files generated by the mapping program, there are numerous ways to obtain information about the number and kinds of reads stored in a BAM file, e.g., using `samtools` or `RSeQC` (see below). The simplest approach to finding out the number of alignments within a BAM file is to do a line count.

```

1 samtools view Aligned.sortedByCoord.out.bam | wc -l

```

Note that if unmapped reads are present in the BAM file, these will also be counted, as well as multiple instances of the same read mapped to different locations if multi-mapped reads were kept. It is therefore more informative to run additional tools that will indicate the counts for specific FLAG values, too.

samtools flagstat This tool assesses the information from the FLAG field (see Section 3.3.2) and prints a summary report to the terminal.

```

1 $ samtools flagstat mat/additionalExamples/alignment/Luce/
2   Randall_1_multimappedAligned.sortedByCoord.out.bam
3 12064054 + 0 in total (QC-passed reads + QC-failed reads)
4 6390243 + 0 secondary
5 0 + 0 supplementary
6 0 + 0 duplicates
7 12064054 + 0 mapped (100.00%:-nan%)
8 0 + 0 paired in sequencing
9 0 + 0 read1
10 0 + 0 read2
10 0 + 0 properly paired (-nan%:-nan%)

```

```

11 0 + 0 with itself and mate mapped
12 0 + 0 singlettons (-nan%:-nan%)
13 0 + 0 with mate mapped to a different chr
14 0 + 0 with mate mapped to a different chr (mapQ>=5)

```

RSeQC's `bam_stats.py` RSeQC is a Python- and R-based suite of tools for various quality controls and visualizations, some of which are specific for RNA-seq experiments (Wang et al., 2012). See Table 9 for the list of all currently available scripts. Although RSeQC is one of the most popular tools for RNA-seq quality control, a recent publication revealed several bugs in the code of RSeQC (Hartley and Mullikin, 2015).

For basic alignment stats, one can use the `bam_stat.py` script:

```

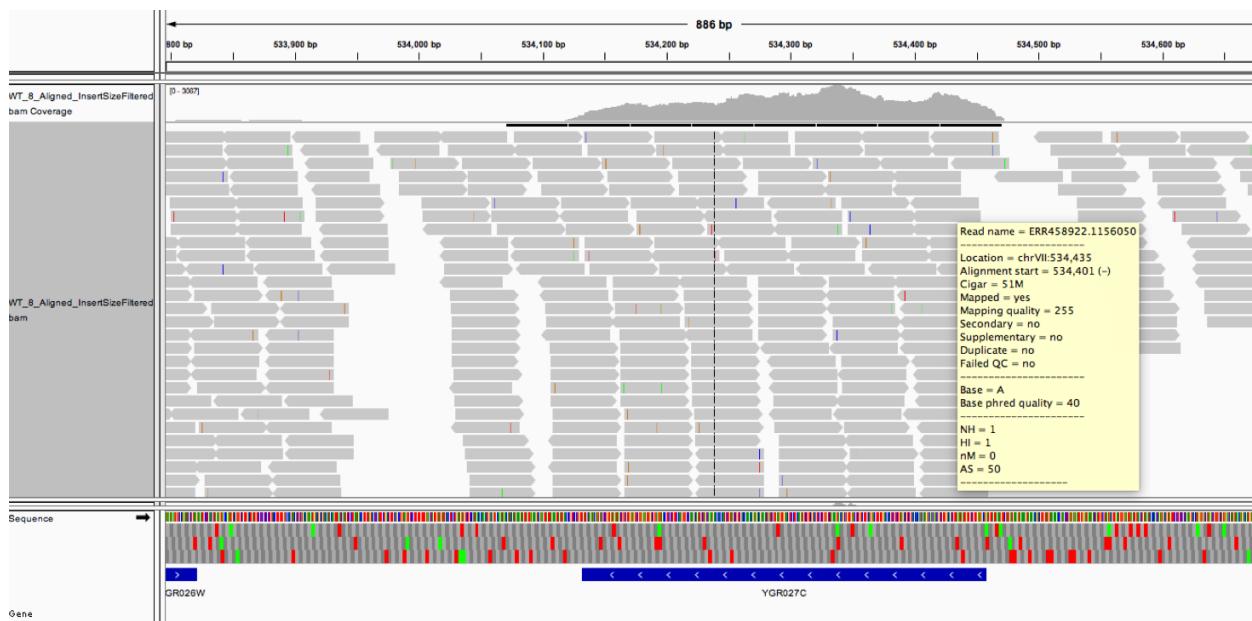
1 # RSeQC is based on Python; add the anaconda installation of Python to your
2 # PATH
3 $ export PATH=/zenodotus/abc/store/courses/2016_rnaseq/software/anaconda2/
4
5 # now, all RSeQC scripts are immediately accessible and you can, for example,
6 # run bam_stat.py
7 $ bam_stat.py -i WT_1_Aligned.sortedByCoord.out.bam
8
9 =====
#All numbers are READ count
#=====
10
11 Total records: 6012470
12
13 QC failed: 0
14 Optical/PCR duplicate: 0
15 Non primary hits 0
16 Unmapped reads: 0
17 mapq < mapq_cut (non-unique): 0
18
19 mapq >= mapq_cut (unique): 6012470
20 Read-1: 0
21 Read-2: 0
22 Reads map to '+': 3014735
23 Reads map to '-': 2997735
24 Non-splice reads: 5962195
25 Splice reads: 50275
26 Reads mapped in proper pairs: 0
27 Proper-paired reads map to different chrom:0

```

Visualization of aligned reads It is always a good idea to visually check the results, i.e., ensure the reads align to the expected regions, preferably without too many mismatches. Here, Genome Browsers come in handy. Different research groups have released different Genome Browsers, and the most well-known browsers are probably those from Ensembl and UCSC. There are some reasons why one may not want to use these web-based options (e.g., HIPAA-protected data or lack of bandwidth to upload all data), and rather resort to stand-alone Genome Browsers (see https://en.wikipedia.org/wiki/Genome_browser for an overview).

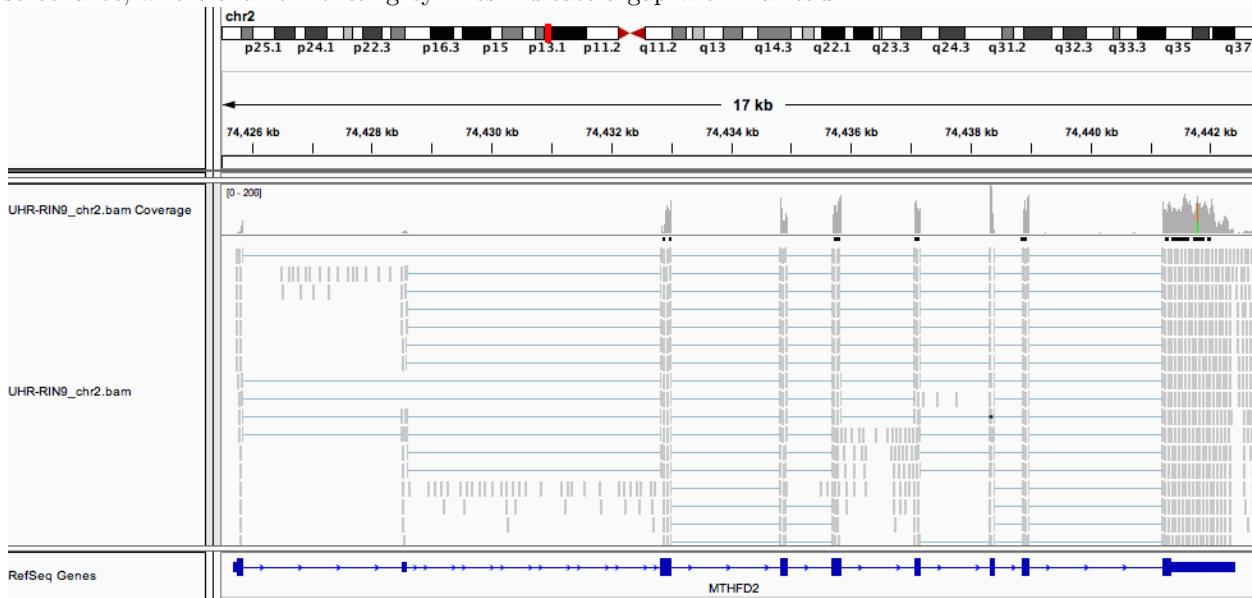
We are going to use the Broad Institute's Integrative Genomics Viewer (IGV) that can be downloaded after a quick registration with an academic email address from <https://www.broadinstitute.org/software/igv/download>. It requires an up-to-date Java installation.

The IGV Genome Browser can display numerous file formats, e.g., indexed BAM files with aligned reads and BED files with information about genomic loci (such as genes). The following IGV snapshot (in IGV, go to "File", then "Save image") shows the region surrounding an arbitrarily chosen yeast gene (blue box) and the reads aligned to it (grey arrows).

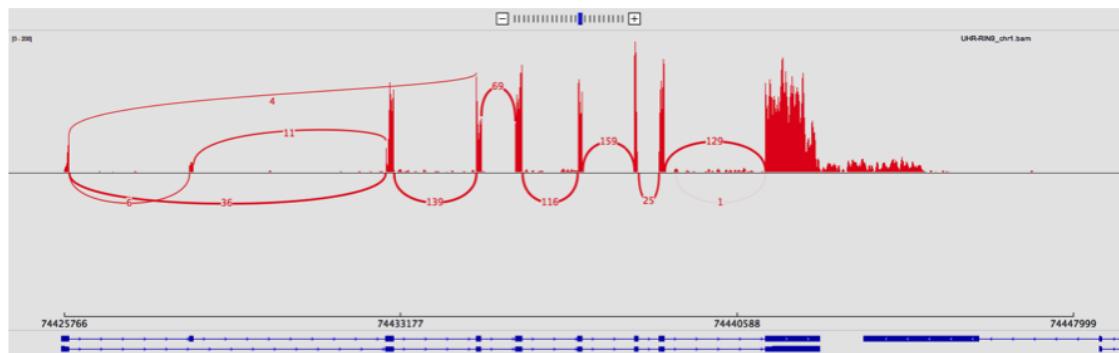


On top of the read alignment display, IGV also produces a coverage line that allows for quick identification of highly covered regions. Blue lines within the reads indicate insertions with respect to the reference genome, red lines indicate deletions. Since yeast genes are often intron-less, the reads can be aligned without gaps.

Human genes, however, tend to have multiple introns, which means that exon-exon-spanning reads must be aligned with often lengthy gaps within them (Figure 7). Examples of this can be seen in the following IGV screenshot, where the horizontal grey lines indicate a gap within a read:



If one is interested in the splice junctions of a particular gene, IGV can generate Sashimi plots (Katz et al., 2015): right-click on the track that contains the BAM file of interest and select “Sashimi plot”. The result will look like this:



In the Sashimi plot, bar graphs indicate read coverage, arcs indicate splice junctions, and numbers represent the number of reads that contain the respective splice junction.

3.4.2 Bias identification

Typical biases of RNA-seq experiments include:

- **Intron coverage:** if many reads align to introns, this is indicative of incomplete poly(A) enrichment or abundant presence of immature transcripts.
- **Intergenic reads:** if a significant portion of reads is aligned outside of annotated gene sequences, this may suggest genomic DNA contamination (or abundant non-coding transcripts).
- **3' bias:** over-representation of 3' portions of transcripts indicates RNA degradation.

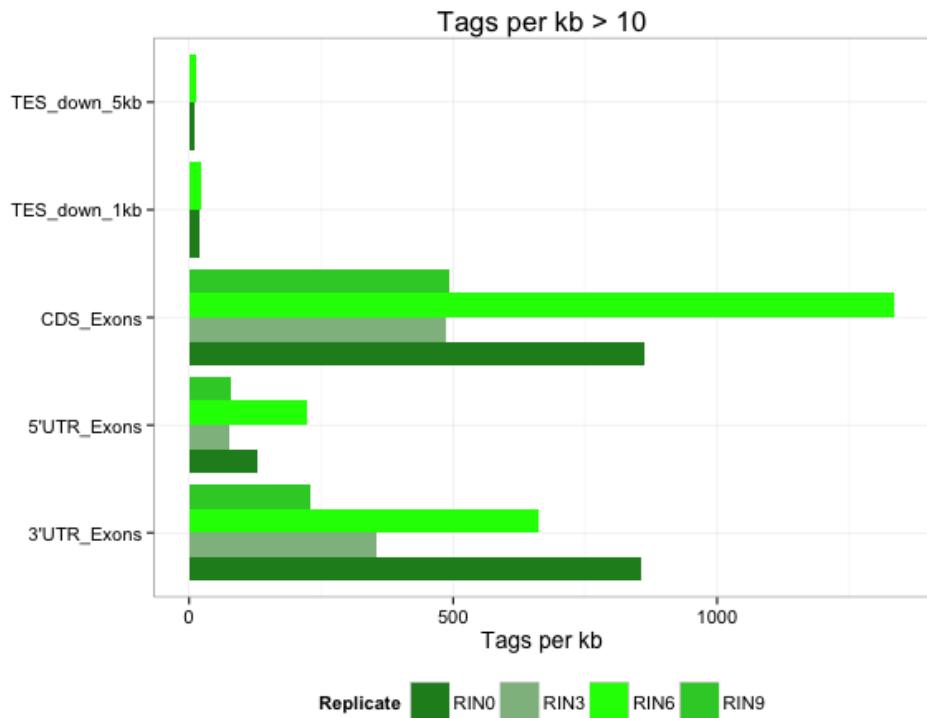
Read distribution For mRNA-seq, one would expect the majority of the aligned reads to overlap with exons. This assumption can be tested using the `read_distribution.py` script, which counts the numbers of reads overlapping with various gene- and transcript-associated genomic regions, such as exons and introns.

```

1 $ read_distribution.py -r ${REF_DIR}/sacCer3.bed \      # annotation file
2   -i WT_1_Aligned.sortedByCoord.out.bam \ # runs only on single files
3
4 Total Reads          7501551
5 Total Tags           7565292
6 Total Assigned Tags  6977808
7 =====
8 Group                Total_bases    Tag_count    Tags/Kb
9 CDS_Exons            8832031       6970376     789.22
10 5' UTR_Exons        0             0           0.00
11 3' UTR_Exons        0             0           0.00
12 Introns              69259         6353        91.73
13 TSS_up_1kb           2421198       309         0.13
14 TSS_up_5kb           3225862       309         0.10
15 TSS_up_10kb          3377251       309         0.09
16 TES_down_1kb         2073978       674         0.32
17 TES_down_5kb         3185496       770         0.24
18 TES_down_10kb        3386705       770         0.23
19 =====

```

To compare the read distribution values for different samples, it is helpful to turn the text-based output of `read_distribution.py` into a bar graph:



The .R script and `read_distribution.py` result files for this plot can be found at https://github.com/friedue/course_RNA-seq2015.

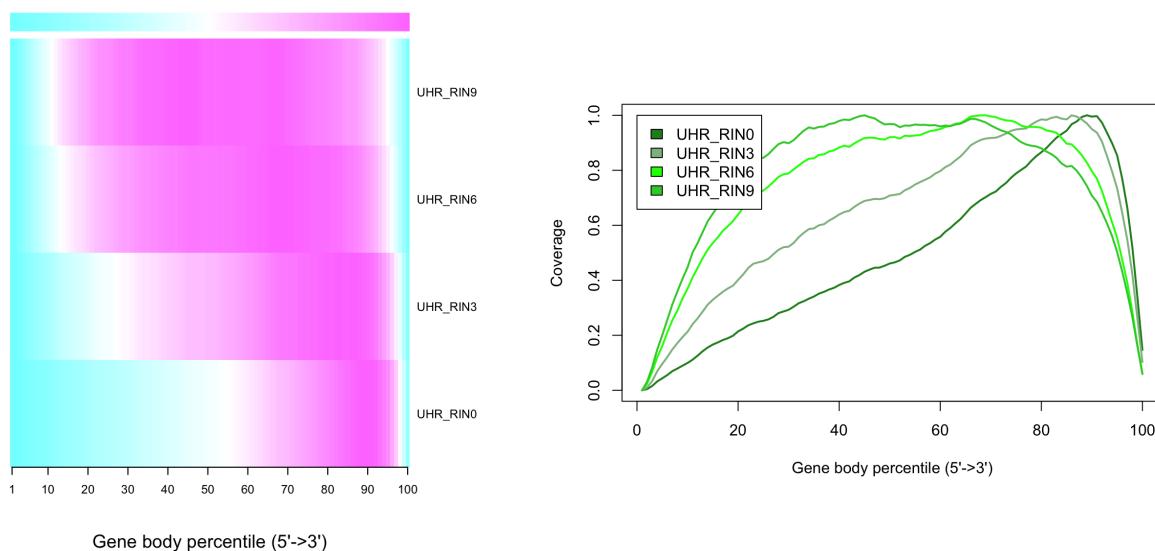
Gene body coverage To assess possible 3' or 5' biases, you can use RSeQC's `geneBody_coverage.py` script. Given an annotation file with the transcript models of your choice, it will divide each transcript into 100 sections, count the reads overlapping with each section and generate two plots visualizing the general abundance of reads across all transcript bodies.

```

1 $ REF_DIR=~/mat/referenceGenomes/S_cerevisiae
2
3 # Generate an index for the BAM file
4 $ samtools index WT_1_Aligned.sortedByCoord.out.bam
5
6 $ geneBody_coverage.py \
7 -i WT_1_Aligned.sortedByCoord.out.bam \ # aligned reads
8 -r ${REF_DIR}/sacCer3.bed \ # annotation file
9 -o geneBodyCoverage_WT_1 # output name
10
11 # if no plots are being generated automatically, the R script produced by the
12 # python script can be run manually:
13 $ <PATH to R installation>/bin/R < geneBodyCoverage_WT_1.geneBodyCoverage.r \
14 --vanilla \ # tells R not to waste time trying to load previous sessions etc.
--slave # makes R run in a less verbose mode

```

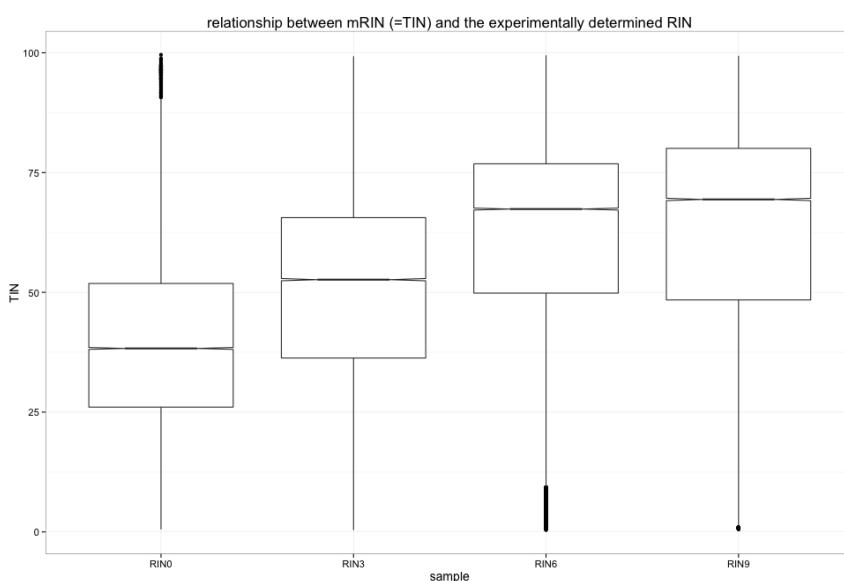
We ran the `geneBody_coverage.py` script on four human samples of RNA-seq with varying degrees of RNA quality, ranging from RIN = 0 (degraded) to RIN = 9 (high quality RNA) (see Section 1.1.1 for details about RIN). The resulting plots show varying degrees of 3' bias where samples with degraded RNA (RIN 0) show a more prominent bias than high-quality RNA (RIN 9).



in silico mRIN calculation The RNA integrity number (RIN, see Section 1.1.1) that is calculated during library preparation to assess the RNA quality is rarely indicated in the public data repositories. It might thus be informative to determine a measure of mRNA degradation *in silico*. RSeQC's `tin.py` script does exactly that, using the deviation from an expected uniform read distribution across the gene body as a proxy (Feng et al., 2015).

```
1 tin.py -i WT_1_Aligned.sortedByCoord.out.bam -r ${REF_DIR}/sacCer3.bed
```

`tin.py` will generate a `.xls` file where the *in silico* mRIN is stored for each gene or transcript from the `BED` file. The second output file, `*summary.txt`, gives a quick overview of the mean and median values across all genes for a given sample. Using human samples with known, experimentally determined RIN numbers, we can see that the *in silico* mRIN does correlate:



You can find the `.R` script and `tin.py` result files underlying these box plots in the following github repository: https://github.com/friedue/course_RNA-seq2015.

3.4.3 Quality control with QoRTs

As an alternative to RSeQC, the Quality of RNA-Seq Toolset (QoRTs) was developed, which is a comprehensive and multifunctional toolset that assists in quality control and data processing of high-throughput RNA sequencing data. It creates many of the same output and plots as RSeQC, but the authors claim it is more accurate (Hartley and Mullikin, 2015).

The following command runs the complete QoRTs QC analysis (refer to Table 10 to see all the individual functions and commands).

```

1 $ REF_DIR=~/mat/referenceGenomes/S_cerevisiae
2
3 # to obtain the total number of raw reads you can make use of the calculator
4 # capabilities of bc
5 $ for FASTQ in mat/rawReads_yeast_Gierlinski/WT_1/ERR45849.gz; do zcat $FASTQ |
6   wc -l ; done | paste -sd+ | bc | awk '{print $1/4}'
7
8 $ java -Xmx4g -jar mat/software/qorts.jar QC \
9   --singleEnded \ # QoRTs assumes the data is paired-end unless this flag is
10  specified
11  --seqReadCt 7014609 \ # total number of starting reads before mapping (see
12  cmd above)
13  --generatePlots WT_1_Aligned.sortedByCoord.out.bam \ # aligned reads
14  ${REF_DIR}/sacCer3.gtf \ # annotation file
15  ./QoRTs_output/ # output folder

```

Note that by default, QoRTs assumes the data is paired end unless otherwise specified.

The run or exclude individual functions:

```

1 $ REF_DIR=~/mat/referenceGenomes/S_cerevisiae
2
3 # to only run a function
4 $ java -Xmx4g -jar qorts.jar QC \
5   --singleEnded \
6   --runFunctions writeGeneBody \ # run only the genebody coverage function
7   --generatePlots WT_1_Aligned.sortedByCoord.out.bam \
8   ${REF_DIR}/sacCer3.gtf \
9   ./QoRTs_output/
10
11 # to exclude a function
12 $ java -Xmx4g -jar QoRTs.jar QC \
13   --singleEnded \
14   --skipFunctions JunctionCalcs \ # run every function except the
15   JunctionCalcs function
16   --generatePlots WT_1_Aligned.sortedByCoord.out.bam \
17   ${REF_DIR}/sacCer3.gtf \
18   ./QoRTs_output/

```

To include or exclude more than one function, use a comma-delimited list (without white spaces) of functions to include or skip.

An example QoRTs report can be found at <http://chagall.med.cornell.edu/RNASEQcourse/>.

4 Read Quantification

4.1 Gene-based read counting

To compare the expression of single genes between different conditions, an essential step is the quantification of reads per gene. In principle, the counting of reads overlapping with genomic features is a fairly simple task, but there are some details that need to be decided on depending on the nature of your experiment and the desired outcome (Figure 10).

When counting reads, make sure you know how the program handles the following:

- overlap size (full read vs. partial overlap)
- multi-mapping reads
- reads overlapping multiple genomic features of the same kind
- reads overlapping introns

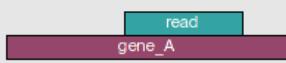
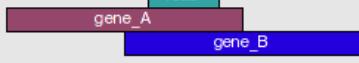
	union	intersection _strict	intersection _nonempty
	gene_A	gene_A	gene_A
	gene_A	no_feature	gene_A
	gene_A	no_feature	gene_A
	gene_A	gene_A	gene_A
	gene_A	gene_A	gene_A
	ambiguous	gene_A	gene_A
	ambiguous	ambiguous	ambiguous

Figure 10: The `htseq-count` script of the HTSeq suite offers three different modes to handle details of read–feature overlaps that are depicted here. The default of `featureCounts` is the behavior of the `union` option. Image taken from <http://www-huber.embl.de/users/anders/HTSeq/doc/count.html>.

The most popular tools for gene quantification are `htseq-count` and `featureCounts`. Both are part of larger tool packages (Anders et al., 2014; Liao et al., 2014). `htseq-count` offers three different modes to tune its behavior to define overlap instances (Figure 10). The recommended mode is `union`, which counts overlaps even if a read only shares parts of its sequence with a genomic feature and disregards reads that overlap more than one feature. This is similar to `featureCounts` that calls a hit if any overlap (1 bp or more) is found between the read and a feature and provides the option to either exclude multi-overlap reads or to count them for each feature that is overlapped.

In addition to the nature and lengths of the reads, gene expression quantification will be strongly affected by the underlying gene models that are usually supplied to the quantification programs via GTF or BED(-like) files (see Section 3.1 for details on the file formats and annotations).

The following commands will count the number of reads overlapping with genes using `featureCounts`.

```
1 # count reads per gene
2 $ subread-1.5.0-p3-Linux-x86_64/bin/featureCounts \
3 -a ${REF_DIR}/sacCer3.gtf \
4 -o featureCounts_results.txt \
5 alignment/*bam # use all BAM files in the folder "alignment"
```

`featureCounts` also allows to count reads overlapping with individual exons.

```
1 # count reads per exon
2 $ subread-1.5.0-p3-Linux-x86_64/bin/featureCounts \
3 -a ${REF_DIR}/sacCer3.gtf \
4 -f \ # count read overlaps on the feature level
5 -t exon \ # feature type
6 -O \ # allow reads to overlap more than one exon
7 -o featCounts_exons.txt \
8 alignment/*bam
```

However, there are (at least) two caveats here:

- If an exon is part of more than one isoform in the annotation file, `featureCounts` will return the read counts for the same exon multiple times ($n = \text{number of transcripts with that exon}$). Make sure you remove those multiple entries in the result file before the differential expression analysis, e.g., using a UNIX command[†] or within R.
- If you want to assess differential expression of exons, it is highly recommended to create an annotation file where overlapping exons of different isoforms are split into artificially disjoint bins before applying `featureCounts`. See, for example, Anders et al. (2012). To create such a “flattened” annotation file from a GTF file (Section 3.1.1), you can use the `dexseq_prepare_annotation.py` script of the DEXSeq package (Anders et al., 2012) and the section “Preparing the annotation” of the corresponding vignette at bioconductor.

4.2 Isoform counting methods

The previously discussed methods count the number of fragments that can be assigned to a gene as a whole. There is another school of thought that insists that quantifying reads that originated from transcripts should also be done on the transcript level[‡]. So far, most comparisons of methods point towards superior results of gene-based quantification (which is why we adhere to it for now) and there is no standard technique for summarizing expression levels of genes with several isoforms (see, for example, Soneson et al. (2015), Dapaz et al. (2016), Germain et al. (2016), and (Teng et al., 2016) for detailed comparisons of transcript-level quantifications).

One trend seems to be clear though: the simple count-based approaches tend to underperform when they are used to determine transcript-level counts because they generally tend to ignore reads that overlap with more than one feature. While this is reasonable when the features are entire genes, this leads to an enormous number of discarded reads since multiple transcripts of the same gene naturally tend to overlap.

In order to quantify isoforms, you should perhaps look into different programs, e.g., Cufflinks (Trapnell et al., 2012), RSEM (Li and Dewey, 2011), eXpress (Roberts and Pachter, 2013) – these tools have been around the longest and are therefore most often cited. These tools typically use a *deBruijn* graph approach to assign reads to a given isoform if they are compatible with that transcript structure (see Figure 11 for a simple example).

[†]e.g., `sort -k2,2n -k3,3n featureCounts_exons.txt | uniq`

[‡]For arguments from followers of the transcript-focused school of thought, see, e.g., Trapnell et al. (2013) and Pimentel’s talk.

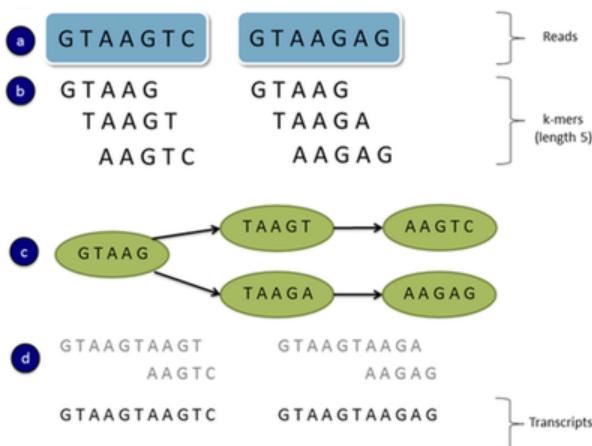


Figure 11: Schema of a simple *deBruijn* graph-based transcript assembly. (A) Read sequences are split into (B) all subsequence k -mers (here: of length 5) from the reads. (C) A *deBruijn* graph is constructed using unique k -mers as the nodes and overlapping k -mers connected by edges (a k -mer shifted by one base overlaps another k -mer by $k-1$ bases). (D) The transcripts are assembled by traversing the two paths in the graph. Figure taken from Moreton et al. (2015)

Very recently, two groups published algorithms that are based on the idea that it may not be important to exactly know *where* within a transcript a certain read originated from. Instead, it may be enough to simply know *which* transcript the read represents. These algorithms therefore do not generate a BAM file because they do not worry about finding the best possible alignment. Instead, they yield a (probabilistic) measure of how many reads indicate the presence of each transcript. The tools are:

- Sailfish (Patro et al., 2014), or the more updated, more accurate version, Salmon (Patro et al., 2015)
- kallisto (Bray et al., 2016)

While these approaches are extremely fast compared to the usual alignment–counting routines that we have described at length, they cannot be used to detect novel isoforms.

Instead of direct isoform quantification, you may be able to glean more accurate answers from alternative approaches, e.g., quantification of exons (Anders et al., 2012)[§] or estimates of alternative splicing events such as exon skipping, intron retention etc. (e.g., MISO (Katz et al., 2010), rMATS (Shen et al., 2014)).

The main take home message here is once again: Know your data and your question, and research the individual strengths and pitfalls of the individual tools before deciding which one to use. For example, one major issue reported for Cufflinks is its inability to handle single-exon transcripts. Therefore, you should avoid using it if you are dealing with a fairly simple transcriptome (Kanitz et al., 2015). On the other hand, transcriptome reconstruction as attempted by Cufflinks generates large amounts of false positives (as well as false negatives) in very complicated transcriptomes, such as the human one (Jänes et al., 2015). In comparison, the novel lightweight quantification algorithms perform well for isoform quantification of known transcriptomes, but they are naturally very sensitive to incomplete or changing annotation. In addition, it is not entirely clear yet whether the resulting values can be used with the established algorithms to determine differential gene expression (Soneson et al., 2015; Pimentel et al., 2016).

The main caveats of assigning reads to transcripts are:

- inconsistent annotation of transcripts
- multiple isoforms of widely differing lengths
- anti-sense/overlapping transcripts of different genes

There is no really good solution yet! Be careful with your conclusions and if possible, limit your analyses to gene-based approaches.

[§]The above shown `featureCounts`-based exon counting should not be used with `DEXSeq` unless exons with varying boundaries have been divided into disjoint bins (Anders et al., 2012; Teng et al., 2016; Soneson et al., 2016).

5 Normalizing and Transforming Read Counts

Given a uniform sampling of a diverse transcript pool, the number of sequenced reads mapped to a gene depends on:

- its own expression level,
- its length,
- the sequencing depth,
- the expression of all other genes within the sample.

In order to compare the gene expression *between two conditions*, we must therefore calculate the fraction of the reads assigned to each gene relative to the total number of reads and with respect to the entire RNA repertoire which may vary drastically from sample to sample. While the number of sequenced reads is known, the total RNA library and its complexity is unknown and variation between samples may be due to contamination as well as biological reasons. The purpose of normalization is to eliminate systematic effects that are not associated with the biological differences of interest. See Table 11 for details on the most commonly used normalization methods that deal with these issues in different ways.

5.1 Normalization for sequencing depth differences

As shown in Figure 12, the size factor method implemented by the R package `DESeq` leads to relatively similar read count distribution between different libraries. We will now use the output of `featureCounts` (= raw read counts), read them into R and normalize the read counts for sequencing depth differences with `DESeq`.

```
1  #### Open an R console, e.g. using RStudio
2  # code lines starting with `>` indicate the R console
3
4  # get the table of read counts
5  > read.counts <- read.table("featureCounts_result.txt", header = TRUE)
6
7  # the gene IDs should be stored as row.names
8  > row.names(read.counts) <- read.counts$Geneid
9
10 # exclude all columns that do not contain read counts
11 > read.counts <- read.counts[, -c(1:6)]
12
13 # give meaningful sample names - this can be achieved via numerous approaches
14 # the one shown here is the least generic and most error-prone one!
15 > names(read.counts) <- c("SNF2_1", "SNF2_2", "SNF2_3", "SNF2_4", "SNF2_5", "WT_1",
   "WT_2", "WT_3", "WT_4", "WT_5")
16
17 # alternative way to assign the sample names, which reduces the
18 # potential for typos as well as for the wrong order:
19 > gsub(".*(WT|SNF2)(_[0-9]+).*", "\\\1\\\2", names(read.counts))
20
21 # ALWAYS CHECK YOUR DATA AFTER YOU MANIPULATE IT!
22 > str(read.counts)
23 > head(read.counts)
24      SNF2_1 SNF2_2 SNF2_3 SNF2_4 SNF2_5 WT_1 WT_2 WT_3 WT_4 WT_5
25 YAL012W    7347    7170   7643   8111   5943  4309  3769  3034  5601  4164
26 YAL069W     0       0       0       0       0       0       0       0       0       0
27 YAL068W-A   0       0       0       0       0       0       0       0       0       0
28 YAL068C     2       2       2       1       0       0       0       0       2       2
29 YAL067W-A   0       0       0       0       0       0       0       0       0       0
30 YAL067C    103     51      44      90      53     12      23      21      30      29
```

Now that we have the read counts, we also need some information about the samples. This should be a `data.frame`, where the rows match the column names of the count data we just generated. In addition, each row should contain information about the condition of each sample (here: WT and SNF2 [knock-out]) [¶].

```

1 # make a data.frame with meta-data where row.names should match the individual
2 # sample names
3 > sample.info <- data.frame(condition = c( rep("SNF2", 5), rep("WT", 5)),
4 #                               row.names = names(read.counts) )
5 > sample.info
6   condition
7 SNF2_1      SNF2
8 SNF2_2      SNF2
9 SNF2_3      SNF2
10 SNF2_4     SNF2
11 SNF2_5     SNF2
12 WT_1       WT
13 WT_2       WT
14 WT_3       WT
15 WT_4       WT
16 WT_5       WT
17
18 # install DESeq2 which is not available via install.packages(), but through
#     bioconductor
19 > source("http://bioconductor.org/biocLite.R")
20 > biocLite("DESeq2")
21 > library(DESeq2)
22
23 # generate the DESeqDataSet
24 > DESeq.ds <- DESeqDataSetFromMatrix(countData = read.counts_noZero,
25 #                                         colData = sample.info,
26 #                                         design = ~ condition)
27
28 # remove genes without any counts
29 > DESeq.ds <- DESeq.ds[ rowSums(counts(DESeq.ds)) > 0, ]
30
31 # investigate different library sizes
32 > colSums(counts(DESeq.ds)) # should be the same as colSums(read.counts)
33
34 # calculate the size factors; normalized read counts can be retrieved via
#     counts(..., normalized = TRUE)
35 > DESeq.ds <- estimateSizeFactors(DESeq.ds)
36 > counts.sf_normalized <- counts(DESeq.ds, normalized = TRUE)
```



While the majority of normalization methods work well, **RPKM** and **total count** normalization should be **avoided** in the context of DE analysis, no matter how often you see them applied in published studies. RPKM, FPKM etc. are only needed if expression values need to be compared *between different genes* within the *same sample* for which the different gene lengths must be taken into consideration.



1. Name two technical reasons why the read count for the same gene may vary *between two samples* although it is not differentially expressed.
2. Name two technical reasons why the read counts of two genes may vary *within the same sample* although they are expressed at the same level.

[¶]These are requirements stipulated by `DESeq2`.

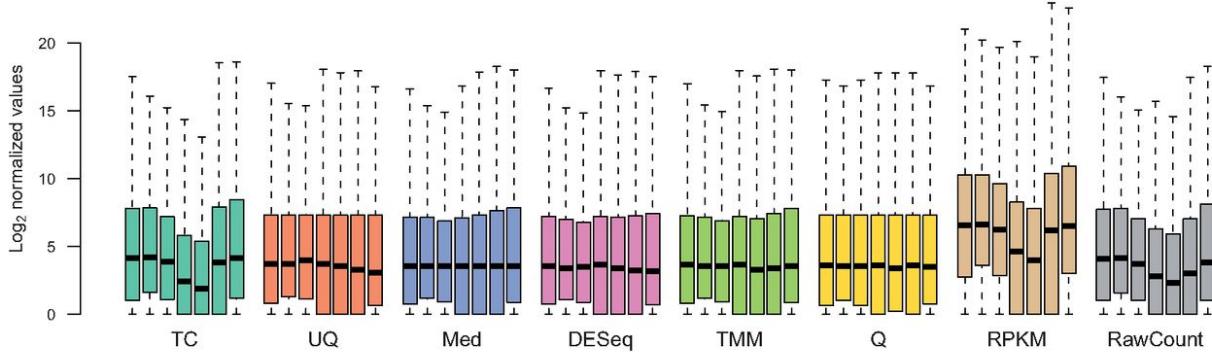


Figure 12: Figure from Dillies et al. (2013) that shows the effects of different approaches to normalize for read count differences due to library sizes (TC, total count; UQ, upper quartile; Med, median; DESeq, size factor; TMM, Trimmed Mean of M-values; Q, quantile) or gene lengths (RPKM). See Tables 11 and 12 for details of the different normalization methods.

5.2 Transformation of sequencing-depth-normalized read counts

While most models used for differential gene expression testing operate on the raw count values, many downstream analyses (including clustering) work better if the read counts are *transformed* to the *log* scale. While you will occasionally see \log_{10} transformed read counts, \log_2 is more commonly used because it is easier to think about doubled values rather than powers of 10. The transformation should be done in addition to sequencing depth normalization.

5.2.1 \log_2 transformation of read counts

```
1 # transform size-factor normalized read counts to log2 scale using a
  pseudocount of 1
2 > log.norm.counts <- log2(counts.sf_normalized + 1)
```

You can see how the \log_2 transformation makes even simple graphs more easily interpretable by generating boxplots of read counts similar to the ones in Figure 12:

```
1 > par(mfrow=c(2,1)) # to plot the following two images underneath each other
2
3 # first, boxplots of non-transformed read counts (one per sample)
4 > boxplot(counts.sf_normalized, notch = TRUE,
5           main = "untransformed read counts", ylab = "read counts")
6
7 # box plots of log2-transformed read counts
8 > boxplot(log.norm.counts, notch = TRUE,
9           main = "log2-transformed read counts",
10          ylab = "log2(read counts)")
```

To get an impression of how similar read counts are between replicates, it is often insightful to simply plot the counts in a pairwise manner (Figure 14, upper panels). This can be achieved with the basic, but versatile `plot()` function:

```
1 plot(log.norm.counts[,1:2], cex=.1, main = "Normalized log2(read counts)")
```

Many statistical tests and analyses assume that data is homoskedastic, i.e. that all variables have similar variance. However, data with large differences among the sizes of the individual observations often shows heteroskedastic behavior. One way to visually check for heteroskedasticity is to plot the mean vs. the standard deviation (Figure 14, lower panel).

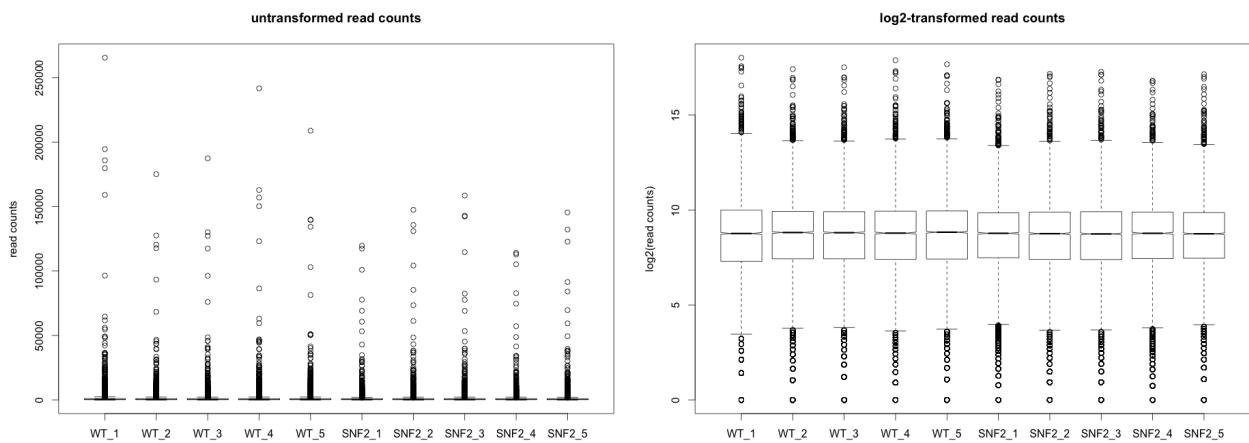


Figure 13: Comparison of the read distribution plots for untransformed and \log_2 -transformed values.

```

1 # install the vsn package
2 > source("http://bioconductor.org/biocLite.R")
3 > biocLite("vsn")
4 > library(vsn)
5 > meanSdPlot(log.norm.counts, ranks = FALSE, ylim = c(0,3),
   main="sequencing depth normalized log2(read counts)")
6

```

Note that the y-axis shows the variance of the read counts across all samples, so some variability is, in fact, expected. The clear hump on the left-hand side indicates that for read counts < 32 ($2^5 = 32$), the variance is higher than for those with greater read counts, so some dependence of the variation on the mean can be observed.

5.2.2 Transformation of read counts including variance shrinkage

To reduce the amount of heteroskedasticity, DESeq2 and also edgeR offer several means to shrink the variance of low read counts by using the dispersion-mean trend seen for the entire data set as a reference. That means that genes with low, but highly variable read counts across the different samples and replicates will be assigned more homogeneous read counts so that their variance resembles the variance of genes with higher read counts. DESeq2's `rlog()` function returns values that are both normalized for sequencing depth and transformed to the \log_2 scale where the values are adjusted to fit the experiment-wide trend of the variance-mean relationship.

```

1 # obtain regularized log-transformed values
2 > rlog.DESeq.sumExp <- rlog(DESeq.ds, blind = TRUE)
3 > rlog.norm.counts <- assay(rlog.DESeq.sumExp)
4
5 # mean-sd plot for rlog-transformed data
6 > meanSdPlot(rlog.norm.counts, ranks = FALSE,
   ylim = c(0,3), main = "rlog-transformed read counts")
7

```

The `rlog()` function's `blind` parameter should be set to `FALSE` if the different conditions lead to strong differences in a large proportion of the genes. If `rlog()` is applied without incorporating the knowledge of the experimental design (`blind = TRUE`, the default setting), the dispersion will be greatly overestimated in such cases.

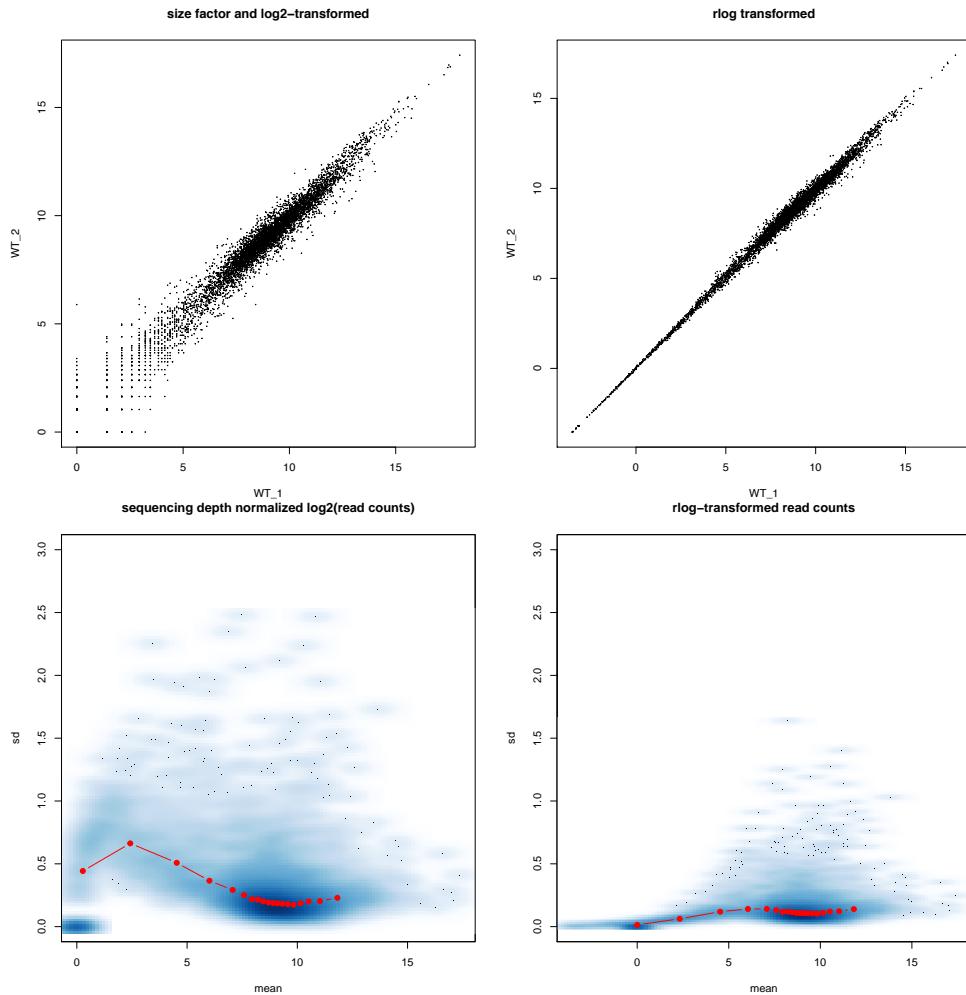


Figure 14: Comparison of \log_2 - and $r\log$ -transformed read counts. The upper panel shows simple pairwise comparisons of replicate samples; the lower panel contains mean-sd-plots based on all samples of the experiment.

5.3 Read count correlations

In order for differential gene expression analysis to succeed, the expression values of individual genes should be fairly similar for biological replicates. The ENCODE consortium recommends that “for messenger RNA, (...) biological replicates [should] display >0.9 correlation for transcripts/features”. The *Pearson correlation coefficient*, r , is a measure of the strength of the linear relationship between two variables and is often used to assess the similarity of RNA-seq samples in a pair-wise fashion. It is defined as the covariance of two variables divided by the product of their standard deviation.

Hierarchical clustering To determine whether the different sample types can be separated in an unsupervised fashion (i.e., samples of different conditions are more dissimilar to each other than replicates within the same condition), hierarchical clustering can be used.

Hierarchical clustering requires two decisions:

1. How should the (dis)similarity between pairs be calculated?
2. How should the (dis)similarity be used for the clustering?

In addition to Pearson correlation coefficient (where the distance is $d = 1 - r$), the Euclidean distance is often used as a measure of distance between two vectors of read counts. The latter is strongly influenced by differences of the scale: if two samples show large differences in sequencing depth, this will affect the Euclidean distance more than the distance based on the Pearson correlation coefficient.

As for the calculation of the similarity, multiple options exist to decide on how the values should be joined into clusters. The most popular choices for the *linkage function* are

- *complete*: intercluster distance \equiv largest distance between any 2 members of either cluster
- *average*: intercluster distance \equiv average distance between any 2 members
- *single*: intercluster distance \equiv shortest distance between any 2 members



Avoid “single” linkage on gene expression data; “complete” and “average” linkage tend to be much more appropriate, with “complete” linkage often outperforming “average” (Gibbons and Roth, 2002).

The result of hierarchical clustering is a *dendrogram* (Figure 15); clusters are obtained by cutting the dendrogram at a level where the jump between two consecutive nodes is large: connected components then form individual clusters.

A dendrogram can be generated in R using the functions `cor()`, `as.dist()`, and `hclust()`:

```

1 # cor() calculates the correlation between columns of a matrix
2 > distance.m_rlog <- as.dist(1 - cor(rlog.norm.counts, method = "pearson" ))
3
4 # plot() can directly interpret the output of hclust()
5 > plot(hclust(distance.m_rlog),
6   labels = colnames(rlog.norm.counts),
7   main = "rlog transformed read counts\ndistance: Pearson correlation")

```



1. Which linkage method was used for the dendrogram generated with the code shown above?
2. Can you make a dendrogram with Euclidean distance and linkage method “average”?

Principal Components Analysis (PCA) A complementary approach to determine whether samples display greater variability between experimental conditions than between replicates of the same treatment is principal components analysis. Principal components represent the directions along which the variation in the data is maximal, so that a few dimensions can be used to represent the information from thousands of mRNAs. Most commonly, the two principal components explaining the majority of the variability are displayed. It is also useful to identify unexpected patterns, such as batch effects or outliers, but PCA is not

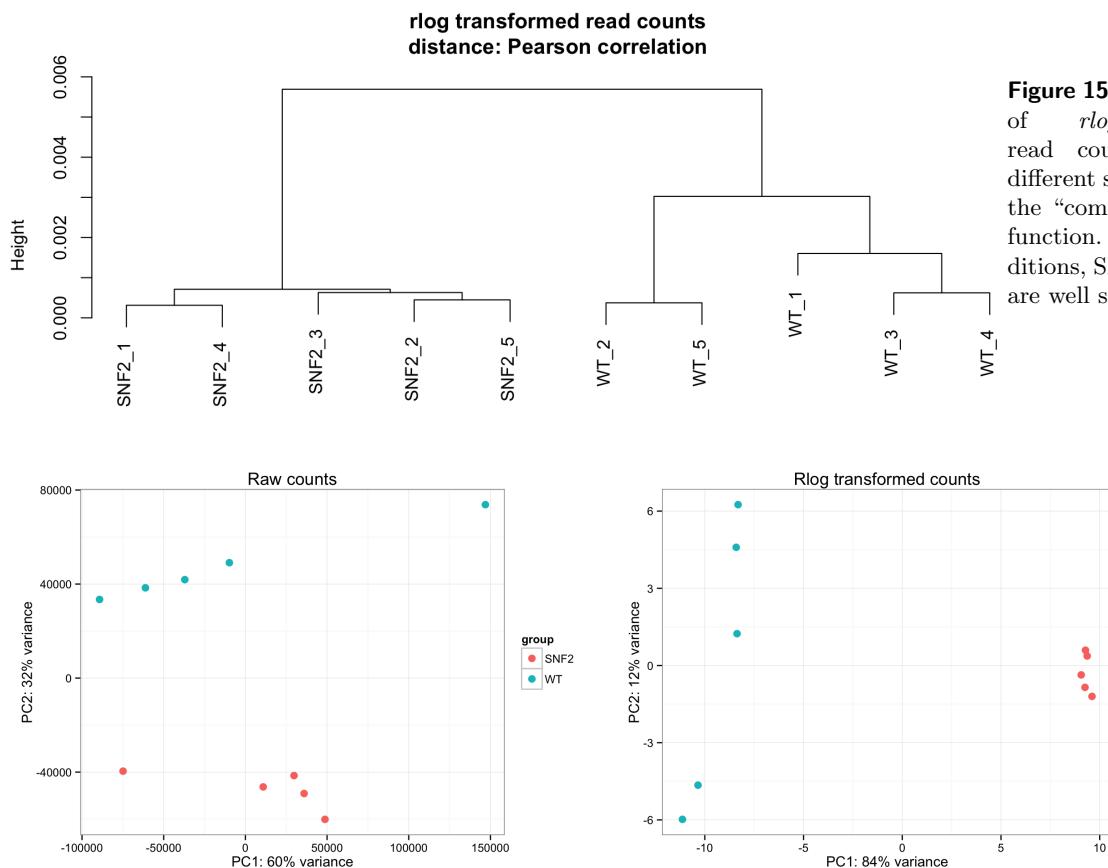


Figure 15: Dendrogram of *rlog*-transformed read counts for ten different samples, using the “complete” linkage function. The two conditions, SNF2 and WT, are well separated.

Figure 16: PCA on raw counts and *rlog*-transformed read counts with the DESeq2 convenience function `plotPCA()`. As indicated by the labels of the axes, the different sample types explain a greater fraction of the variance for *rlog*-transformed values than for the raw counts.

designed to discover unknown groupings; it is up to the researcher to identify the experimental or technical reason underlying the principal components.

PCA can be performed in base R using the function `prcomp()`.

```

1 > pc <- prcomp(t(rlog.norm.counts))
2 > plot(pc$x[,1], pc$x[,2],
3   col = colData(DESeq.ds)[,1],
4   main = "PCA of seq.depth normalized\n and rlog-transformed read counts")

```

DESeq2 also offers a convenience function based on `ggplot2` to do PCA directly on a `DESeqDataSet`:

```

1 > library(DESeq2)
2 > library(ggplot2)
3
4 # PCA
5 > P <- plotPCA(rlog.DESeq.sumExp)
6
7 # plot cosmetics
8 > P <- P + theme_bw() + ggtitle("Rlog transformed counts")
9 > print(P)

```

! PCA and clustering should be done on normalized and preferably transformed read counts, so that the high variability of low read counts does not occlude potentially informative data trends.

6 Differential Gene Expression Analysis

The two basic tasks of all DGE tools are:

1. Estimate the *magnitude* of differential expression between two or more conditions based on read counts from replicated samples, i.e., calculate the fold change of read counts, taking into account the differences in sequencing depth and variability.
2. Estimate the *significance* of the difference and correct for multiple testing.

To determine the genes whose read count differences between two conditions are greater than expected by chance, DGE tools must make assumptions about the distribution of read counts. The null hypothesis – that the mean read counts of the samples of condition A are equal to the mean read counts of the samples of condition B – is tested for each gene individually. One of the most popular choices to model the read counts is the *Poisson* distribution because:

- individual reads can be interpreted as binary data (Bernoulli trials): they either originate from gene i or not.
- we are trying to model the discrete probability distribution of the number of successes (success = read is present in the sequenced library).
- the pool of possible reads that could be present is large, while the proportion of reads belonging to gene i is quite small.

The nice feature of a Poisson distribution is that $\text{variance} = \text{mean}$. Thus, if the RNA-seq experiment gives us a precise estimate of the mean read counts per condition, we implicitly know what kind of variance to expect for read counts that are not truly changing between two conditions. This, in turn, then allows us to identify those genes that show greater differences between the two conditions than expected by chance.

While read counts of the same library preparation (= technical replicates) can indeed be well approximated by the Poisson distribution, it has been shown that biological replicates have greater variance (noise) than expected. This *overdispersion* can be captured with the *negative binomial* distribution, which is a more general form of the Poisson distribution that allows the variance to exceed the mean. The square root of the dispersion is the coefficient of variation – $\frac{SD}{\text{mean}}$ – after subtracting the variance we expect due to Poisson sampling.

In contrast to the Poisson distribution, we now need to estimate two parameters from the read counts: the mean as well as the dispersion. The precision of these estimates strongly depends on the number (and variation) of replicates – the more replicates, the better the grasp on the underlying mean expression values of unchanged genes and the variance that is due to biological variation rather than the experimental treatment. For most RNA-seq experiments, only two to three replicates are available, which is not enough for reliable mean and variance estimates. Some tools therefore compensate for the lack of replication by borrowing information across genes with similar expression values and shrink a given gene's variance towards the regressed values. These fitted values of the mean and dispersion are then used instead of the raw estimates to test for differential gene expression.

The best performing tools tend to be `edgeR` (Robinson et al., 2010), `DESeq/DESeq2` (Anders and Huber, 2010; Love et al., 2014), and `limma-voom` (Ritchie et al., 2015) (see Rapaport et al. (2013); Soneson and Delorenzi (2013); Schurch et al. (2015) for reviews of DGE tools). `DESeq` and `limma-voom` tend to be more conservative than `edgeR` (better control of false positives), but `edgeR` is recommended for experiments with fewer than 12 replicates (Schurch et al., 2015).



All statistical methods developed for read counts rely on approximations of various kinds, so that assumptions must be made about the data properties. `edgeR` and `DESeq`, for example, assume that the majority of the transcriptome is *unchanged* between the two conditions. If this assumption is not met by the data, both \log_2 fold change and the significance indicators are most likely incorrect!

Table 5: Comparison of programs for differential gene expression identification (Rapaport et al., 2013; Seyednasrollah et al., 2015; Schurch et al., 2015).

Feature	DESeq2	edgeR	limmaVoom	Cuffdiff
Seq. depth normalization	Sample-wise size factor	Gene-wise trimmed median of means (TMM)	Gene-wise trimmed median of means (TMM)	FPKM-like or DESeq-like
Assumed distribution	Neg. binomial	Neg. binomial	<i>log</i> -normal	Neg. binomial
Test for DE	Exact test (Wald)	Exact test for over-dispersed data	Generalized linear model	<i>t</i> -test
False positives	Low	Low	Low	High
Detection of differential isoforms	No	No	No	Yes
Support for multi-factored experiments	Yes	Yes	Yes	No
Runtime (3-5 replicates)	Seconds to minutes	Seconds to minutes	Seconds to minutes	Hours

6.1 Running DGE analysis tools

6.1.1 DESeq2 workflow

For our example data set, we would like to compare the effect of the *snf2* mutants versus the wildtype samples, with the wildtype values used as the denominator for the fold change calculation.

```
1 # DESeq uses the levels of the condition to determine the order of the
   comparison
2 > str(colData(DESeq.ds)$condition)
3
4 # set WT as the first-level-factor
5 > colData(DESeq.ds)$condition <- relevel(colData(DESeq.ds)$condition, "WT")
```

Now, running the DGE analysis is very simple:

```
1 > DESeq.ds <- DESeq(DESeq.ds)
```

The `DESeq()` function is basically a wrapper around the following three individual functions:

```
1 > DESeq.ds <- estimateSizeFactors(DESeq.ds) # sequencing depth normalization
   between the samples
2 > DESeq.ds <- estimateDispersions(DESeq.ds) # gene-wise dispersion estimates
   across all samples
3 > DESeq.ds <- nbinomWaldTest(DESeq.ds) # this fits a negative binomial GLM and
   applies Wald statistics to each gene
```



Note that the DGE analysis is performed on the *raw* read counts (untransformed, not normalized for sequencing depth), so supplying anything but raw read counts to either `DESeq` or `edgeR` will result in nonsensical results.

The `results()` function lets you extract the base means across samples, moderated \log_2 fold changes, standard errors, test statistics etc. for every gene.

```

1 > DGE.results <- results(DESeq.ds, pAdjustMethod = "BH")
2 > summary(DGE.results)
3
4 # the DESeqResult object can basically be handled like a data.frame
5 > head(DGE.results)
6 > table(DGE.results$padj < 0.05)
7 > rownames(subset(DGE.results, padj < 0.05))

```

6.1.2 Exploratory plots following DGE analysis

Histograms Histograms are a simple and fast way of getting a feeling for how frequently certain values are present in a data set. A common example is a histogram of p-values (Figure 17).

```

1 > hist(DGE.results$pvalue,
2   col = "grey", border = "white", xlab = "", ylab = "",
3   main = "frequencies of p-values")

```

MA plot MA plots were originally developed for microarray visualization, but they are also useful for RNA-seq analyses. The MA plot provides a global view of the relationship between the expression change between conditions (log ratios, M), the average expression strength of the genes (average mean, A) and the ability of the algorithm to detect differential gene expression: genes that pass the significance threshold (adjusted p-value <0.05) are colored in red (Figure 17).

```

1 > plotMA(DGE.results, alpha = 0.05, main = "WT vs. SNF2 mutants",
2   ylim = c(-4,4))

```

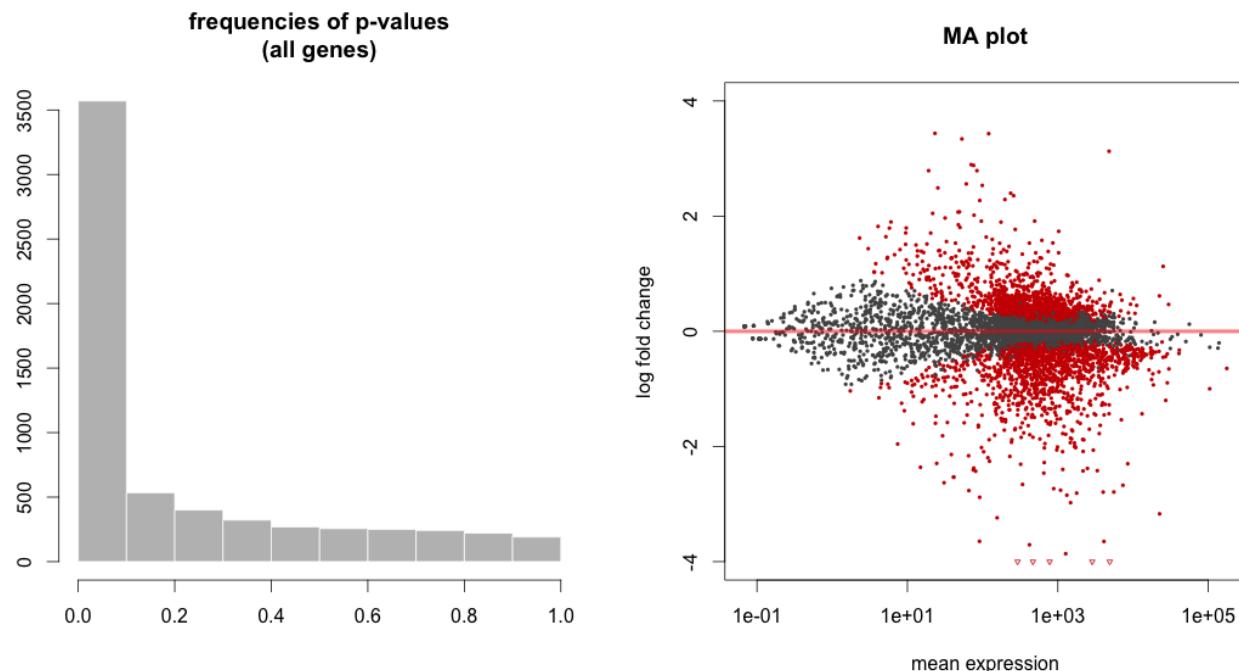


Figure 17: Left: Histogram of p -values for all genes tested for no differential expression between the two conditions, SNF2 and WT. Right: The MA plot shows the relationship between the expression change (M) and average expression strength (A); genes with adjusted p -values <0.05 are marked in red.

Heatmaps Heatmaps are a popular means to visualize the expression values across the individual samples. The following commands can be used to obtain heatmaps for *rlog*-normalized read counts for genes that show differential expression with adjusted p-values <0.05. Note that there are numerous functions for generating heatmaps in R; here we will use a fairly new function called `aheatmap()`, which is similar to `gplots::heatmap.2()` and `pheatmap::pheatmap()`.

```

1 # load the library with the aheatmap() function
2 > library(NMF)
3
4 # aheatmap needs a matrix of values, e.g., a matrix of DE genes with the
5 # transformed read counts for each replicate
6 # sort the results according to the adjusted p-value
7 > DGE.results.sorted <- DGE.results[order(DGE.results$padj), ]
8
9 # identify genes with the desired adjusted p-value cut-off
10 > DGEgenes <- rownames(subset(DGE.results.sorted, padj < 0.05))
11
12 # extract the normalized read counts for DE genes into a matrix
13 > hm.mat_DGEgenes <- log.norm.counts[DGEgenes, ]
14
15 # plot the normalized read counts of DE genes sorted by the adjusted p-value
16 > aheatmap(hm.mat_DGEgenes, Rowv = NA, Colv = NA)
17
18 # combine the heatmap with hierarchical clustering
19 > aheatmap(hm.mat_DGEgenes,
20   Rowv = TRUE, Colv = TRUE, # add dendrograms to rows and columns
21   distfun = "euclidean", hclustfun = "average")
22
23 # scale the read counts per gene to emphasize the sample-type-specific
24 # differences
25 > aheatmap(hm.mat_DGEgenes,
26   Rowv = TRUE, Colv = TRUE,
27   distfun = "euclidean", hclustfun = "average",
28   scale = "row") # values are transformed into distances from the center
29   # of the row-specific average: (actual value - mean of the group) /
30   # standard deviation

```

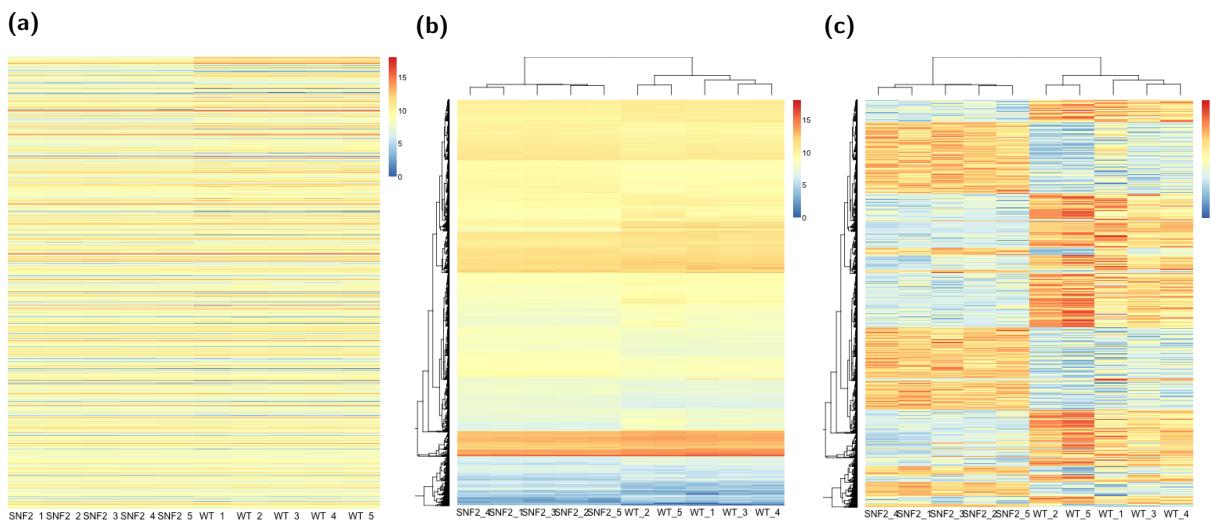


Figure 18: Heatmaps of *rlog*-transformed read counts for genes with adjusted *p*-values <0.05 in the DGE analysis. a) Genes sorted according to the adjusted *p*-values of the DGE analysis. b) Genes sorted according to hierarchical clustering. c) Same as for (b), but the read count values are scaled per row so that the colors actually represent *z*-scores rather than the underlying read counts.

Read counts of single genes An important sanity check of your data and the DGE analysis is to see whether genes about which you have prior knowledge behave as expected. For example, the samples named “SNF2” were generated from a mutant yeast strain where the *snf2* gene was deleted, so *snf2* should be among the most strongly downregulated genes in this DGE analysis.

To check whether *snf2* expression is absent in the mutant strain, we first need to map the ORF identifiers that we used for generating the read count matrix to the gene name so that we can retrieve the *rlog*-transformed read counts and the moderated \log_2 fold changes. There is more than one way to obtain annotation data, here we will use a data base that can be directly accessed from within R. The website <http://www.bioconductor.org/packages/3.1/data/annotation/> lists all annotation packages that are available through bioconductor. For our yeast samples, we will go with `org.Sc.sgd.db`.

```

1 > source("http://bioconductor.org/biocLite.R")
2 > biocLite("org.Sc.sgd.db")
3 > library(org.Sc.sgd.db)
4
5 # list the words (keytypes) that are available to query the annotation data
# base
6 > keytypes(org.Sc.sgd.db)
7
8 # list columns that can be retrieved from the annotation data base
9 > columns(org.Sc.sgd.db)
10
11 # make a batch retrieval for all DE genes
12 > anno <- select(org.Sc.sgd.db,
13                     keys = DGEgenes, keytype = "ORF",
14                     columns = c("SGD", "GENENAME", "CHR"))
15
16 # check whether SNF2 pops up among the top downregulated genes
17 > DGE.results.sorted_logFC <- DGE.results[order(DGE.results$log2FoldChange), ]
18 > DGEgenes_logFC <- rownames(subset(DGE.results.sorted_logFC, padj < 0.05))
19 > head(anno[match(DGEgenes_logFC, anno$ORF), ])
20
21 # find the ORF corresponding to SNF2
22 subset(anno, GENENAME == "SNF2")
23
24 # DESeq2 offers a wrapper function to plot read counts for single genes
25 > library(grDevices) # for italicizing the gene name
26 > plotCounts(dds = DESeq.ds,
27               gene = "YOR290C",
28               normalized = TRUE, transform = FALSE,
29               main = expression(atop("Expression of " *italic("snf2"), "(YOR290C)"))
23 )

```

While R offers myriad possibilities to perform downstream analyses on the lists of DE genes, you may also need to export the results into a simple text file that can be opened with other programs.

```

1 # merge the information of the DGE analysis with the information about the
# genes
2 > out.df <- merge(as.data.frame(DGE.results), anno,
3                     by.x = "row.names", by.y = "ORF")
4
5 # export all values for all genes into a tab-separated text file
6 > write.table(out.df, file = "DESeq2results_WT-vs-SNF2.txt",
7               sep = "\t", quote = FALSE, row.names = FALSE)

```

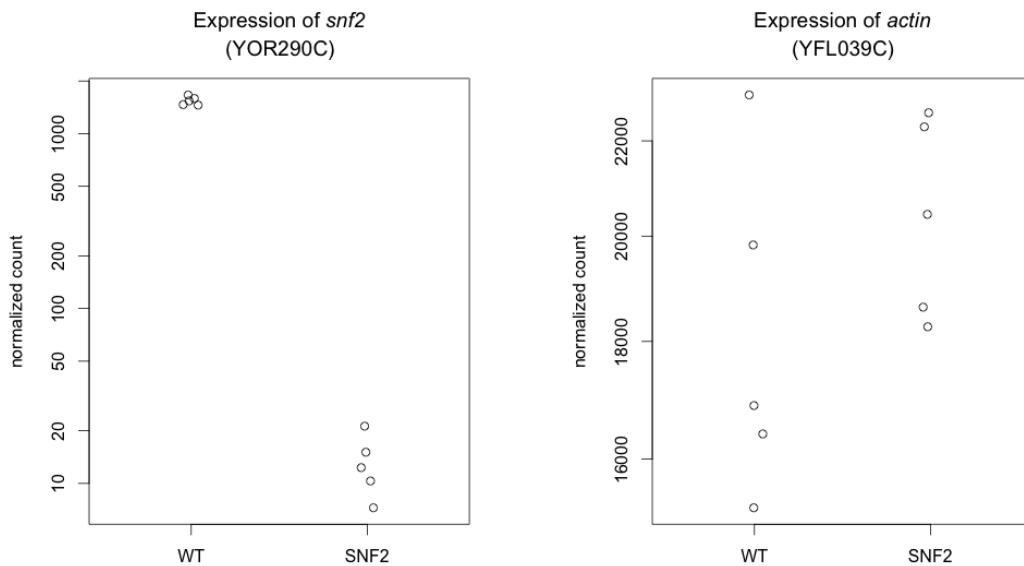


Figure 19: Read counts for *snf2* and *actin* in the replicates of both conditions.

- ?
- What is the difference between the moderated *log*-transformed values reported by either `rlog(DESeqDataSet)` or `results(DESeqDataSet)`?
 - Which analyses should be preferably performed on *log*-transformed read counts, which ones on *log* fold changes?

6.1.3 Exercise suggestions

The following exercises will help you to familiarize yourself with the handling of the data objects generated by `DESeq2`:

- Make a heatmap with the 50 genes that show the strongest change between the conditions. (the cut-off for the adjusted p-value should remain in place)
- Plot the read counts for a gene that is not changing between the two conditions, e.g., *actin*.
- Write a function that will plot the *rlog*-transformed values for a single gene, as in Figure 19. (Hint: aim for a boxplot via `plot()`, then add individual dots via `points()`)

6.1.4 edgeR

`edgeR` is very similar in spirit to `DESeq2`: both packages rely on the negative binomial distribution to model the raw read counts in a gene-wise manner while adjusting the dispersion estimates based on trends seen across all samples and genes (Table 5). The methods are, however, not identical, and results may vary. The following commands should help you perform a basic differential gene expression analysis, analogous to the one we have shown you for `DESeq2`, where five replicates from two conditions ("SNF2", "WT") were compared.

```

1 # install edgeR from bioconductor
2 > source("http://www.bioconductor.org/biocLite.R")
3 > biocLite("edgeR")
4 > library(edgeR)

```

`edgeR` requires a matrix of read counts where the row names = gene IDs and the column names = sample IDs. Thus, we can use the same object that we used for `DESeq2`.

```

1 > head(read.counts)

```

In addition, we need to specify the sample types, similarly to what we did for `DESeq2`.

```
1 > sample.type <- factor(c( rep("WT", 5), rep("SNF2", 5) ))
```

Now, `DGEList()` is the function that converts the count matrix into an `edgeR` object.

```
1 > edgeR.DGEList <- DGEList(counts = read.counts,
2                               group = sample.type)
3
4 # check the result
5 > head(edgeR.DGEList$counts)
6 > edgeR.DGEList$samples
```

To filter out genes with almost no coverage, we first plot a histogram of counts per million calculated by `edgeR`'s `cpm()` function.

```
1 # get an impression of the coverage across samples
2 > hist(log2(rowSums(cpm(edgeR.DGEList))))
```



```
3
4 # remove genes that do not have one count per million in at least 5 samples
5 # (adjust this to your sample!)
6 > keep <- rowSums(cpm(edgeR.DGEList) > 1) >= 5
7 > edgeR.DGEList <- edgeR.DGEList[keep,]
```



```
8
9 # recompute library sizes after filtering
10 edgeR.DGEList$samples$lib.size <- colSums(edgeR.DGEList$counts)
```

Calculate normalization factors for the library sizes:

```
1 > edgeR.DGEList <- calcNormFactors(edgeR.DGEList, method = "TMM")
2 > edgeR.DGEList$samples
```

To determine the differential expression in a gene-wise manner, `edgeR` first estimates the dispersion and subsequently tests whether the observed gene counts fit the respective negative binomial model. Note that the following commands are only appropriate if your data is based on an experiment with a single factor (e.g., mouse strain A vs. B; untreated cell culture vs. treated cell culture). For details on more complicated experimental set-ups, see the vignette of `edgeR` which can be found at the [bioconductor](#) website.

```
1 # first, estimate the dispersion for all read counts across all samples
2 > edgeR.DGEList <- estimateDisp(edgeR.DGEList, design.robust = TRUE)
3
4 # perform DE testing
5 > fit <- glmFit(edgeR.DGEList, sample.info)
6 > lrt <- glMLRT(fit, coef = 2)
```

In contrast to `DESeq`, `edgeR` does not produce any values similar to the *rlog*-transformed read count values. You will, however, get normalized *log₂* fold changes.

```
1 # extract results
2 > DGE.results_edgeR <- topTags(lrt, n = Inf, # to retrieve all genes
3                                   sort.by = "PValue", adjust.method = "BH")
```

6.1.5 limma-voom

Like `DESeq` and `edgeR`, `limma` starts with a matrix of raw read counts where each gene is represented by a row and the columns represent samples. `limma` assumes that rows with zero or very low counts have been removed. In addition, reads can be normalized for sequencing depth using `edgeR`'s `calcNormFactors()` function.

```

1 > library(limma)
2 > library(edgeR)
3
4 # use edgeR to normalize reads for sequencing depth
5 > edgeR.DGEList <- DGEList(counts = read.counts, group = sample.type)
6 > edgeR.DGEList <- calcNormFactors(edgeR.DGEList, method = "TMM")
7
8
9 # transform the count data to log2-counts-per-million and estimate the mean-
  variance relationship, which is used to compute weights for each count
10 > design <- model.matrix(~ sample.type)
11 > rownames(design) <- colnames(edgeR.DGEList)
12 > voomTransformed <- voom(edgeR.DGEList, design, plot=FALSE)
13
14 # fit a linear model for each gene
15 > voomed.fitted <- lmFit(voomTransformed, design = design)
16
17 # compute moderated t-statistics, moderated F-statistics, and log-odds of
  differential expression by empirical Bayes moderation of the standard errors
  towards a common value
18 > voomed.fitted <- eBayes(voomed.fitted)
19
20 # extract gene list with logFC and statistical measures
21 > colnames(design)
22 > topTable(voomed.fitted, coef = "sample.typeWT",
23             number = 10, adjust.method = "BH",
24             sort.by = "logFC")

```

6.2 Judging DGE results

Once you have obtained a table of genes that show signs of differential expression, you have reached one of the most important milestones of RNA-seq analysis! To evaluate how confident you can be in that list of DE genes, you should look at several aspects of the analyses you did and perform basic checks on your results:

1. Did the unsupervised clustering and PCA analyses reproduce the major trends of the initial experiment? For example, did replicates of the same condition cluster together and were they well separated from the replicates of the other condition(s)?
2. How well do different DGE programs agree on the final list of DE genes? You may want to consider performing downstream analyses only on the list of genes that were identified as DE by more than one tool.
3. Do the results of the DGE analysis agree with results from small-scale experiments? Can you reproduce qPCR results (and vice versa: can you reproduce the results of the DGE analysis with qPCR)?
4. How robust are the observed fold changes? Can they explain the effects you see on a phenotypic level?



If your RNA-seq results are suggesting expression changes that differ dramatically from everything you would have expected based on prior knowledge, you should be very cautious!

Most likely, this list will be the starting point for various downstream exploratory analyses, such as GO term enrichment analysis and correlation with other data sets. While the following list is by no means exhaustive, it may give you an idea of possible directions and tools to explore.

- RNA-seq with **more complex experimental designs** – the vignettes of `DESeq2` and `edgeR` have a good introduction and examples, e.g., for time course experiments, paired samples etc. as well as filtering genes with the `genefilter` package (Bourgon et al., 2010)

- GO term enrichment analysis
 - `goana()` function of the `limma` package
 - `goseq` R package
 - GOrilla (<http://cbl-gorilla.cs.technion.ac.il/>)
 - REVIGO (<http://revigo.irb.hr/>)
 - g:profiler (<http://biit.cs.ut.ee/gprofiler/>)
 - Gene Set Enrichment Analysis (GSEA; <https://www.broadinstitute.org/gsea/index.jsp>)
 - Ingenuity Pathway Analysis Studio (commercial software)
 - ...

7 Appendix

7.1 Additional tables

Table 6: Biases and artifacts of Illumina sequencing data. All high-throughput sequencing data will suffer from some degree of bias due to the biochemistry of the sequencing, the detection technique and bioinformatics processing.

Problem	Reasons	Solutions
Batch effects	<ul style="list-style-type: none">• variation in the sample processing (e.g., reagent batches, experimenters, pipetting accuracy)• flowcell inconsistencies• differences between sequencing runs (e.g., machine calibration)	<ul style="list-style-type: none">• appropriate experimental design (e.g., proper randomization)• samples of the same experiment should have similar quality and quantity• optimal experimental conditions (use of master mixes etc.)
Sequencing errors and errors in base calling	<ul style="list-style-type: none">• loss of synchronized base incorporation into the single molecules within one cluster of clonally amplified DNA fragments (phasing and pre-phasing)• mixed clusters• signal intensity decay over time due to unstable reagents• uneven signal intensities depending on the position on the flowcell• overlapping emission frequency spectra of the four fluorescently-labelled nucleotides	<ul style="list-style-type: none">• improvement of the sequencing chemistry and detection• optimized software for base calling• computational removal of bases with low base calling scores
GC bias and duplicate reads	<ul style="list-style-type: none">• GC-rich regions are preferably amplified by standard PCR• small fragments are preferably hybridized to the flowcell• low number of founder DNA fragments	<ul style="list-style-type: none">• optimizing cross-linking, sonication, and the mRNA enrichment to ensure that the majority of the transcriptome is present in the sample• limiting PCR cycles during library preparation to a minimum• computational correction for GC content and elimination of reads from identical DNA fragments
Copy number variations and mappability	<ul style="list-style-type: none">• incomplete genome assemblies• strain-specific differences from the reference assembly may lead to misrepresentation of individual loci• repetitiveness of genomes and shortness of sequencing reads hinder unique read alignment	<ul style="list-style-type: none">• longer sequencing reads• paired-end sequencing• exclusion of blacklisted regions that are known to attract artificially high read numbers (Kundaje, 2013)• computational correction for mappability

Table 7: FASTQC modules.

Plot title	Details	Warning (Failure)	Solution
Per base sequence quality	This plot is based on the quality scores reported and stored by the sequencing platform (Phred scores, see above). For Illumina sequencing, the reagents degrade throughout the sequencing run which is why the last bases tend to get worse scores.	Lower quartile for any base <10 (<5), or median for any base <25 (<20).	Trimming reads based on their average quality.
Per tile sequence quality	Check whether a part of the flowcell failed (e.g., due to bubbles or dirt on the flowcell). The plot shows the deviation from the average quality.	Any tile with a mean Phred score >2 (>5) for that base across all tiles.	The FASTQ file contains the tile IDs for each read which can be used to filter out reads from flawed tiles. Note though that variation in the Phred score of a flowcell is also a sign of overloading; tiles that are affected for only few cycles can be ignored.
Per sequence quality scores	Identify putative subsets of poor sequences.	Most frequent mean quality <27 (<20).	Quality trimming.
Per base sequence content	Expectation: all bases should be sequenced as often as they are represented in the genome. Reasons why this assumption may not hold: <ul style="list-style-type: none"> • random hexamer priming during library preparation • contamination (e.g., adapter dimers) • bisulfite treatment (loss of cytosines) 	Difference between A and T, or G and C $>10\%$ ($>20\%$) in any position.	If overrepresented sequences are the cause of a failure, these can be removed.
Per sequence GC content	The GC content of each read should be roughly normally distributed (maximum should correspond to the average GC content of the genome). Note that the reference shown here is also based on the supplied FASTQ file, therefore it will not be able to detect a global shift of GC content.	Deviations from the normal distribution for $>15\%$ (30%) of the reads	Sharp peaks on an otherwise smooth distribution usually indicate a specific contaminant that should also be reported as an overrepresented sequence. Broader peaks outside the expected distribution may represent contaminating sequences from a species with a different GC genome content.
Per base N content	Percentage of base calls at each position for which an N was called by the sequencer indicating lack of confidence to make a base call.	Any position with an N content of $>5\%$ (20%)	A common reason for large numbers of N is lack of library complexity, i.e., if all reads start with the same bases, the sequencer may not be able to differentiate them sufficiently.

Continued on next page

Table 7 – Continued from previous page

Plot title	Details	Warning (Failure)	Solution
Sequence length distribution	Determines the lengths of the sequences of the FASTQ file. The distribution of read sizes should be uniform for Illumina sequencing.	Warning is raised if not all sequences are the same length; failure is issued if any sequence has zero length.	For Illumina sequencing, different read lengths should only occur if some sort of bioinformatic trimming has happened prior to the FASTQC analysis.
Duplicate sequences	The sequenced library should contain a random and complete representation of the genome or transcriptome, i.e., most sequences should occur only once. High duplication rates are indicative of PCR over-amplification which may be caused by an initial lack of starting material. Note that this module only takes the first 100,000 sequences of each file into consideration.	Non-unique sequences make up more than 20% (50%) of all reads.	Unless you have reasons to expect sequences to be duplicated (i.e., specific enrichments of certain sequences), this plot is a strong indicator of suboptimal sample preparation. Duplication due to excessive sequencing will be reflected by flattened lines in the plot; PCR overamplification of low complexity libraries are indicated by sharp peaks towards the right-hand side of the plot.
Over-represented sequences	All sequences which make up more than 0.1% of the first 100,000 sequences are listed.	Any sequence representing >0.1% (1%) of the total.	Bioinformatic removal of contaminating sequences.
Adapter content	This module specifically searches for a set of known adapters. The plot itself shows a cumulative percentage count of the proportion of your library which has seen each of the adapter sequences at each position.	Any adapter present in more than 5% (10%) of all reads.	Bioinformatic removal of adapter sequences.
K-mer content	The number of each 7-mer at each position is counted; then a binomial test is used to identify significant deviations from an even coverage at all positions (only 2% of the whole library is analyzed and the results are extrapolated).	Any k-mer represented with a binomial p-value <0.01 ($<10^{-5}$)	Libraries which derive from random priming will nearly always show k-mer bias at the start of the library due to an incomplete sampling of the possible random primers. Always check the results of the adapter content and overrepresented sequences modules, too.

Table 8: Types of optional entries that can be stored in the header section of a **SAM** file following the format @<Section><Tag>:<Value>. The asterisk indicates which tags are required if a section is set.

Section	Tag	Description
HD (header)	VN*	File format version
	SO	Sort order (“unsorted”, “queryname”, or “coordinate”)
SQ (sequence dictionary)	SN*	Sequence name (e.g., chromosome name)
	LN*	Sequence length
	AS	Genome assembly identifier (e.g., mm9)
	M5	MD5 checksum of the sequence
	UR	URI of the sequence
RG (read group)	SP	Species
	ID*	Read group identifier (e.g. “Lane1”)
	SM*	Sample
	LB	Library
	DS	Description
	PU	Platform unit (e.g., lane identifier)
	PI	Predicted median insert size
	CN	Name of the sequencing center
	DT	Date of the sequencing run
	PL	Sequencing platform
PG (program)	ID*	Name of the program that produced the SAM file
	VN	Program version
	CL	Command line used to generate the SAM file
CO (comment)	Unstructured one-line comment lines (can be used multiple times)	

Table 9: Overview of RSeQC scripts. See the online documentation of RSeQC (<http://rseqc.sourceforge.net>) and Wang et al. (2012) for more details. Note that tools marked in red have been shown to return erroneous results (Hartley and Mullikin, 2015).

Script	Purpose
Basic read quality	
<code>read_duplication</code>	Determine read duplication rate, based on the sequence only (<code>output.dup.seq.DupRate.xls</code>) as well as on the alignment positions (<code>output.dup.pos.DupRate.xls</code>).
<code>read_hexamer</code>	Calculates hexamer frequencies from FASTA or FASTQ files. Similar to FASTQC's k-mer content analysis.
<code>read_GC</code>	Calculates % GC for all reads. Similar to FASTQC's GC distribution plot, the peak of the resulting histogram should coincide with the average GC content of the genome.
<code>read_NVC</code>	Calculates the nucleotide composition across the reads; similar to FASTQC's per base sequence content analysis.
<code>read_quality</code>	Calculates distributions for base qualities across reads; similar to FASTQC's per base sequence quality analysis.
Alignment QC	
<code>bam_stat</code>	Calculates reads mapping statistics for a BAM (or SAM) file. Note that uniquely mapped reads are determined on the basis of the mapping quality.
<code>clipping_profile</code>	Estimates clipping profile of RNA-seq reads from BAM or SAM file. This will fail if the aligner that was used does not support clipped mapping (CIGAR strings must have S operation).
<code>mismatch_profile</code>	Calculates distribution of mismatches along reads based on the MD tag.
<code>insertion_profile</code> , <code>deletion_profile</code>	Calculates the distribution of insertions or deletions along reads.
<code>read_distribution</code>	Calculates fractions of reads mapping to transcript features such as exons, 5'UTR, 3' UTR, introns.
<code>RPKM_saturation</code>	The full read set is sequentially downsampled and RPKMs are calculated for each subset. The resulting plot helps determine how well the RPKM values can be estimated. The visualized percent relative error is calculated as $\frac{ RPKM_{subsample} - RPKM_{max} }{RPKM_{max}} * 100$
RNA-seq-specific QC	
<code>geneBody_coverage</code>	Scales all transcripts to 100 bp, then calculates the coverage of each base. The read coverage should be uniform and ideally not show 5' or 3' bias since that would suggest problems with degraded RNA input or with cDNA synthesis.
<code>infer_experiment</code>	Speculates how RNA-seq sequencing was configured, i.e., PE or SR and strand-specificity. This is done by subsampling reads and comparing their genome coordinates and strands with those of the transcripts from the reference gene model. For non-strand-specific libraries, the strandedness of the reads and the transcripts should be independent. See http://rseqc.sourceforge.net/#infer-experiment-py for details.
<code>junction_annotation</code>	Compares detected splice junctions to the reference gene model, classifying them into 3 groups: annotated, complete novel, partial novel (only one of the splice sites is unannotated). The script differentiates between splice events (single read level) and splice junctions (multiple reads show the same splicing event).
<code>junction_saturation</code>	Similar concept to <code>RPKM_saturation</code> : splice junctions are detected for each sampled subset of reads. The detection of annotated splice sites should be saturated with the maximum coverage (= all supplied reads), otherwise alternative splicing analyses are not recommended because low abundance splice sites will not be detected.

Continued on next page

Table 9 – *Continued from previous page*

Script	Purpose
<code>tin</code>	Calculates the transcript integrity number (TIN) (not to be confused with the RNA integrity number, RIN, that is calculated before sequencing based on the $28S/18S$ ratio). TIN is calculated for each transcript and represents the fraction of the transcript with uniform read coverage.
	General read file handling and processing
<code>bam2fq</code>	Converts BAM to FASTQ.
<code>bam2wig</code>	Converts read positions stored in a BAM file into read coverage measures all types of RNA-seq data in BAM format into wiggle file.
<code>divide_bam</code>	Equally divides a given BAM file (m alignments) into n parts. Each part contains roughly m/n alignments that are randomly sampled from the entire alignment file.
<code>inner_distance</code>	Estimates the inner distance (or insert size) between two paired RNA reads (requires PE reads). The results should be consistent with the gel size selection during library preparation.
<code>overlay_bigwig</code>	Allows the manipulation of two BIGWIG files, e.g., to calculate the sum of coverages. See the <code>--action</code> option for all possible operations.
<code>normalize_bigwig</code>	Normalizes all samples to the same wigsum (= number of bases covered by <code>read length * total no. of reads</code>).
<code>split_bam,</code> <code>split_paired_bam</code>	Provided with a gene list and a BAM file, this module will split the original BAM file into 3 smaller ones: <ol style="list-style-type: none"> 1. <code>*.in.bam</code>: reads overlapping with regions specified in the gene list 2. <code>*.ex.bam</code>: reads that do not overlap with the supplied regions 3. <code>*.junk.bam</code>: reads that failed the QC or are unaligned
<code>RPKM_count</code>	Calculates the raw count and RPKM values for each exon, intron and mRNA region defined by the provided annotation file.

Table 10: Overview of QoRTs QC functions. See the online documentation of QoRTs (<http://hartleys.github.io/QoRTs/index.html>) and Hartley and Mullikin (2015) for more details.

QC Function	Purpose
Basic read quality	
<code>GCdistribution</code>	Calculates GC content distribution for all reads. Similar to FASTQC's GC distribution plot, the peak of the resulting histogram should coincide with the average GC content of the genome.
<code>NVC</code>	Calculates the nucleotide composition across the length of the reads; similar to FASTQC's per base sequence content analysis.
<code>QualityScoreDistribution</code>	Calculates distributions for base qualities across reads; similar to FASTQC's per base sequence quality analysis.
Alignment QC	
<code>chromCounts</code>	Calculates number of reads mapping to each category of chromosome (autosomes, allosomes, mtDNA).
<code>CigarOpDistribution</code>	Calculates rate of various CIGAR operations as a function of read length, including insertions, deletions, splicing, hard and soft clipping, padding, and alignment to reference. See Section 3.3.2 for details about CIGAR operations.
<code>GeneCalcs</code>	Calculates fractions of reads mapping to genomic features such as unique genes, UTRs, ambiguous genes, introns, and intergenic regions.
RNA-seq-specific QC	
<code>writeGeneBody</code>	Breaks up all genes into 40 equal-length counting bins and determines the number of reads that overlap with each counting bin. The read coverage should be uniform and ideally not show 5' or 3' bias since that would suggest degraded RNA input or problems with cDNA synthesis.
<code>writeGenewiseGeneBody</code>	Writes file containing gene-body distributions for each gene.
<code>StrandCheck</code>	Checks if the data is strand-specific by calculating the rate at which reads appear to follow the two possible library-type strandedness rules (fr-firststrand and fr-secondstrand, described by the CuffLinks documentation at http://cufflinks.ccb.umd.edu/manual.html#library).
<code>JunctionCalcs</code>	Calculates the number of novel and known splice junctions. Splice junctions are split into 4 groups, first by whether the splice junction appears in the gene annotation GTF (known vs. novel), and then by whether the splice junction has fewer or ≥ 4 reads covering it.
General read file handling and processing	
<code>InsertSize</code>	Estimates the inner distance (or insert size) between two paired RNA reads (requires PE reads). The results should be consistent with the gel size selection during library preparation.
<code>makeWiggles</code>	Converts read positions stored in a BAM file into wiggle files with 100-bp window size.
<code>makeJunctionBed</code>	Creates a BED file for splice-junction counts.
<code>writeGeneCounts</code>	Calculates raw number of reads mapping to each gene in the annotation file. Also creates a cumulative gene diversity plot, which shows the percent of the total proportion of reads as a function of the number of genes sequenced. This is useful as an indicator of whether a large proportion of the reads stem from a small number of genes (as a result of ribosomal RNA or hemoglobin contamination, for example).
<code>calcDetailedGeneCounts</code>	Calculates more detailed read counts for each gene, including the number of reads mapping to coding regions, UTR, and intronic regions.

Continued on next page

Table 10 – *Continued from previous page*

Function	Purpose
<code>writeBiotypeCounts</code>	Write a table listing read counts for each biotype, which is a classification of genes into broader categories (e.g., protein coding, pseudogene, processed pseudogene, miRNA, rRNA, scRNA, snoRNA, snRNA). Note that, in order for this function to succeed, the optional “gene.biotype” attribute is required to be present in the gene annotation file (GTF).
<code>FPKM</code>	Calculates FPKM values for each gene in the annotation file.

Table 11: Normalization methods for the comparison of gene read counts between different conditions. See, for example, Bullard et al. (2010) and Dillies et al. (2013) for comprehensive assessments of the individual methods.

Name	Details	Comment
Total Count	All read counts are divided by the total number of reads (library size) and multiplied by the mean total count across all samples.	<ul style="list-style-type: none"> biased by highly expressed genes cannot account for different RNA repertoire between samples poor detection sensitivity when benchmarked against qRT-PCR (Bullard et al., 2010)
Counts Per Million	Each gene count is divided by the corresponding library size (in millions).	<ul style="list-style-type: none"> see Total Count
DESeq's size factor	<ol style="list-style-type: none"> For each gene, the geometric mean of read counts across all samples is calculated. Every gene count is divided by the geometric mean. A sample's size factor is the median of these ratios (skipping the genes with a geometric mean of zero). 	<ul style="list-style-type: none"> the size factor is applied to all read counts of a sample more robust than total count normalization implemented by the <code>DESeq</code> R library (<code>estimateSizeFactors()</code> function), also available in <code>edgeR</code> (<code>calcNormFactors()</code> function with option <code>method = "RLE"</code>) details in Anders and Huber (2010)
Trimmed Mean of M-values (TMM)	<p>TMM is always calculated as the weighted mean of log ratios between two samples:</p> <ol style="list-style-type: none"> Calculate gene-wise \log_2 fold changes (= M-values): $M_g = \log_2\left(\frac{Y_{gk}}{N_k}\right)/\log_2\left(\frac{Y_{gk'}}{N_{k'}}\right)$ <p>where Y is the observed number of reads per gene g in library k and N is the total number of reads.</p> Trimming: removal of upper and lower 30%. Precision weighing: the inverse of the estimated variance is used to account for lower variance of genes with larger counts. 	<ul style="list-style-type: none"> the size factor is applied to every sample's library size; normalized read counts are obtained by dividing raw read counts by the TMM-adjusted library sizes more robust than total count normalization implemented in <code>edgeR</code> via <code>calcNormFactors()</code> with the default <code>method = "TMM"</code> details in Robinson and Oshlack (2010)
Upper quartile	<ol style="list-style-type: none"> Find the upper quartile value (top 75% read counts after removal of genes with 0 reads). Divide all read counts by this value. 	<ul style="list-style-type: none"> similar to total count normalization, thus it also suffers from a great influence of highly-expressed DE genes can be calculated with <code>edgeR</code>'s <code>calcNormFactors()</code> function (<code>method = "upperquartile"</code>)

Table 12: Normalization methods for the comparison of gene read counts within the same sample.

Name	Details	Comment
RPKM (reads per kilobase of exons per million mapped reads)	<p>1. For each gene, count the number of reads mapping to it (X_i).</p> <p>2. Divide that count by: the length of the gene, l_i, in base pairs divided by 1,000 multiplied by the total number of mapped reads, N, divided by 10^6.</p> $RPKM_i = \frac{X_i}{\left(\frac{l_i}{10^3}\right)\left(\frac{N}{10^6}\right)}$	<ul style="list-style-type: none"> introduces a bias in the per-gene variances, in particular for lowly expressed genes (Oshlack and Wakefield, 2009) implemented in edgeR's <code>rpk()</code> function
FPKM (fragments per kilobase...)	<p>1. Same as RPKM, but for paired-end reads:</p> <p>2. The number of fragments (defined by two reads each) is used.</p>	<ul style="list-style-type: none"> implemented in DESeq2's <code>fpkm()</code> function
TPM	<p>Instead of normalizing to the total library size, TPM represents the abundance of an individual gene i in relation to the abundances of the other transcripts (e.g., j) in the sample.</p> <p>1. For each gene, count the number of reads mapping to it and divide by its length in base pairs (= counts per base).</p> <p>2. Multiply that value by 1 divided by the sum of all counts per base of every gene.</p> <p>3. Multiply that number by 10^6.</p> $TPM_i = \frac{X_i}{l_i} * \frac{1}{\sum_j \frac{X_j}{l_k}} * 10^6$	<ul style="list-style-type: none"> details in Wagner et al. (2012)

7.2 Installing bioinformatics tools on a UNIX server

Tools are shown in order of usage throughout the script.

FASTQC

1. Download.

```
1 $ wget http://www.bioinformatics.babraham.ac.uk/projects/fastqc/fastqc_v0
    .11.3.zip
```

2. Unzip (specifying a folder to extract into) and make executable.

```
1 $ unzip fastqc_v0.11.5.zip -d ../software/
2 $ cd FastQC
3 $ chmod 755 fastq
```

samtools

1. Download source code and unzip it.

```
1 $ wget https://github.com/samtools/samtools/releases/download/1.3.1/
      samtools-1.3.1.tar.bz2 -O samtools-1.3.1.tar.bz2
2 $ tar jxf samtools-1.3.1.tar.bz2
3 $ cd samtools-1.2
```

2. Compile.

```
1 $ make
```

3. Check whether the tool is running.

```
1 $ ./samtools
```

4. Add the location where samtools was installed to your PATH variable; this way you will not need to specify the exact location everytime you want to run the tool.

```
1 $ export PATH=/zenodotus/abc/store/courses/2016_rnaseq/software/samtools
      -1.3.1:$PATH
```

RSeQC

1. Install Anaconda, a package which helps manage Python installs:

```
1 $ wget https://3230d63b5fc54e62148e-c95ac804525aac4b6dba79b00b39d1d3.ssl.
      cf1.rackcdn.com/Anaconda2-2.4.0-Linux-x86_64.sh
2 $ bash Anaconda2-2.4.0-Linux-x86_64.sh
```

You will need to accept the license terms, specify where to install anaconda, and specify whether you want anaconda's install location to be prepended to your PATH variable.

2. Make sure anaconda's install location is prepended to your PATH variable:

```
1 $ export PATH=/zenodotus/abc/store/courses/2016_rnaseq/software/anaconda2/
      bin:$PATH
```

3. Install RSeQC using anaconda's installer

```
1 $ pip install RSeQC
```

QoRTs

1. Install the R component:

```
1 > install.packages("http://hartleys.github.io/QoRTs/QoRTs_LATEST.tar.gz",
2                      repos=NULL,
3                      type="source");
```

2. Download the java component.

```
1 $ wget -O qorts.jar "https://github.com/hartleys/QoRTs/releases/download/v1
.1.8/QoRTs.jar"
```

STAR

1. Download.

```
1 $ wget https://github.com/alexdobin/STAR/archive/2.5.2a.tar.gz -O STAR_2
.5.2a.tar.gz
```

2. Unzip.

```
1 $ tar -zxf STAR_2.5.2a.tar.gz
```

3. Compile. From version 2.5.0 onwards, STAR requires the use of gcc version 4.7.0 or above. On the ABC servers, this can be achieved with the softlib library command `slchoose`.

```
1 $ cd STAR_2.5.2a
2 $ slchoose gcc 4.7.4 gcc4_64
3 $ make STAR
```

To run STAR:

```
1 $ ./bin/Linux_x86_64_static/STAR
```

UCSC tools aka Kent tools

1. Figure out which operating system version you have

```
1 $ uname -a
```

2. Download the already compiled binaries from the corresponding folder (shown here for the Linux server) and make them executable. The programs indicated here are the ones most commonly used for typical NGS analyses, but feel free to download more (or fewer) tools.

```
1 $ mkdir UCSCtools
2 $ cd UCSCtools
3 $ for PROGRAM in bedGraphToBigWig bedClip bigWigAverageOverBed
   bigWigCorrelate bigWigInfo bigWigSummary faToTwoBit fetchChromSizes
   genePredToBed gff3ToGenePred liftOver wigToBigWig
   do
5   wget http://hgdownload.soe.ucsc.edu/admin/exe/linux.x86_64/${PROGRAM}
6   chmod +x ${PROGRAM}
7   done
```

featureCounts (subread package)

1. Download.

```
1 $ wget https://sourceforge.net/projects/subread/files/subread-1.5.0-p3/
    subread-1.5.0-p3-Linux-x86_64.tar.gz
```

2. Unzip.

```
1 $ tar -zxf subread-1.5.0-p3-Linux-x86_64.tar.gz
```

R

1. Download.

```
1 $ wget --no-check-certificate https://cran.r-project.org/src/base/R-3/R-
    -3.2.1.tar.gz
```

2. Unzip.

```
1 $ tar zxvf R-3.2.1.tar.gz
```

3. Compile. You can use the `--prefix` option to specify the folder where you would like to install the program.

```
1 $ cd R-3.2.1/
2 $ ./configure --prefix=<path to folder of choice>
3 $ make
4 $ make install
```

References

- Anders S and Huber W. DESeq: Differential expression analysis for sequence count data. *Genome Biology*, **11**:R106, 2010. doi:10.1186/gb-2010-11-10-r106.
- Anders S, Pyl PT, and Huber W. HTSeq - A Python framework to work with high-throughput sequencing data. *Bioinformatics*, **31**(2):166–169, 2014. doi:10.1093/bioinformatics/btu638.
- Anders S, Reyes A, and Huber W. Detecting differential usage of exons from RNA-seq data. *Genome Research*, **22**(10):2008–2017, 2012. doi:10.1101/gr.133744.111.
- Bourgon R, Gentleman R, and Huber W. Independent filtering increases detection power for high-throughput experiments. *PNAS*, **107**(21):9546–9551, 2010. doi:10.1073/pnas.0914005107.
- Bray NL, Pimentel H, Melsted P, and Pachter L. Near-optimal probabilistic RNA-seq quantification. *Nat Biotech*, **34**(5):525–527, 2016. doi:10.1038/nbt.3519.
- Bullard JH, Purdom E, Hansen KD, and Dudoit S. Evaluation of statistical methods for normalization and differential expression in mRNA-Seq experiments. *BMC Bioinformatics*, **11**:94, 2010. doi:10.1186/1471-2105-11-94.
- Cathala G, Savouret JF, Mendez B, West BL, Karin M, Martial JA, and Baxter JD. A method for isolation of intact, translationally active ribonucleic acid. *DNA*, **2**(4):329–335, 1983. doi:10.1089/dna.1983.2.329.
- Ching T, Huang S, and Garmire LX. Power analysis and sample size estimation for RNA-Seq differential expression. *RNA*, pp. rna.046011.114–, 2014. doi:10.1261/rna.046011.114.
- Dapas M, Kandpal M, Bi Y, and Davuluri RV. Comparative evaluation of isoform-level gene expression estimation algorithms for RNA-seq and exon-array platforms. *Briefings in Bioinformatics*, (Oct 2015), 2016. doi:10.1093/bib/bbw016.
- Dillies MA, Rau A, Aubert J, Hennequet-Antier C, Jeanmougin M, Servant N, Keime C, Marot NS, Castel D, Estelle J, Guernec G, Jagla B, Jouneau L, Laloë D, Le Gall C, Schäffer B, Le Crom S, Guedj M, and Jaffrézic F. A comprehensive evaluation of normalization methods for Illumina high-throughput RNA sequencing data analysis. *Briefings in Bioinformatics*, **14**(6):671–683, 2013. doi:10.1093/bib/bbs046.
- Dobin A, Davis CA, Schlesinger F, Drenkow J, Zaleski C, Jha S, Batut P, Chaisson M, and Gingeras TR. STAR: Ultrafast universal RNA-seq aligner. *Bioinformatics*, **29**(1):15–21, 2013. doi:10.1093/bioinformatics/bts635.
- ENCODE. Standards, Guidelines and Best Practices for RNA-Seq, 2011. URL https://genome.ucsc.edu/ENCODE/protocols/dataStandards/ENCODE_RNAseq_Standards_V1.0.pdf.
- Engström PG, Steijger T, Sipos B, Grant GR, Kahles A, Rätsch G, Goldman N, Hubbard TJ, Harrow J, Guigó R, and Bertone P. Systematic evaluation of spliced alignment programs for RNA-seq data. *Nature Methods*, **10**(12):1185–1191, 2013. doi:10.1038/nmeth.2722.
- Farrell R. *RNA methodologies laboratory guide for isolation and characterization*. Elsevier/Academic Press, Amsterdam Boston, 2010.
- Feng H, Zhang X, and Zhang C. mRIN for direct assessment of genome-wide and gene-specific mRNA integrity from large-scale RNA-sequencing data. *Nature Communications*, **6**(May):7816, 2015. doi:10.1038/ncomms8816.
- Genohub. Depth of Coverage (RNA), 2015. URL <https://genohub.com/next-generation-sequencing-guide/#depth2>.
- Germain PL, Vitriolo A, Adamo A, Laise P, Das V, and Testa G. RNAontheBENCH: computational and empirical resources for benchmarking RNAseq quantification and differential expression methods. *Nucleic Acids Research*, **44**(11), 2016. doi:10.1093/nar/gkw448.
- Gibbons FD and Roth FP. Judging the quality of gene expression-based clustering methods using gene annotation. *Genome Research*, **12**(10):1574–1581, 2002. doi:10.1101/gr.397002.
- Gierliński M, Cole C, Schofield P, Schurch NJ, Sherstnev A, Singh V, Wrobel N, Gharbi K, Simpson G, Owen-Hughes T, Blaxter M, and Barton GJ. Statistical models for RNA-seq data derived from a two-condition 48-replicate experiment. *Bioinformatics*, **31**(22):1–15, 2015. doi:10.1093/bioinformatics/btv425.
- Hartley SW and Mullikin JC. QoRTs: a comprehensive toolset for quality control and data processing of RNA-Seq experiments. *BMC Bioinformatics*, **16**(1):224, 2015. doi:10.1186/s12859-015-0670-5.
- Head SR, Kiyomi Komori H, LaMere SA, Whisenant T, Van Nieuwerburgh F, Salomon DR, and Ordoukhani P. Library construction for next-generation sequencing: Overviews and challenges. *BioTechniques*,

- 56(2):61–77, 2014. doi:10.2144/000114133.
- Illumina. RNA-Seq Data Comparison with Gene Expression Microarrays, 2011. URL http://www.europeanpharmaceuticalreview.com/wp-content/uploads/Illumina_whitepaper.pdf.
- Jänes J, Hu F, Lewin A, and Turro E. A comparative study of RNA-seq analysis strategies. *Briefings in Bioinformatics*, (January):1–9, 2015. doi:10.1093/bib/bbv007.
- Kanitz A, Gypas F, Gruber AJ, Gruber AR, Martin G, and Zavolan M. Comparative assessment of methods for the computational inference of transcript isoform abundance from RNA-seq data. *Genome Biology*, 16(1), 2015. doi:10.1186/s13059-015-0702-5.
- Katz Y, Wang ET, Airoldi EM, and Burge CB. Analysis and design of RNA sequencing experiments for identifying isoform regulation. *Nature Methods*, 7(12):1009–1015, 2010. doi:10.1038/nmeth.1528.
- Katz Y, Wang ET, Silterra J, Schwartz S, Wong B, Thorvaldsdóttir H, Robinson JT, Mesirov JP, Airoldi EM, and Burge CB. Quantitative visualization of alternative exon expression from RNA-seq data. *Bioinformatics (Oxford, England)*, 2015. doi:10.1093/bioinformatics/btv034.
- Kim D, Pertea G, Trapnell C, Pimentel H, Kelley R, and Salzberg SL. TopHat2: accurate alignment of transcriptomes in the presence of insertions, deletions and gene fusions. *Genome Biology*, 14(4):R36, 2013. doi:10.1186/gb-2013-14-4-r36.
- Kundaje A. A comprehensive collection of signal artifact blacklist regions in the human genome. Tech. rep., 2013. URL <https://sites.google.com/site/anshulkundaje/projects/blacklists>.
- Levin JZ, Yassour M, Adiconis X, Nusbaum C, Thompson DA, Friedman N, Gnirke A, and Regev A. Comprehensive comparative analysis of strand-specific RNA sequencing methods. *Nature Methods*, 7(9):709–715, 2010. doi:10.1038/nmeth.1491.
- Li B and Dewey C. RSEM: accurate transcript quantification from RNA-Seq data with or without a reference genome. *BMC Bioinformatics*, 12(1):323, 2011. doi:10.1186/1471-2105-12-323.
- Li H, Handsaker B, Wysoker A, Fennell T, Ruan J, Homer N, Marth G, Abecasis G, and Durbin R. The Sequence Alignment/Map format and SAMtools. *Bioinformatics*, 25(16):2078–9, 2009. doi:10.1093/bioinformatics/btp352.
- Liao Y, Smyth GK, and Shi W. featureCounts: an efficient general purpose program for assigning sequence reads to genomic features. *Bioinformatics*, 30(7):923–30, 2014. doi:10.1093/bioinformatics/btt656.
- Liu Y, Zhou J, and White KP. RNA-seq differential expression studies: more sequence or more replication? *Bioinformatics*, 30(3):301–4, 2014. doi:10.1093/bioinformatics/btt688.
- Love MI, Huber W, and Anders S. Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2. *Genome Biology*, 15(12):550, 2014. doi:10.1186/s13059-014-0550-8.
- Moreton J, Izquierdo A, and Emes RD. Assembly, assessment and availability of de novo generated eukaryotic transcriptomes. *Frontiers in Genetics*, 6(January):1–9, 2015. doi:10.3389/fgene.2015.00361.
- Nookaei I, Papini M, Pornputtapong N, Scalcinati G, Fagerberg L, Uhlén M, and Nielsen J. A comprehensive comparison of RNA-Seq-based transcriptome analysis from reads to differential gene expression and cross-comparison with microarrays: A case study in *Saccharomyces cerevisiae*. *Nucleic Acids Research*, 40(20):10084–10097, 2012. doi:10.1093/nar/gks804.
- NuGEN. Detection of Genomic DNA in Human RNA Samples for RNA-Seq, 2013. URL http://www.nugen.com/sites/default/files/M01355_v1-TechnicalReport, DetectionofGenomicDNAinHumanRNASamplesforRNA-Seq.pdf.
- Oshlack A and Wakefield MJ. Transcript length bias in RNA-seq data confounds systems biology. *Biology Direct*, 4:14, 2009. doi:10.1186/1745-6150-4-14.
- Patro R, Guddal G, and Kingsford C. Accurate, fast, and model-aware transcript expression quantification with Salmon, 2015. doi:10.1101/021592.
- Patro R, Mount SM, and Kingsford C. Sailfish enables alignment-free isoform quantification from RNA-seq reads using lightweight algorithms. *Nature Biotechnology*, 32(5):462–464, 2014. doi:10.1038/nbt.2862.
- Pimentel HJ. RNA-Seq Methods and Algorithms (Part III Quantification), . URL https://www.youtube.com/watch?v=ztyjiCCt_1M.
- Pimentel HJ, Bray N, Puente S, Melsted P, and Pachter L. Differential analysis of RNA-Seq incorporating quantification uncertainty. *bioRxiv*, p. 058164, 2016. doi:10.1101/058164.
- Rapaport F, Khanin R, Liang Y, Pirun M, Krek A, Zumbo P, Mason CE, Soccia ND, and Betel D. Comprehensive evaluation of differential gene expression analysis methods for RNA-seq data. *Genome Biology*,

- 14(9):R95, 2013. doi:10.1186/gb-2013-14-9-r95.
- Risso D, Ngai J, Speed TP, and Dudoit S. Normalization of RNA-seq data using factor analysis of control genes or samples. *Nature Biotechnology*, pp. 1–10, 2014. doi:10.1038/nbt.2931.
- Ritchie ME, Phipson B, Wu D, Hu Y, Law CW, Shi W, and Smyth GK. limma powers differential expression analyses for RNA-sequencing and microarray studies. *Nucleic Acids Research*, 43(7):e47–, 2015. doi:10.1093/nar/gkv007.
- Roberts A and Pachter L. Streaming fragment assignment for real-time analysis of sequencing experiments. *Nature Methods*, 10(1):71–73, 2013. doi:10.1038/nmeth.2251.
- Robinson MD, McCarthy DJ, and Smyth GK. edgeR: a Bioconductor package for differential expression analysis of digital gene expression data. *Bioinformatics*, 26(1), 2010. doi:10.1093/bioinformatics/btp616.
- Robinson MD and Oshlack A. A scaling normalization method for differential expression analysis of RNA-seq data. *Genome Biology*, 11(3):R25, 2010. doi:10.1186/gb-2010-11-3-r25.
- Schurch NJ, Schofield P, Gierliński M, Cole C, Sherstnev A, Singh V, Wrobel N, Gharbi K, Simpson GG, Owen-Hughes T, Blaxter M, and Barton GJ. How many biological replicates are needed in an RNA-seq experiment and which differential expression tool should you use? *RNA*, pp. 1–13, 2016. doi:10.1261/rna.053959.115.
- Schurch NJ, Schofield P, Gierliński M, Cole C, Simpson GG, Hughes TO, Blaxter M, and Barton GJ. Evaluation of tools for differential gene expression analysis by RNA-seq on a 48 biological replicate experiment. *ArXiv e-prints*, 2015. URL <http://arxiv.org/abs/1505.02017>.
- Seyednasrollah F, Laiho A, and Elo LL. Comparison of software packages for detecting differential expression in RNA-seq studies. *Briefings in Bioinformatics*, 16(1):59–70, 2015. doi:10.1093/bib/bbt086.
- Shen S, Park JW, Lu Zx, Lin L, Henry MD, Wu YN, Zhou Q, and Xing Y. rMATS: Robust and flexible detection of differential alternative splicing from replicate RNA-Seq data. *PNAS*, 2014. doi:10.1073/pnas.1419161111.
- Sims D, Sudbery I, Ilott NE, Heger A, and Ponting CP. Sequencing depth and coverage: key considerations in genomic analyses. *Nature Reviews Genetics*, 15(2):121–32, 2014. doi:10.1038/nrg3642.
- Soneson C and Delorenzi M. A comparison of methods for differential expression analysis of RNA-seq data. *BMC Bioinformatics*, 14(1):91, 2013. doi:10.1186/1471-2105-14-91.
- Soneson C, Love MI, and Robinson MD. Differential analyses for RNA-seq: transcript-level estimates improve gene-level inferences. *F1000Research*, 4(0):1521, 2015. doi:10.12688/f1000research.7563.2.
- Soneson C, Matthes KL, Nowicka M, Law CW, and Robinson MD. Isoform prefiltering improves performance of count-based methods for analysis of differential transcript usage. *Genome Biology*, 17(1):12, 2016. doi:10.1186/s13059-015-0862-3.
- Sultan M, Amstislavskiy V, Risch T, Schuette M, Dökel S, Ralser M, Balzereit D, Lehrach H, and Yaspo ML. Influence of RNA extraction methods and library selection schemes on RNA-seq data. *BMC Genomics*, 15(1):675, 2014. doi:10.1186/1471-2164-15-675.
- Teng M, Love MI, Davis CA, Djebali S, Dobin A, Graveley BR, Li S, Mason CE, Olson S, Pervouchine D, Sloan CA, Wei X, Zhan L, and Irizarry RA. A benchmark for RNA-seq quantification pipelines. *Genome Biology*, 17(1), 2016. doi:10.1186/s13059-016-0940-1.
- Trapnell C, Hendrickson DG, Sauvageau M, Goff L, Rinn JL, and Pachter L. Differential analysis of gene regulation at transcript resolution with RNA-seq. *Nat Biotech*, 31(1):46–53, 2013. doi:10.1038/nbt.2450.
- Trapnell C, Roberts A, Goff L, Pertea G, Kim D, Kelley DR, Pimentel H, Salzberg SL, Rinn JL, and Pachter L. Differential gene and transcript expression analysis of RNA-seq experiments with TopHat and Cufflinks. *Nature Protocols*, 7(3):562–78, 2012. doi:10.1038/nprot.2012.016.
- Wagner GP, Kin K, and Lynch VJ. Measurement of mRNA abundance using RNA-seq data: RPKM measure is inconsistent among samples. *Theory in Biosciences*, 131(4):281–285, 2012. doi:10.1007/s12064-012-0162-3.
- Wang L, Wang S, and Li W. RSeQC: Quality control of RNA-seq experiments. *Bioinformatics*, 28(16):2184–2185, 2012. doi:10.1093/bioinformatics/bts356.
- Zeng W and Mortazavi A. Technical considerations for functional sequencing assays. *Nature Immunology*, 13(9):802–807, 2012. doi:10.1038/ni.2407.
- Zhao S, Fung-Leung WP, Bittner A, Ngo K, and Liu X. Comparison of RNA-Seq and microarray in transcriptome profiling of activated T cells. *PLoS ONE*, 9(1), 2014. doi:10.1371/journal.pone.0078644.